

Advanced Cloud Computing

Hadoop

Wei Wang
CSE@HKUST
Spring 2025



THE DEPARTMENT OF

COMPUTER SCIENCE & ENGINEERING

計算機科學及工程學系

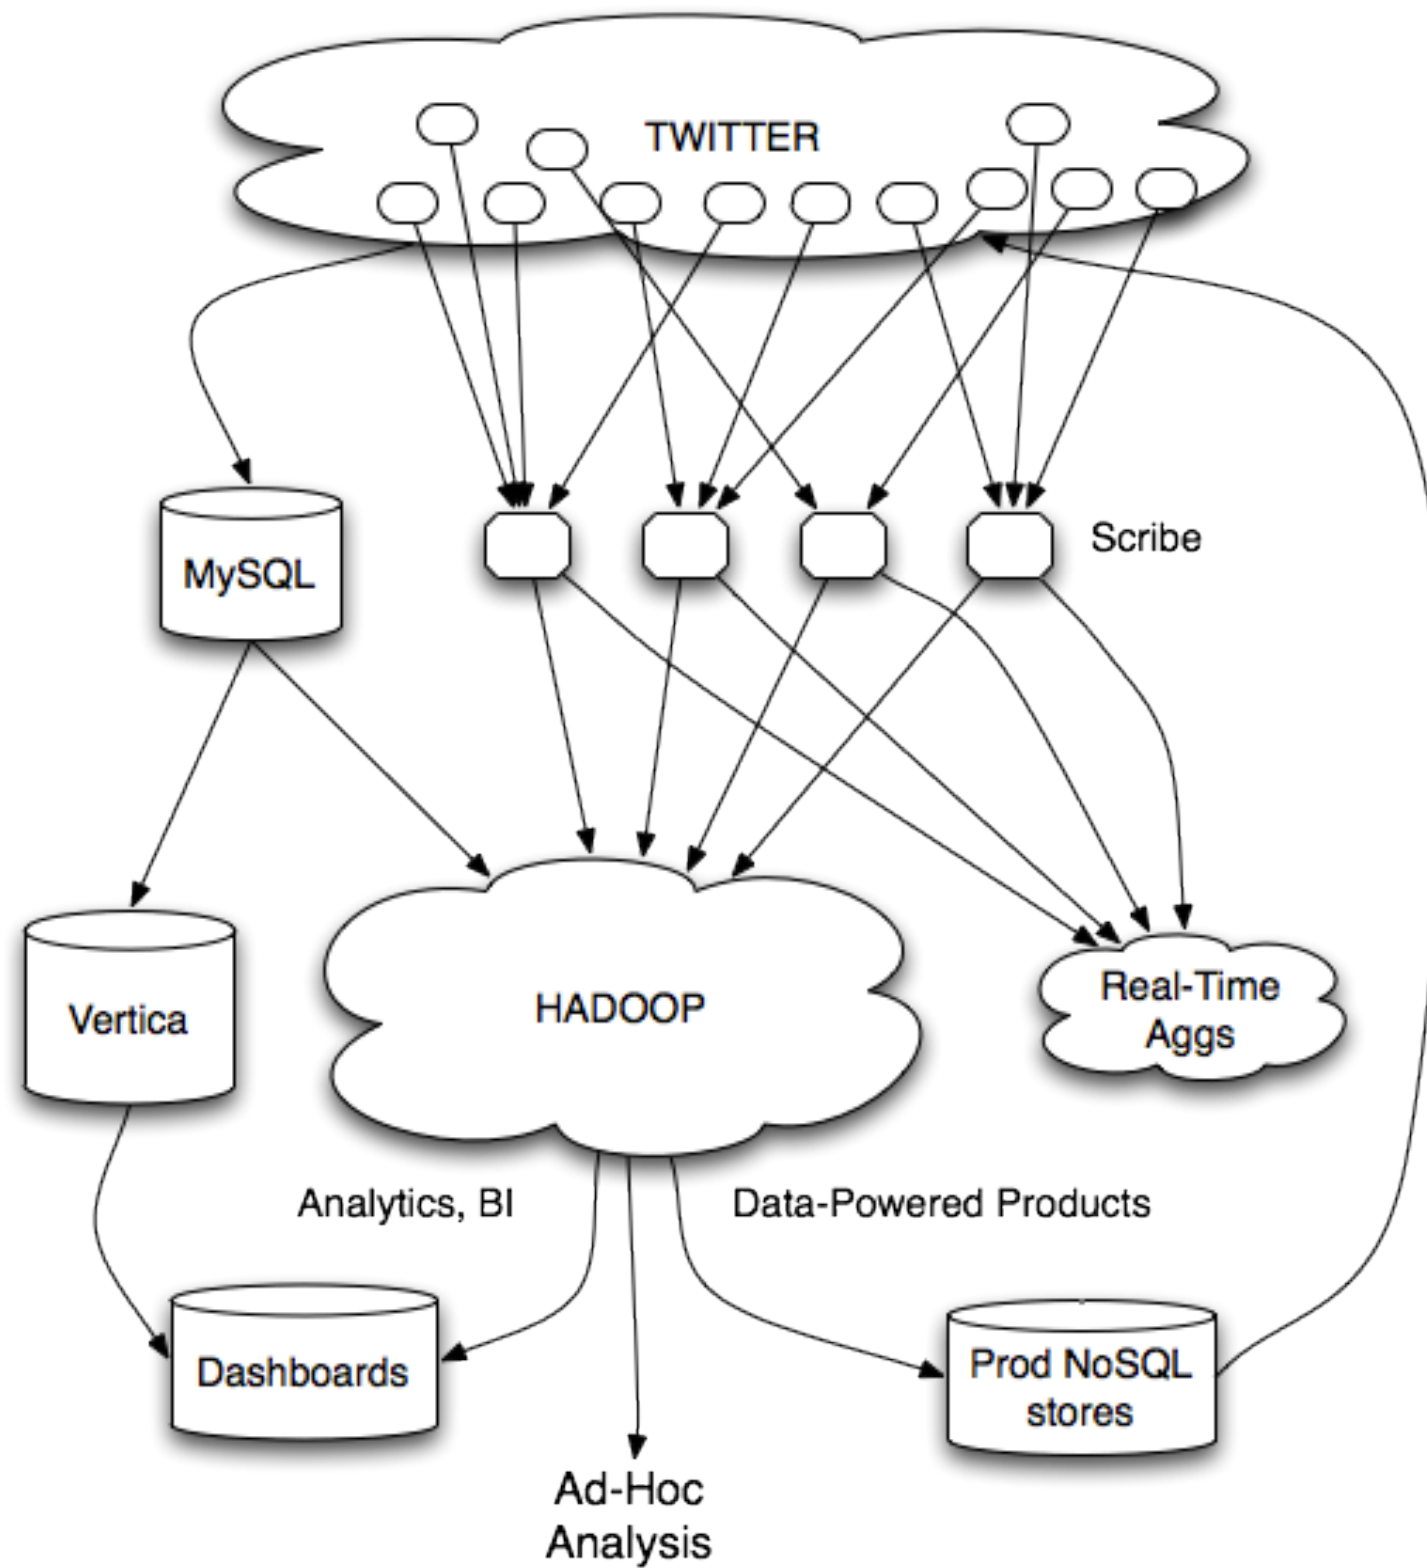
You say, “tomato...”

Google's proprietary implementation	Open-source equivalent
MapReduce	Hadoop
GFS	HDFS
BigTable	HBase
Chubby	ZooKeeper



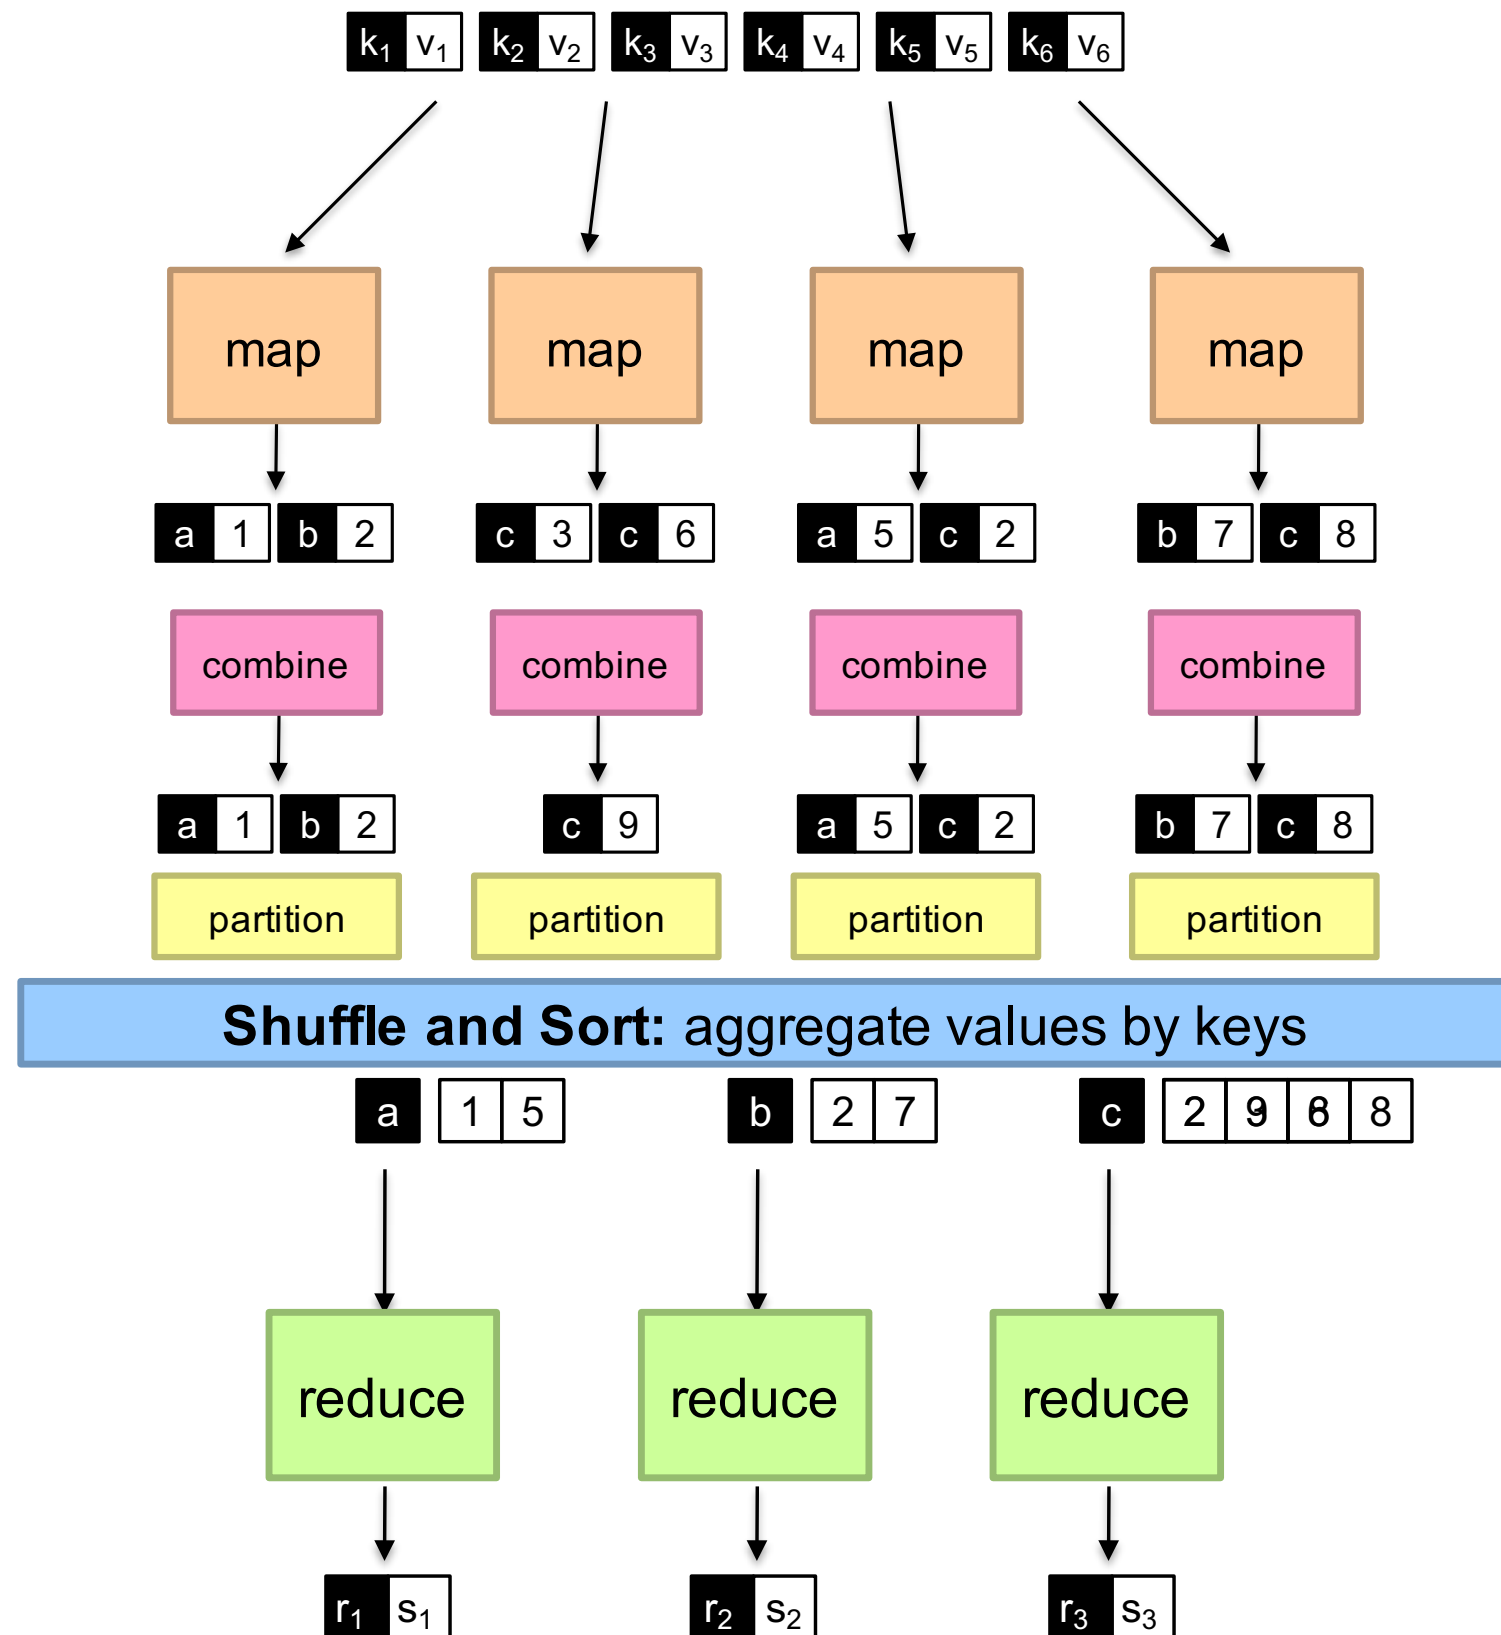
An open-source implementation of MapReduce in Java

- ▶ development led by Yahoo!, now an Apache project
- ▶ used in production at Yahoo!, Facebook, Twitter, LinkedIn, Netflix, ...
- ▶ the *de facto* big data processing platform
- ▶ large and expanding software ecosystem
- ▶ lots of custom research implementations



Twitter's data warehousing architecture





Hadoop has two versions of API

org.apache.hadoop.mapreduce

org.apache.hadoop.mapred



Basic Hadoop API

Mapper

- ▶ **void setup(Mapper.Context context)**
called once at the start of the task
- ▶ **void map(K key, V value, Mapper.Context context)**
called once for each key/value pair in the input split
- ▶ **void cleanup(Mapper.Context context)**
called once at the end of the task

Basic Hadoop API

Reducer/Combiner

- ▶ **void setup(Reducer.Context context)**
called once at the start of the task
- ▶ **void reduce(K key, Iterable<V> value, Reducer.Context context)**
called once for each key
- ▶ **void cleanup(Reducer.Context context)**
called once at the end of the task

Basic Hadoop API

Partitioner

- ▶ `int getPartition(K key, V value, int numPartitions)`
get the partition number given total number of partitions

Hadoop terminology

Job

- ▶ a packaged Hadoop program for submission to cluster
- ▶ need to specify input and output paths
- ▶ need to specify input and output formats
- ▶ need to specify mapper, reducer, combiner, partitioner
- ▶ need to specify intermediate/final key/value classes
- ▶ need to specify number of reducers (but not mappers, why?)

Hadoop terminology

Task

- ▶ an execution of a mapper or a reducer on a slice of data, a.k.a., Task-In-Progress (TIP)

Task attempt

- ▶ a particular instance of an attempt to execute a task on a machine
- ▶ a particular task will be attempted at least once, possibly more times if it crashes

Terminology example

Running a WordCount across 20 files is one **job**

20 files to be mapped imply 20 **map tasks** + some number of **reduce tasks**

At least 20 map **task attempts** will be performed

- ▶ more if a machine crashes, etc.

Data types in Hadoop

Writable



WritableComparable



IntWritable

LongWritable

Text

...

SequenceFile

Defines a **de/serialization** protocol. Every data type in Hadoop is a **Writable**.

Defines a sort order. All keys must be of this type (but not values).

Concrete classes for different data types.

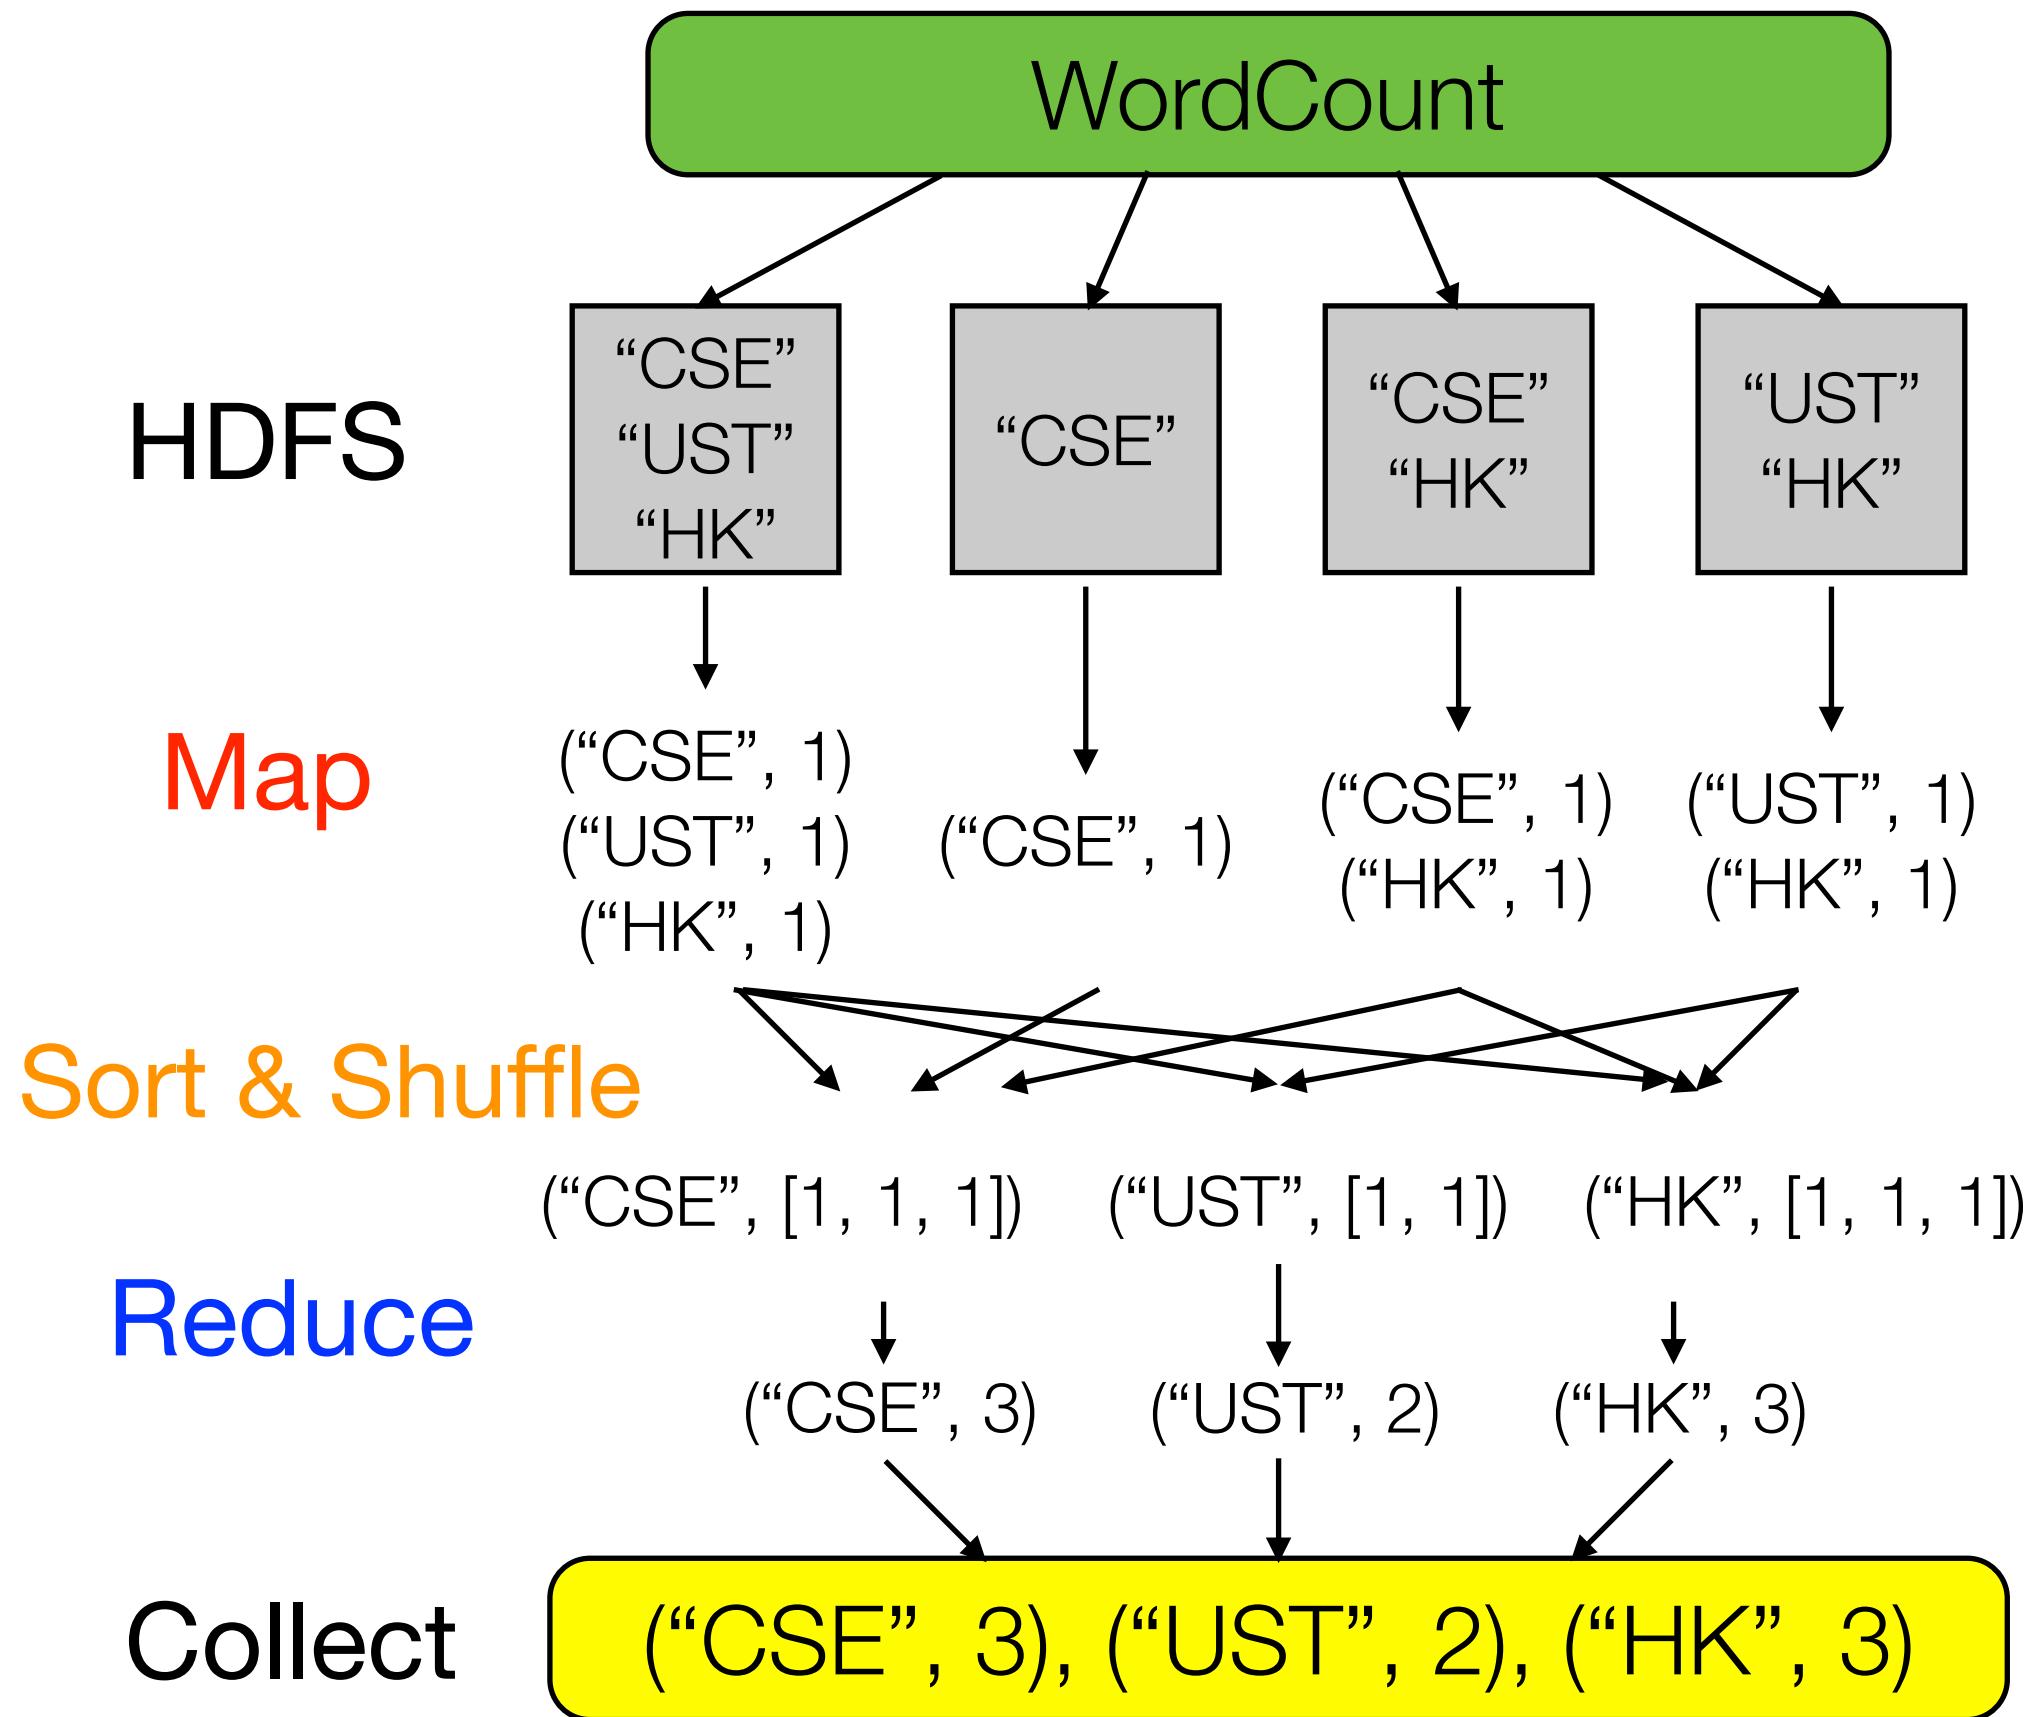
Binary encoded of a sequence of key/value pairs

Why not use Java Serialization?

Java comes with its own serialization mechanism. Why reinvent the wheel?

“Because it (Java Serialization) looked **big** and **hairy** and I thought we needed something **lean** and **mean**, where we had precise control over exactly how objects are written and read, since that is central to Hadoop.”

—Dough Cutting



WordCount: pseudo code

Map(String docid, String text):

for each word w in text:

Emit(w, 1);

Reduce(String term, Iterator<Int> values):

int sum = 0;

for each v in values:

sum += v;

Emit(term, sum)

WordCount: Mapper

Custom mapper inherits from the Mapper class:
 $\text{map}(k, v) \rightarrow [\langle k', v' \rangle]$

```
private static class MyMapper
    extends Mapper<LongWritable, Text, Text, IntWritable> {
    // avoid creating objects on the fly
    private final static IntWritable ONE = new IntWritable(1);
    private final static Text WORD = new Text();

    @Override    // key = byte offset of each line; value = line text
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String line = ((Text) value).toString();
        String[] words = line.trim().split("\\s+");
        for (String w: words) {
            WORD.set(w);
            context.write(WORD, ONE);
        }
    }
}
```

WordCount: Reducer

Custom reducer inherits from the Reducer class:
 $\text{reduce}(k, [v]) \rightarrow \langle k', v' \rangle$

```
private static class MyReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {
    // avoid creating objects on the fly
    private final static IntWritable SUM = new IntWritable();

    @Override // values with the same key are reduced together
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable v: values) {
            sum += v.get();
        }
        SUM.set(sum);
        context.write(key, SUM);
    }
}
```

Three Gotchas

Avoid object creation whenever possible

- ▶ reuse **Writable** objects, change the payload

Execution framework reuses value object in reducer

Passing parameters via class statics

Configure the job and run it

```
// Create and configure a MapReduce job
```

```
Configuration conf = getConf();  
Job job = Job.getInstance(conf);  
job.setJobName("Word Count");  
job.setJarByClass(WordCount.class);
```

```
job.setNumReduceTasks(reduceTasks); // Optional
```

```
// Specify inputs, outputs
```

```
FileInputFormat.setInputPaths(job, new Path(inputPath));  
FileOutputFormat.setOutputPath(job, new Path(outputPath));
```

```
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);
```

```
// Specify mapper, combiner, and reducer class
```

```
job.setMapperClass(WordCountMapper.class);  
job.setCombinerClass(WordCountReducer.class);  
job.setReducerClass(WordCountReducer.class);
```

```
// Run job and wait for its completion
```

```
System.exit(job.waitForCompletion(true) ? 0 : 1);
```


Try it in Assignment-3

Sometimes, you may need complex data types, e.g., key as a pair of strings

Complex data types

The easiest way:

- ▶ encode it as **Text**, e.g., $(a, b) = \text{"a:b"}$
- ▶ use regular expressions to parse and extract data
- ▶ works but pretty "hack-ish" and hard to read

Complex data types

The standard (and hard) way:

- ▶ define a custom implementation of **Writable(Comparable)**
- ▶ must implement: **readFields**, **write**, (**compareTo**)
- ▶ computationally efficient, but slow for rapid prototyping
- ▶ implement **WritableComparator** hook for performance

Somewhere in the middle

- ▶ third-party implementations: there are plenty of them!

Example: PairOfStrings

Anatomy of Hadoop

Basic cluster components

One of each cluster:

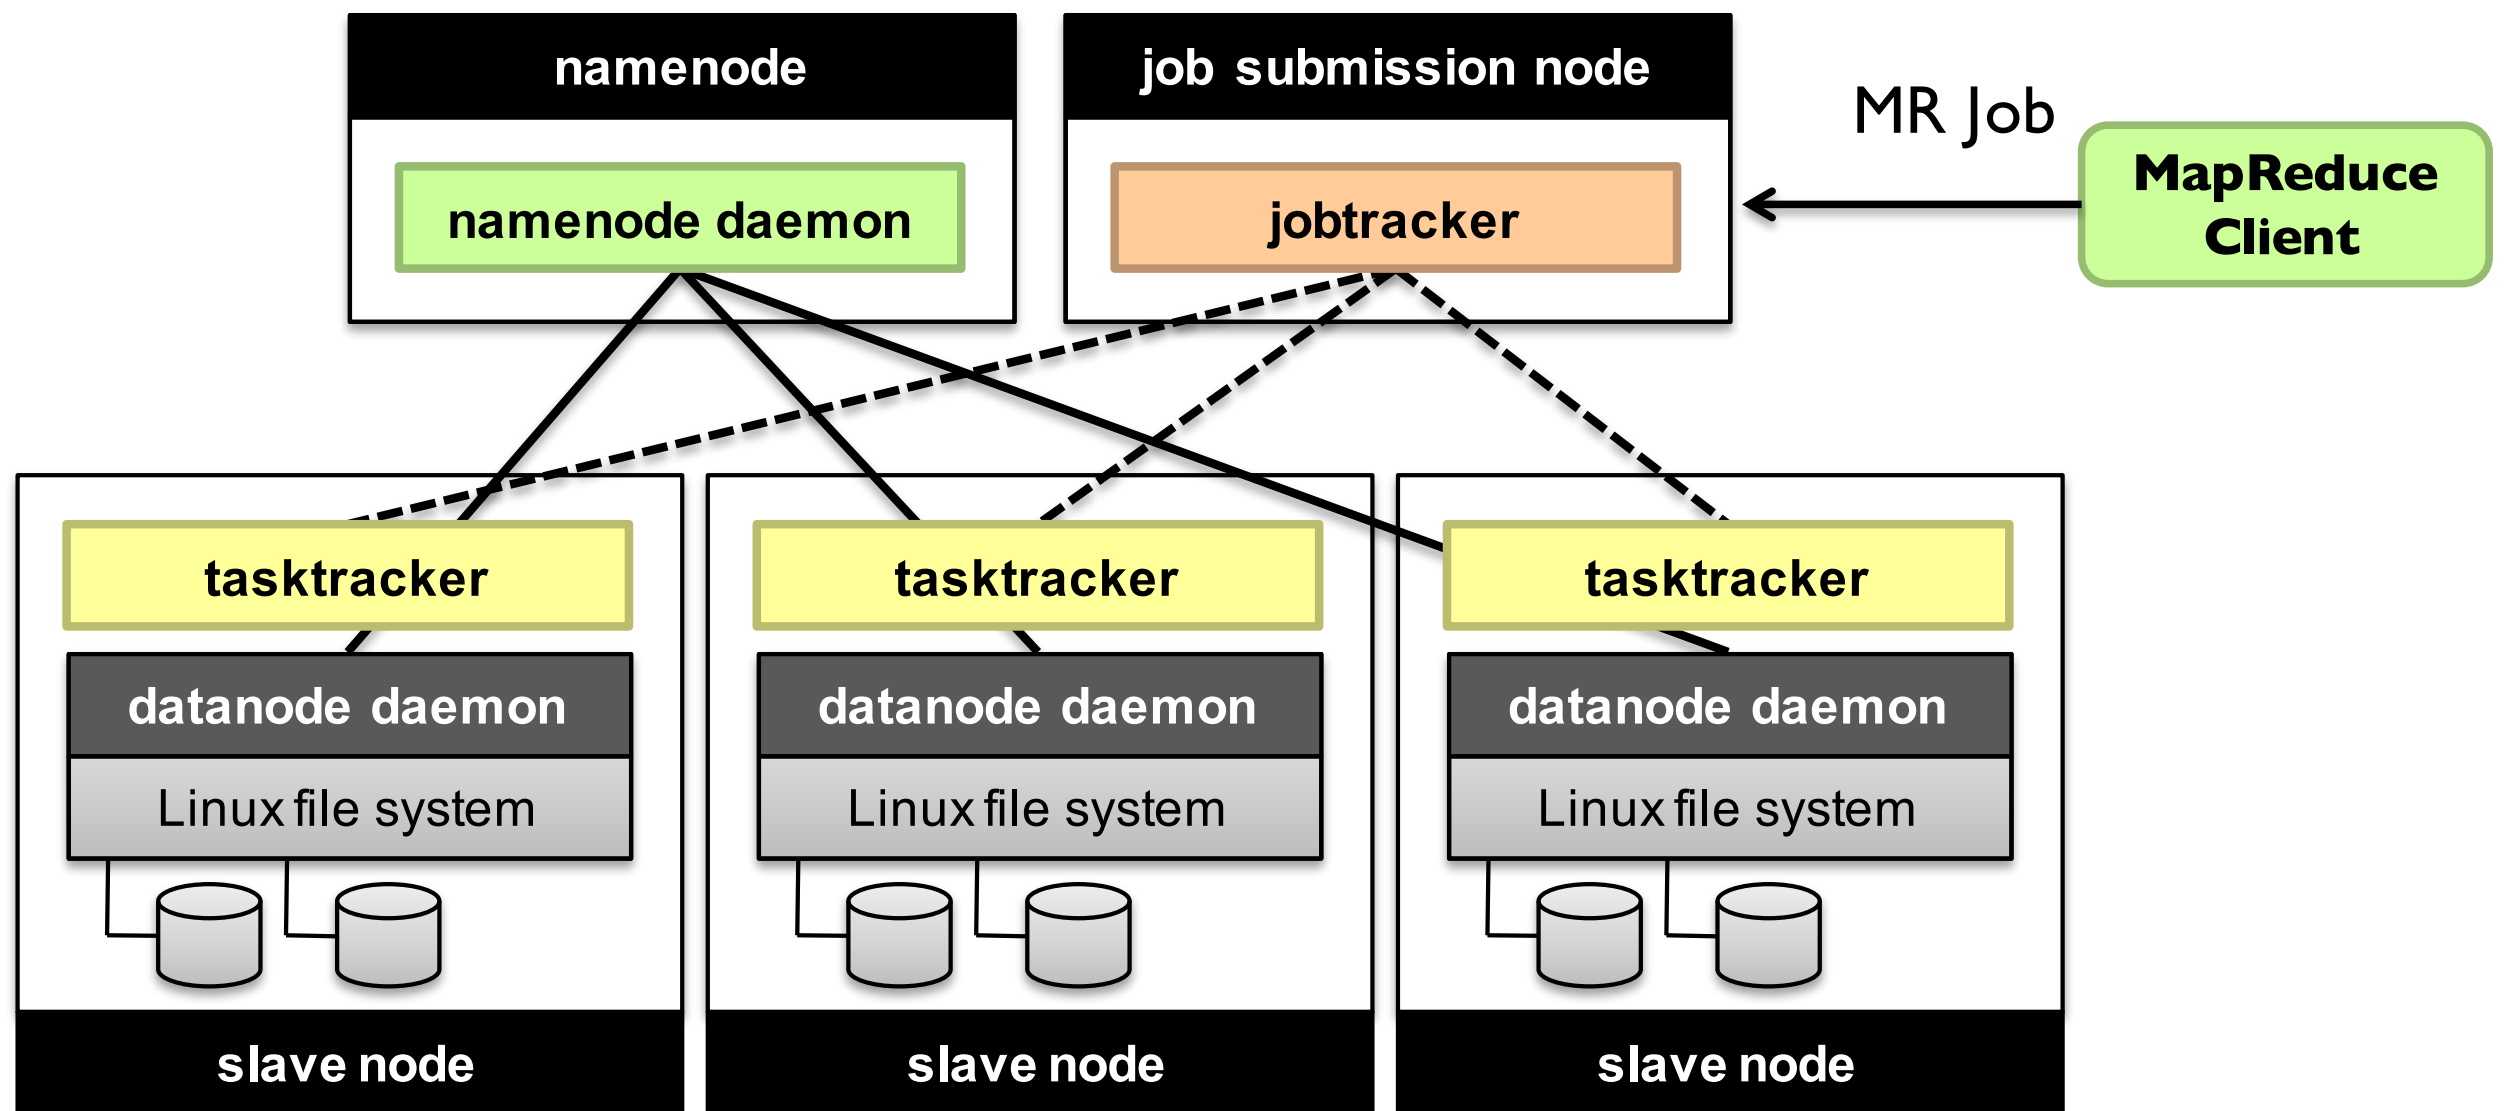
- ▶ **NameNode** (NN): master node for HDFS
- ▶ **JobTracker** (JT): master node for job submission

Set of each per slave machine:

- ▶ **DataNode** (DN): serves HDFS data blocks
- ▶ **TaskTracker** (TT): contains multiple task slots

Hadoop = HDFS + MapReduce

When HDFS meets MapReduce



Node-to-node communications

Hadoop uses its own RPC protocol

All communication begins in slave nodes

- ▶ prevents circular-wait deadlock
- ▶ slaves periodically poll for “status” message

Classes must provide explicit serialization

- ▶ that's why Hadoop data type must inherits from **Writable**

Nodes, trackers, tasks

Master node runs **JobTracker** instance, which accepts **Job** requests from clients

TaskTracker instances run on slave nodes

TaskTracker forks separate Java process for task instances

Anatomy of a job

MapReduce program in Hadoop = Hadoop job

- ▶ jobs are divided into map and reduce tasks
- ▶ multiple jobs can be composed into a **workflow**
 - ▶ map->reduce->map->reduce->...

Job distribution

Job submission:

- ▶ client (i.e., driver program) creates a job, configures it, and submits it to **JobTracker**
- ▶ “jar” file + an XML file containing serialized program configuration options

Running a MapReduce job

- ▶ places “jar” file and XML file into the HDFS
- ▶ notifies **TaskTrackers** where to retrieve the relevant program code

Under the hood

Input splits are computed (on the client end)

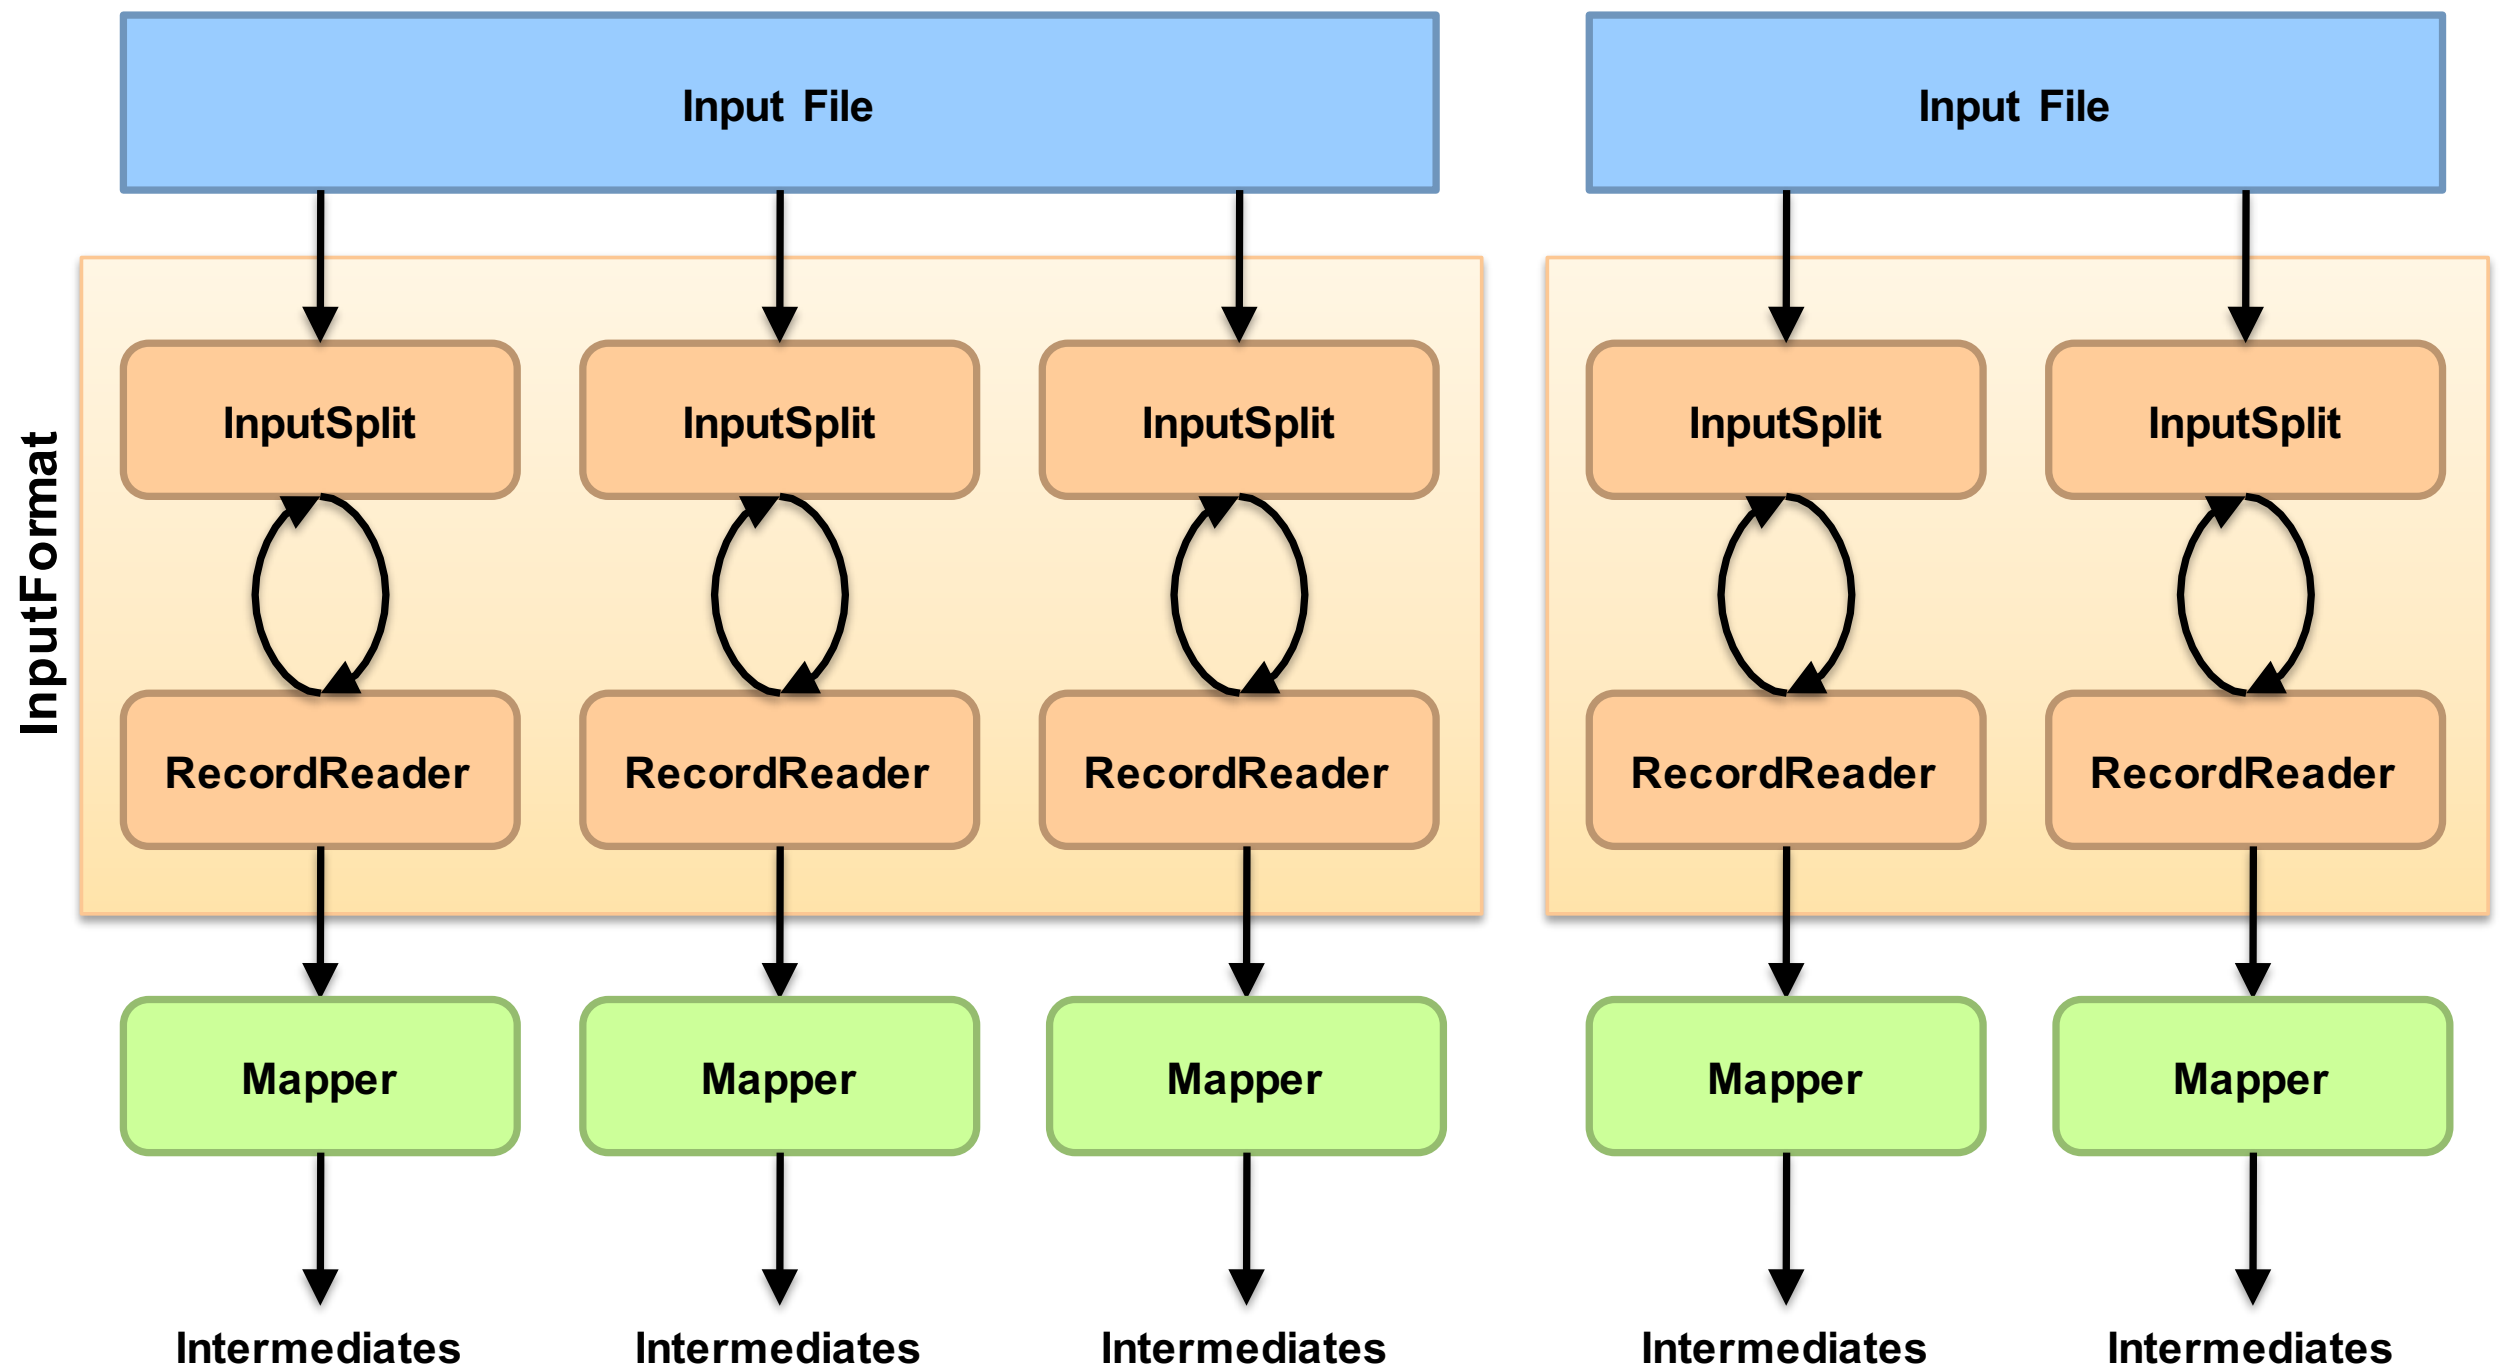
Job data sent to **JobTracker**

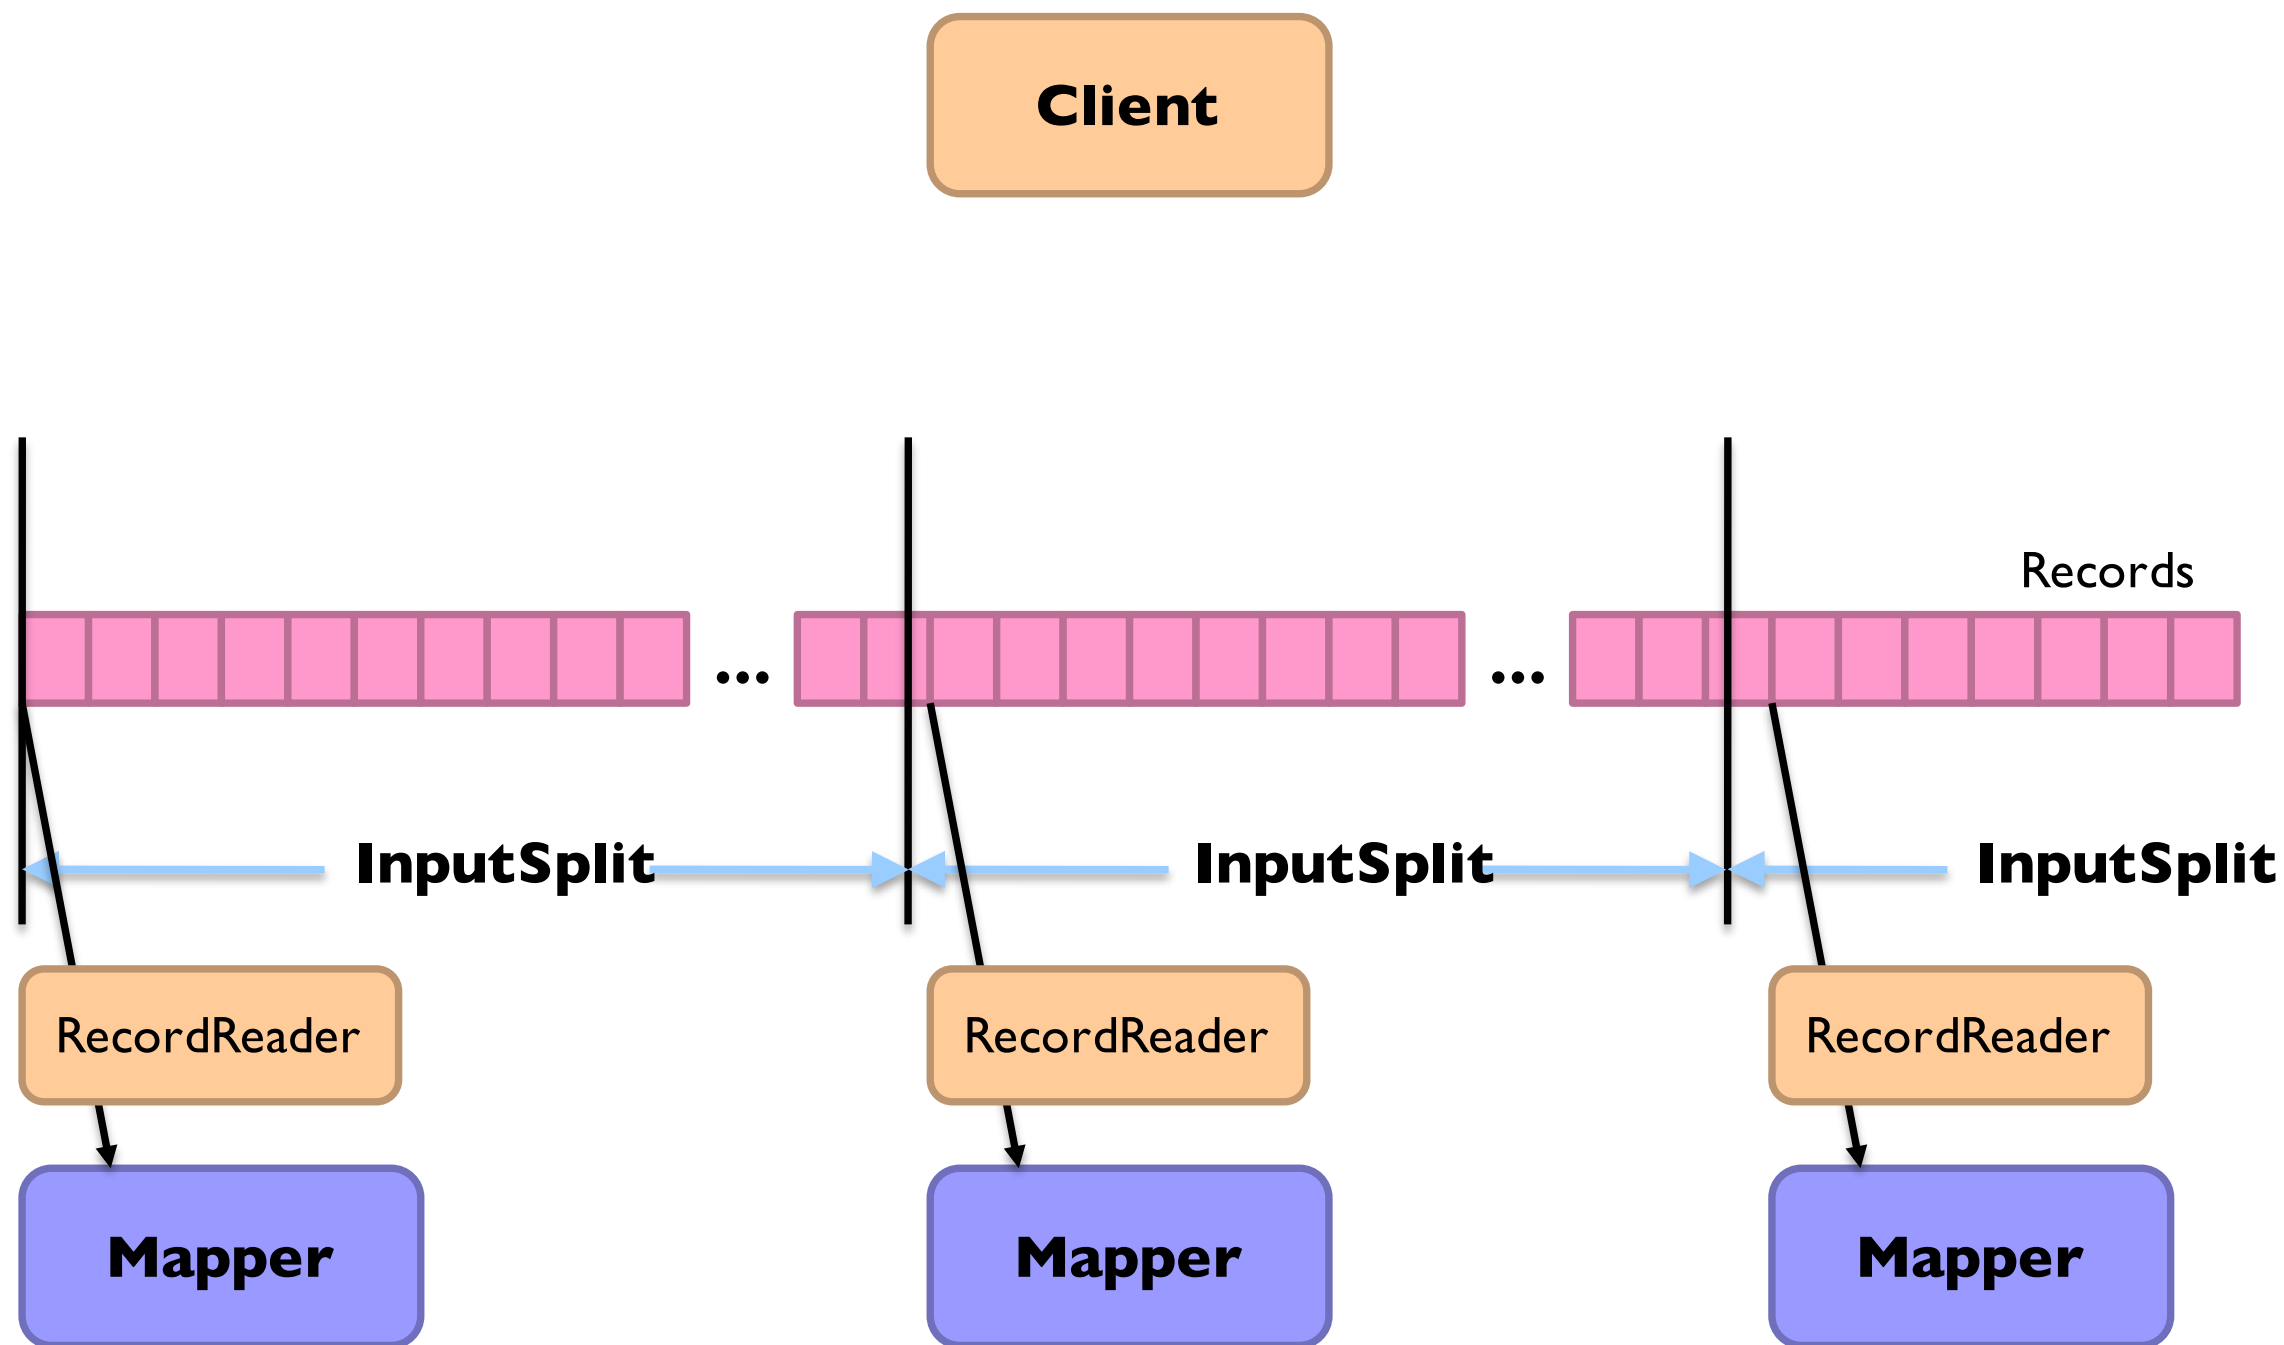
- ▶ jar + configuration XML

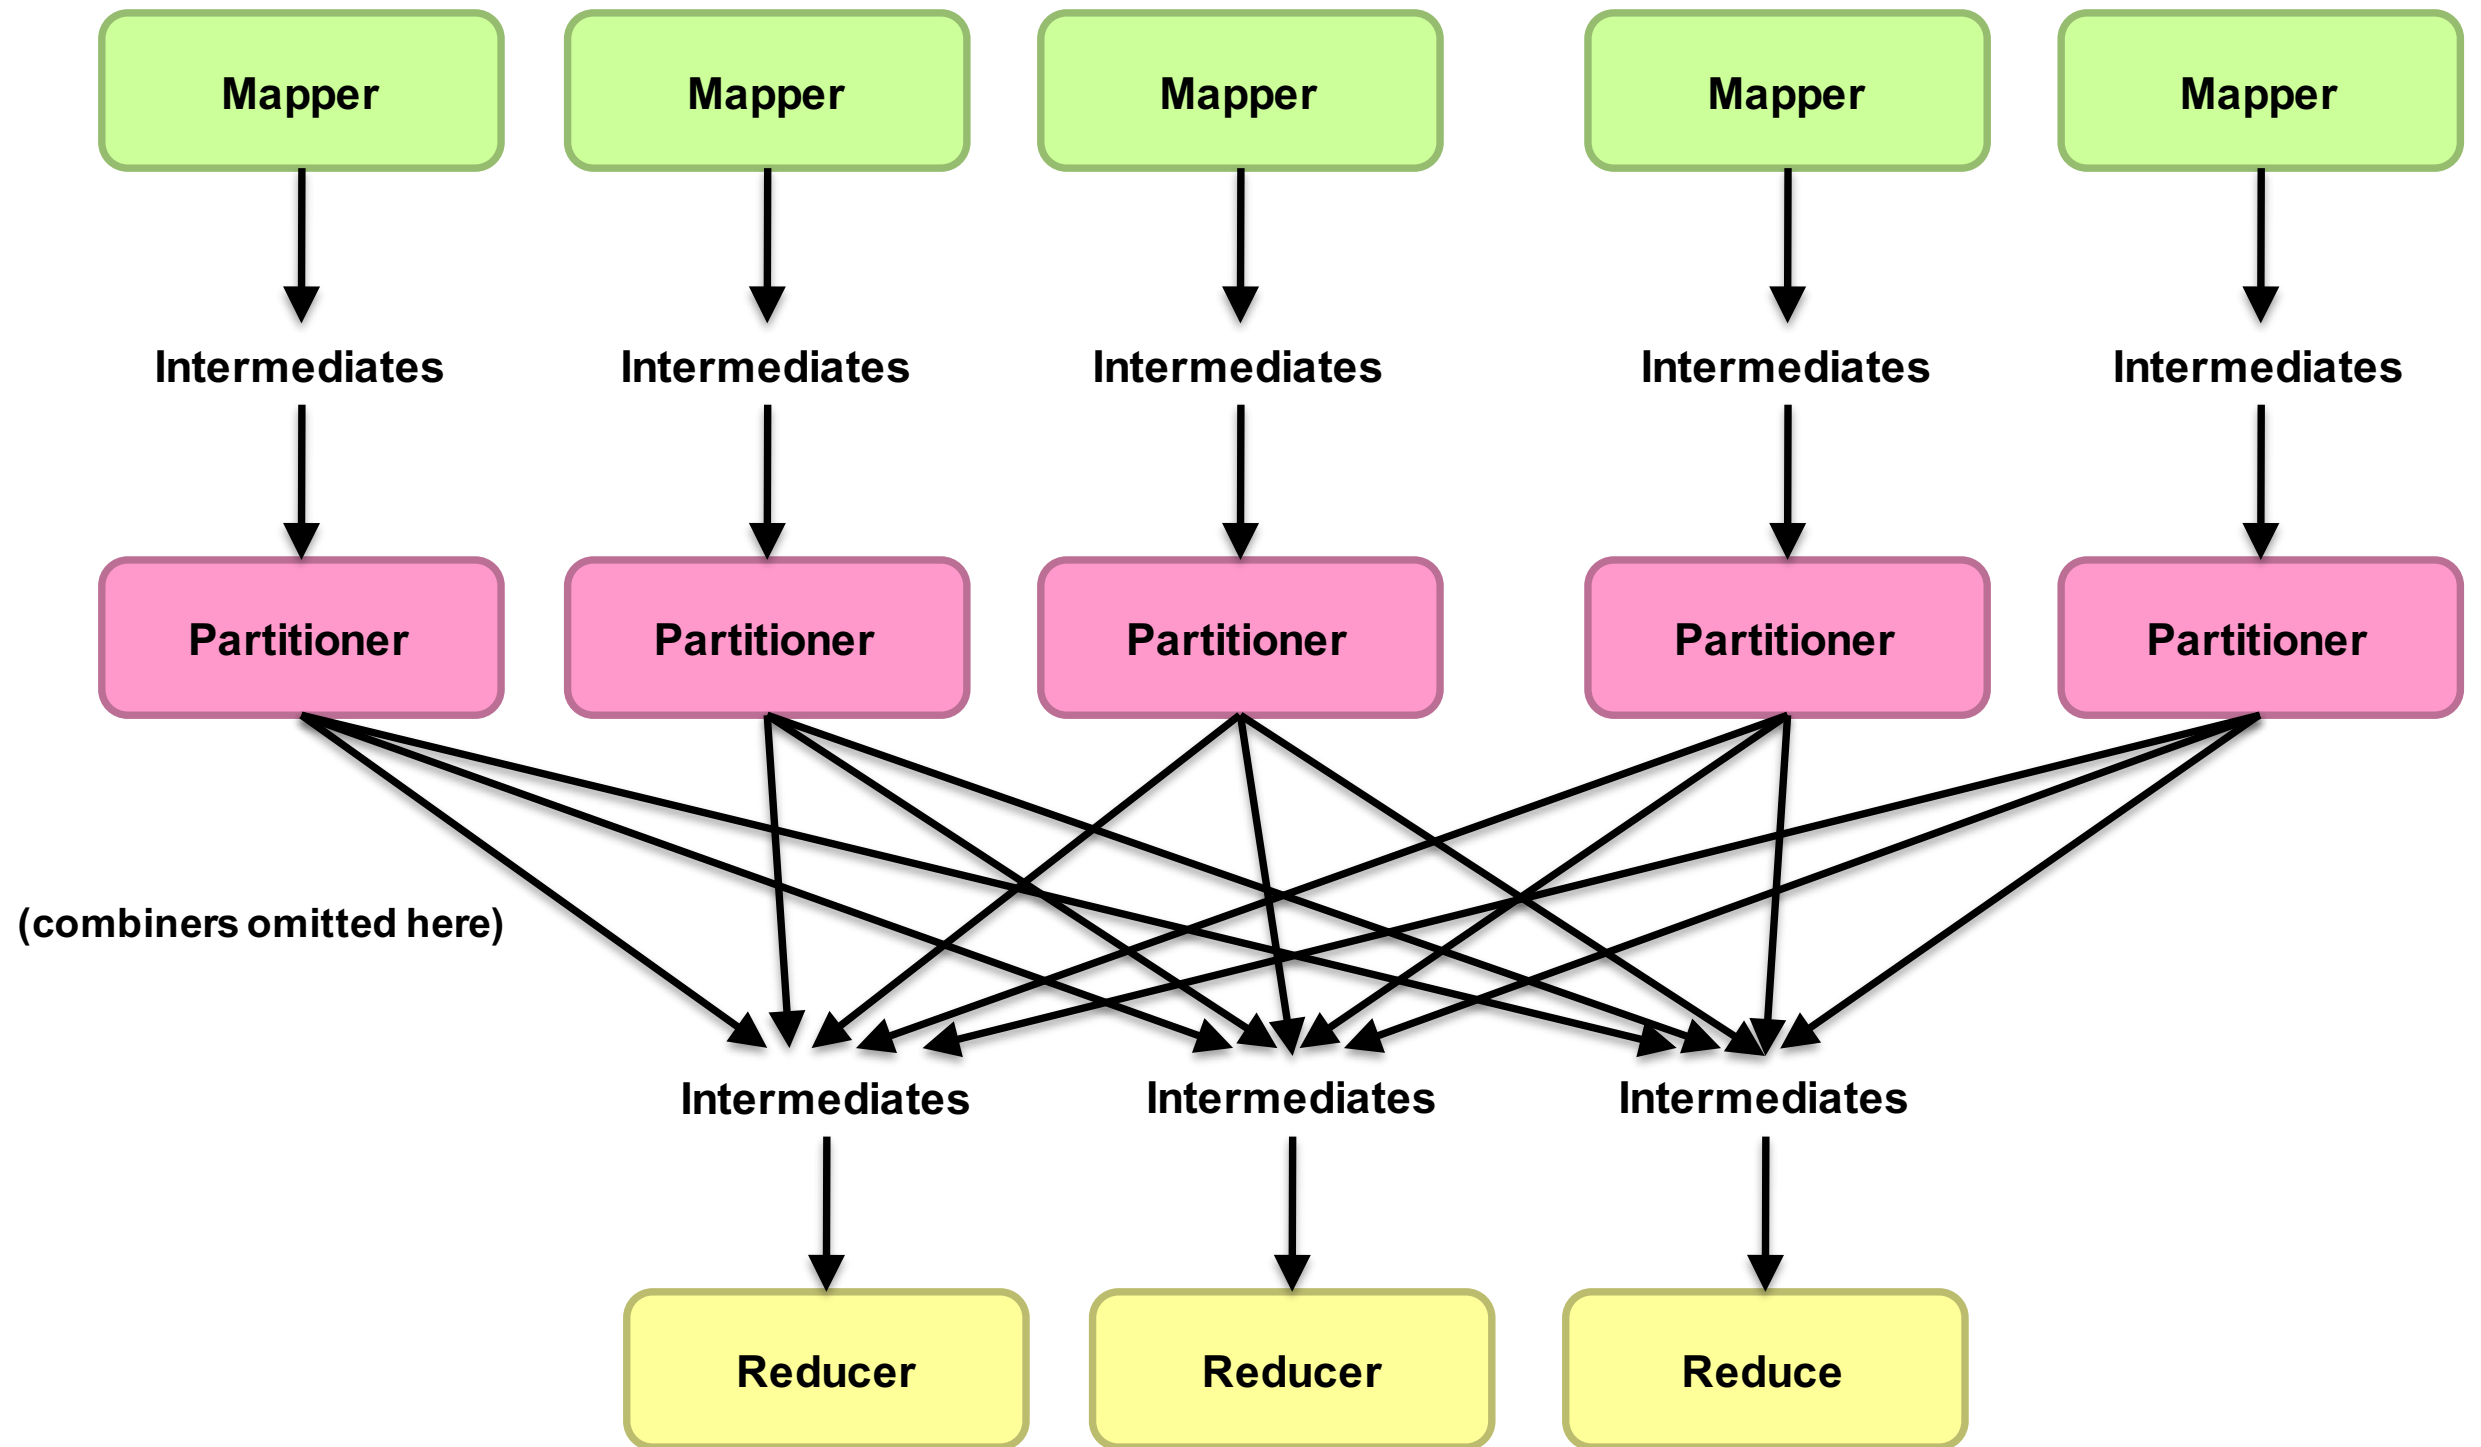
JobTracker puts job data in shared location (HDFS),
enqueues tasks

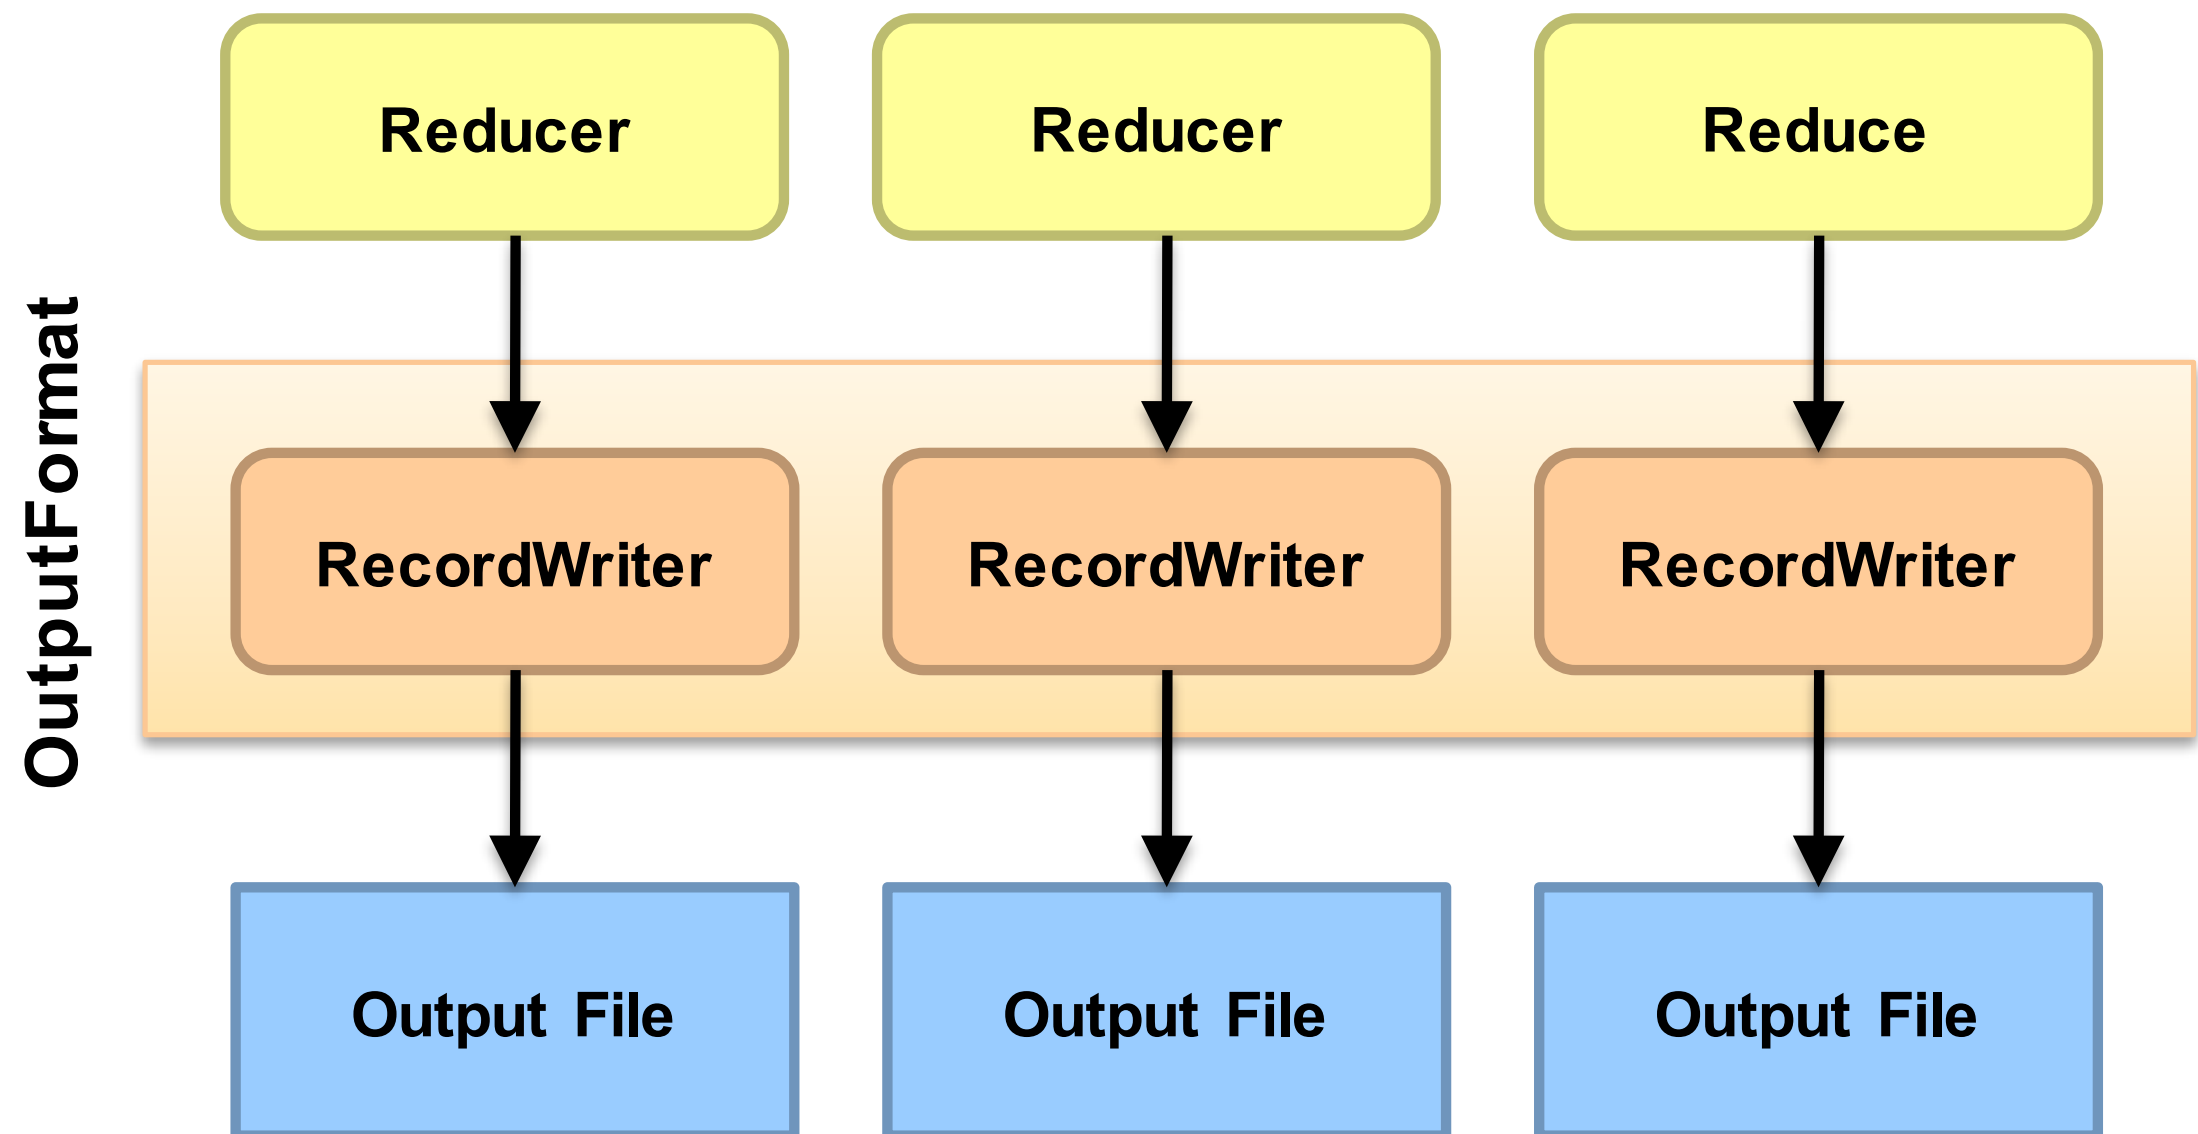
TaskTrackers poll for tasks

Off to the races...









Hadoop I/O

InputFormat:

- ▶ **TextInputFormat**: treats each ‘\n’-terminated line of a file as a value
- ▶ **KeyValueTextInputFormat**: maps ‘\n’-terminated text lines of “k v”
- ▶ **SequenceFileInputFormat**: binary file of (k, v) pairs
- ▶ ...

OutputFormat:

- ▶ **TextOutputFormat**: writes “key val\n” strings to output file
- ▶ **SequenceFileOutputFormat**: uses a binary format to pack (k, v) pairs
- ▶ ...

Shuffle and sort in Hadoop

Probably the most complex aspect of MapReduce

Map

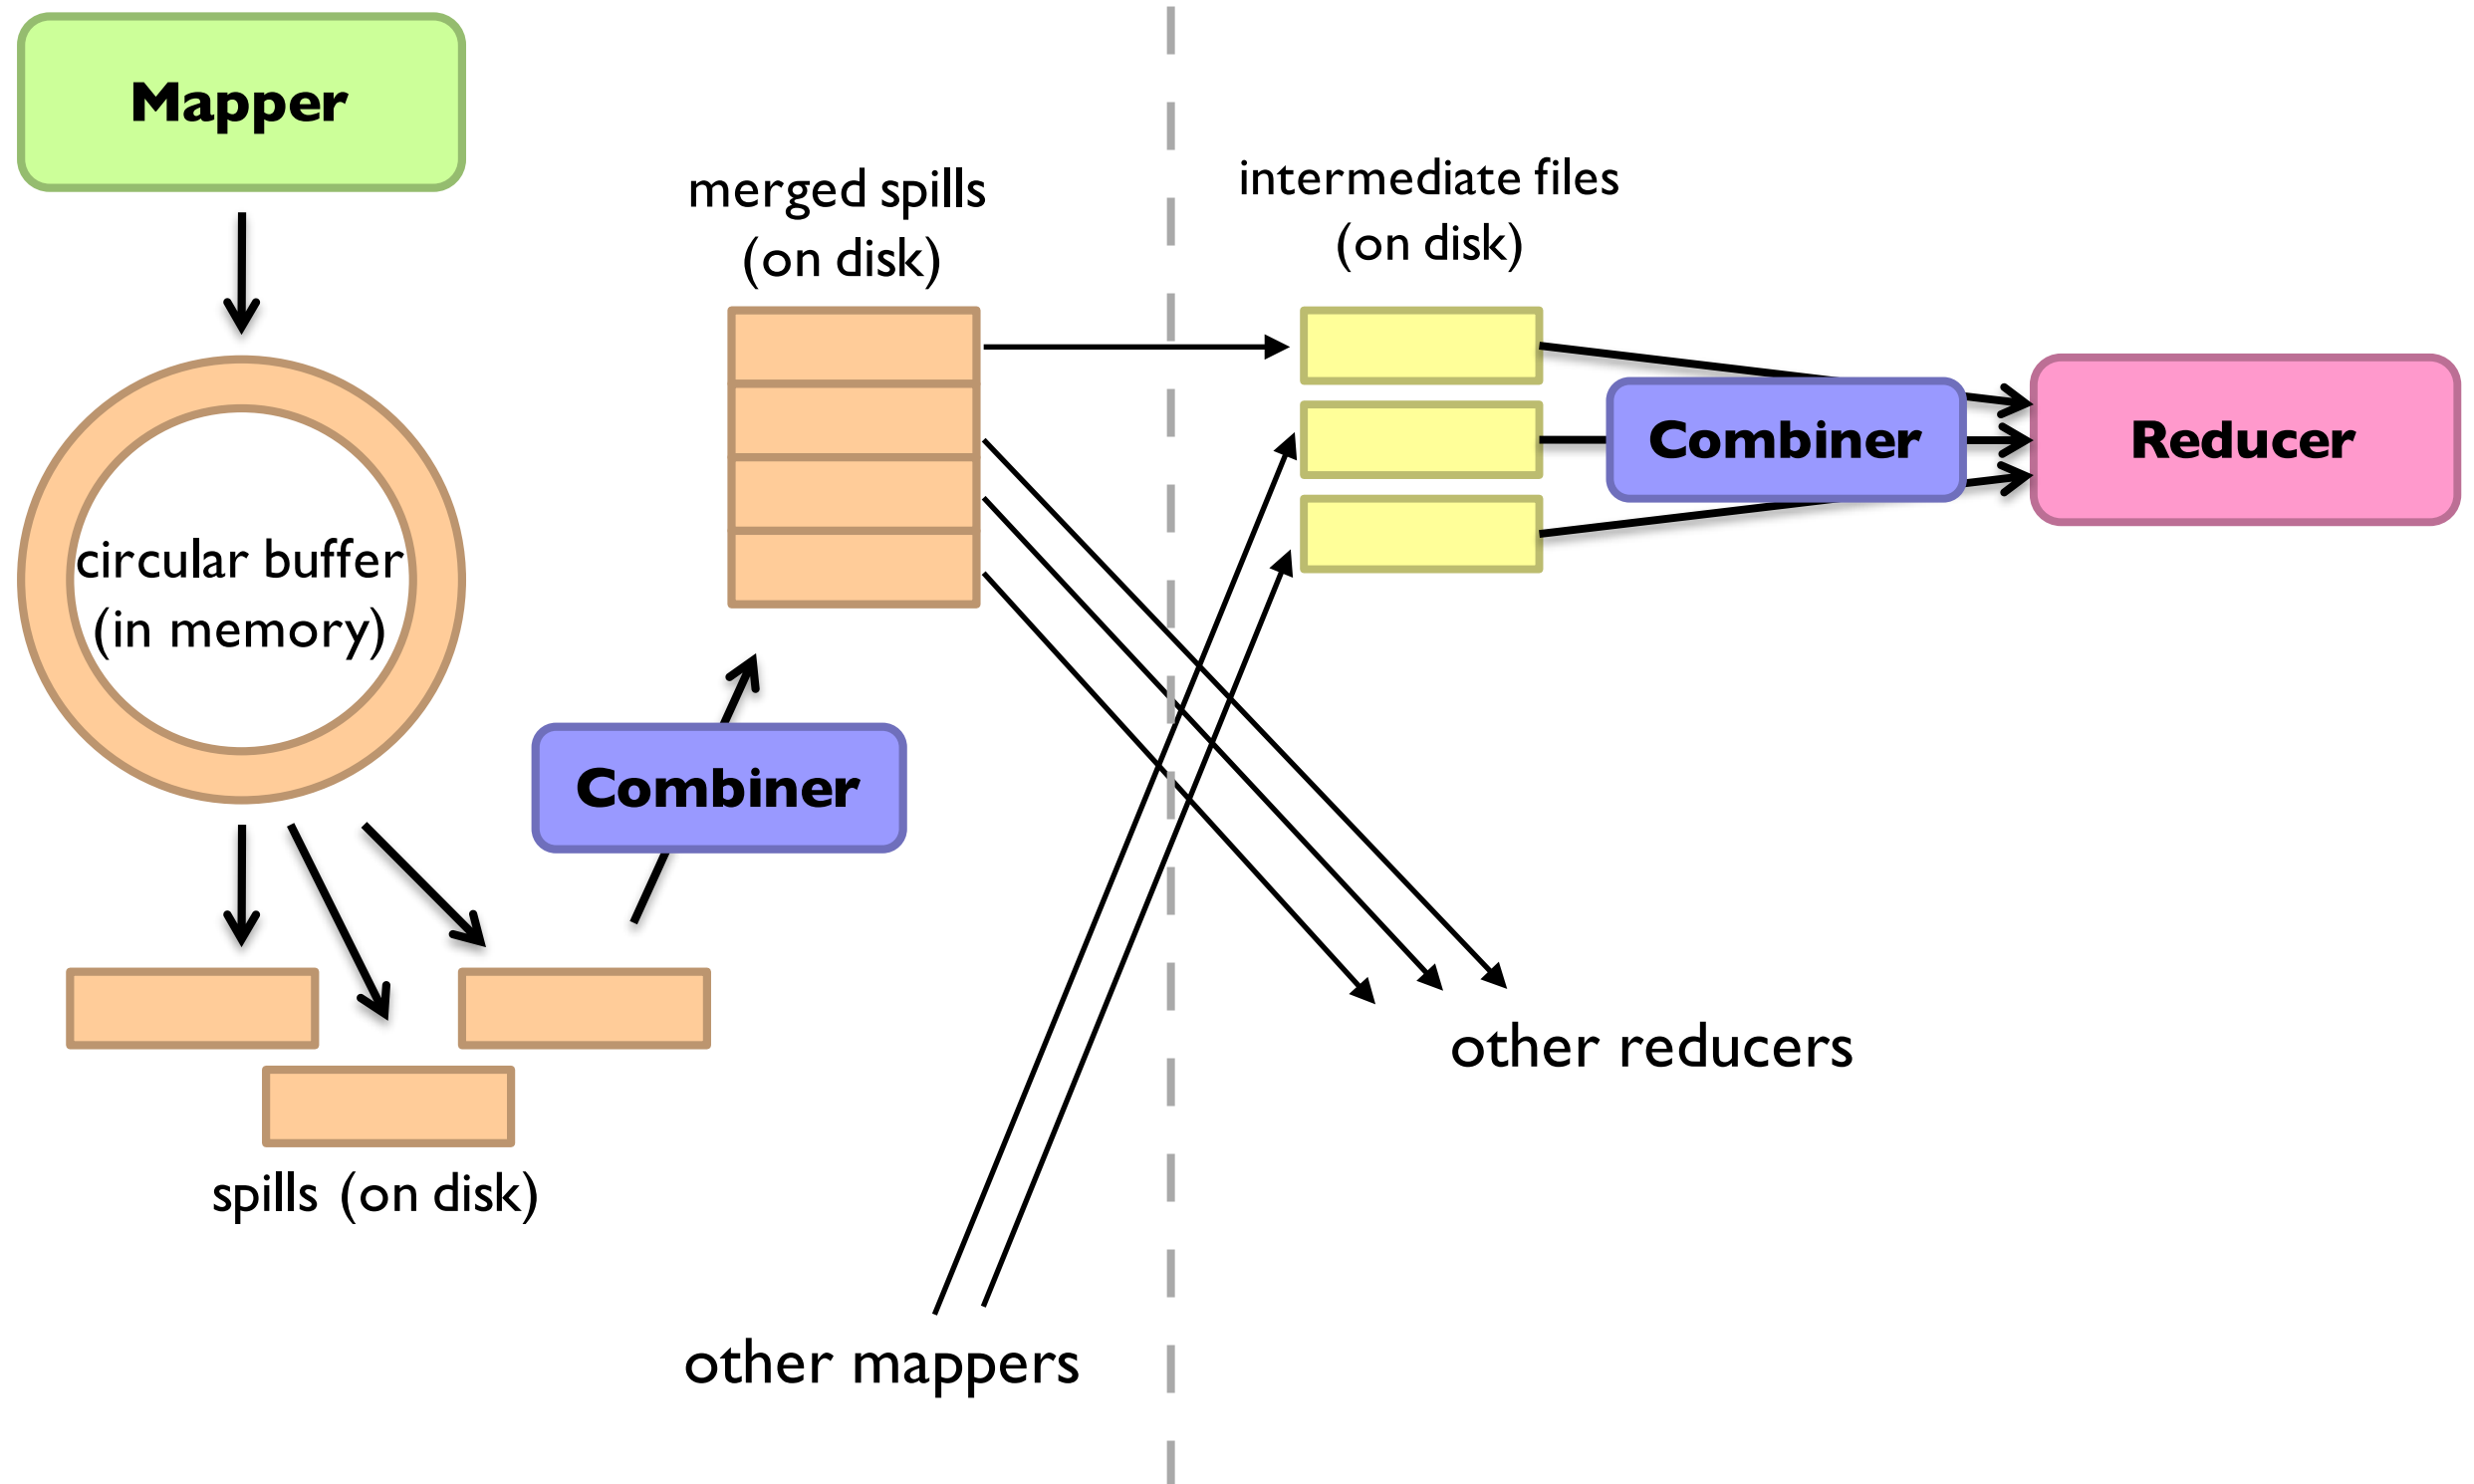
- ▶ Map outputs are buffered in memory in a circular buffer
- ▶ when buffer reaches threshold, contents are spilled to disk
- ▶ spills merged in a single, partitioned file (sorted within each partition): combiner runs during the merges

Shuffle and sort in Hadoop


Reduce

- ▶ map outputs are copied over to reducer machine
- ▶ “sort” is a multi-pass merge of map outputs (happens in memory and on disk): combiner runs during the merges
- ▶ final merge pass goes directly into reducer

Network



Hadoop workflow



Hadoop workflow

Develop code in Eclipse on local machine

Build distribution on local machine

Check out copy of code on cluster

Compile and package

Run job on cluster

Iterate

Code execution environments

Different ways to run code:

- ▶ plain Java
- ▶ local (standalone) mode
- ▶ pseudo-distributed mode (emulate cluster nodes using multiple processes)
- ▶ fully-distributed mode

Debugging

- ▶ Start small, start locally
- ▶ Build incrementally

Credits

Slides are adapted from Prof. Jimmy Lin's slides at the University of Waterloo