

# Natural Language Processing

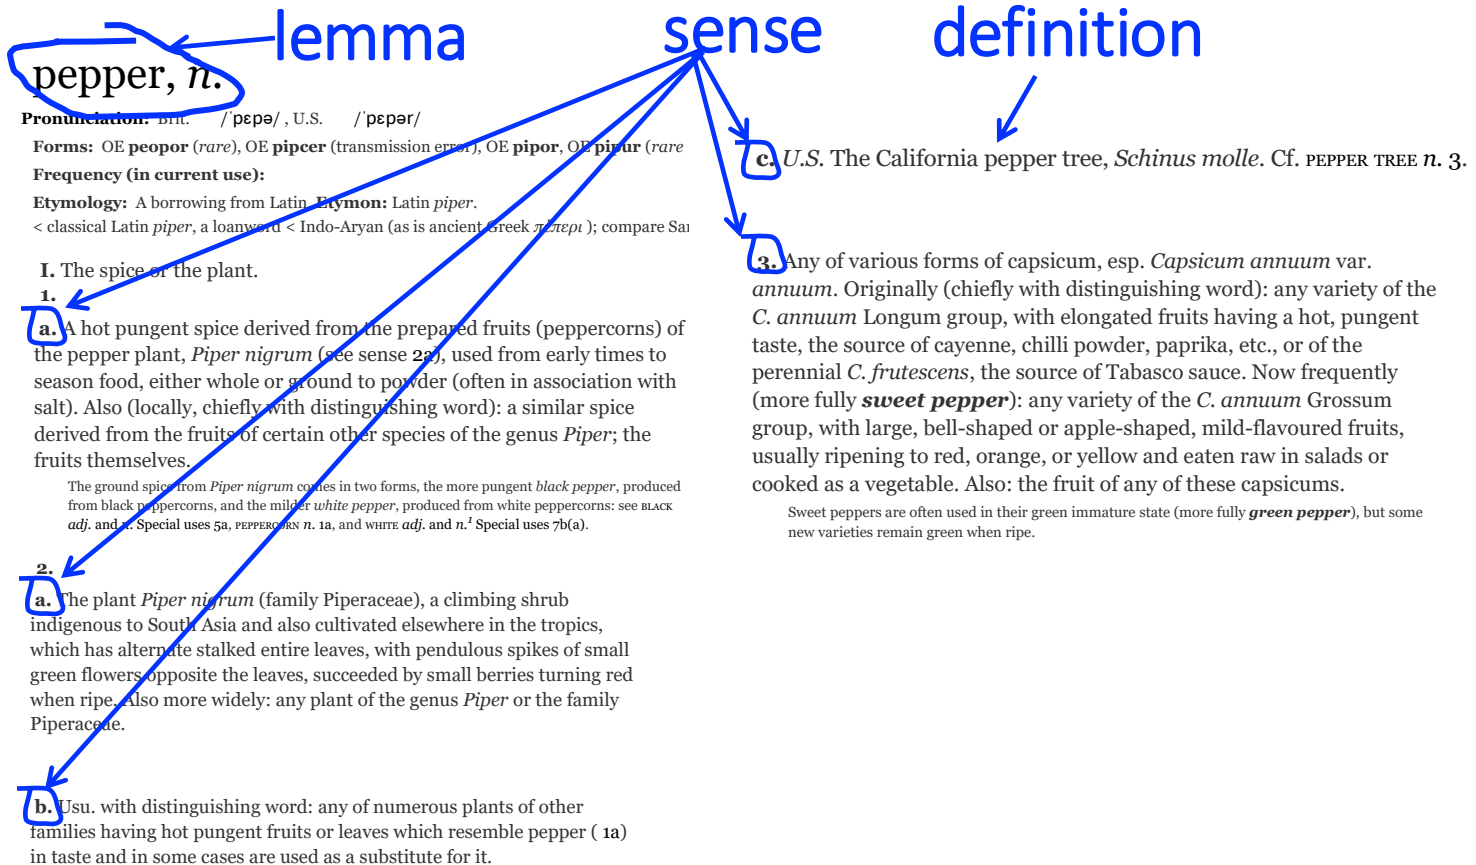
Word Embeddings

Instructor: Yangqiu Song

# What do words mean?

- First thought: look in a dictionary
- <http://www.oed.com/>

# Words, Lemmas, Senses, Definitions



# Lemma pepper

Sense 1: spice from pepper plant

Sense 2: the pepper plant itself

Sense 3: another similar plant (Jamaican pepper)

Sense 4: another plant with peppercorns (California pepper)

Sense 5: *capsicum* (i.e. chili, paprika, bell pepper, etc)

# There are relations between senses

- Synonymy
- Antonymy
- Similarity
- ...

# Relation: Similarity

Words with similar meanings. Not synonyms, but sharing some element of meaning

car, bicycle

cow, horse

# Ask humans how similar 2 words are

| word1  | word2      | similarity |
|--------|------------|------------|
| vanish | disappear  | 9.8        |
| behave | obey       | 7.3        |
| belief | impression | 5.95       |
| muscle | bone       | 3.65       |
| modest | flexible   | 0.98       |
| hole   | agreement  | 0.3        |

SimLex-999 dataset (Hill et al., 2015)

# Relation: Word relatedness

- Also called "word association"
- Words be related in any way, perhaps via a semantic frame or field
  - car, bicycle: **similar**
  - car, gasoline: **related**, not similar
- Similarity is a specific type of relatedness: graded
  - car vs. automobile -> 1.0
  - car vs. vehicle -> 0.6
  - car vs. tire -> 0.2
  - car vs. street -> 0.1
- Similarity: **synonyms**, **hyponyms/hyperonyms**, and **siblings** are highly similar
  - doctor vs. surgeon, bike vs. bicycle
- Relatedness: **topically related** or based on any other **semantic relation**
  - heart vs. surgeon, tire vs. car



- What is the computational approach to evaluate word similarities/relatedness?

# Computational Approaches

- Knowledge base based
  - WordNet Similarity
  - ...
- Corpus based
  - Distributional similarity
  - Deep learning

# Corpus based Approach

- Distributional semantics
  - The basic idea of **distributional semantics** can be summed up in the so-called **distributional hypothesis**: *linguistic items with similar distributions have similar meanings*.
    - You shall know a word by the company it keeps." (Firth (1957))
  - The **distributional hypothesis** in linguistics is derived from the semantic theory of language usage, i.e. words that are used and occur in the same contexts tend to purport similar meanings.
    - Distributional hypothesis: Semantically similar words occur in similar contexts (Harris (1954))

We will mention **distributed representation** based neural language models in this class

# Corpus based Approach


## 1) Corpus



## 2) Preprocessing



## 3) Dimensionality Reduction



## 4) Post Processing



# Let's try to keep the kitchen \_\_\_\_\_ .

- Observation: context can tell us a lot about word meaning
- **Context**: local window around a word occurrence (for now)
- Pros: data-driven, easy to implement
- Cons: ambiguity

# Window based co-occurrence matrix

- Window length 1 (more common: 5 - 10)
- Symmetric (irrelevant whether left or right context)
- Example corpus:
  - I like deep learning.
  - I like NLP.
  - I enjoy flying.

# Window based co-occurrence matrix

- Example corpus:
  - I like deep learning.
  - I like NLP.
  - I enjoy flying.

| counts   | I | like | enjoy | deep | learning | NLP | flying | . |
|----------|---|------|-------|------|----------|-----|--------|---|
| I        | 0 | 2    | 1     | 0    | 0        | 0   | 0      | 0 |
| like     | 2 | 0    | 0     | 1    | 0        | 1   | 0      | 0 |
| enjoy    | 1 | 0    | 0     | 0    | 0        | 0   | 1      | 0 |
| deep     | 0 | 1    | 0     | 0    | 1        | 0   | 0      | 0 |
| learning | 0 | 0    | 0     | 1    | 0        | 0   | 0      | 1 |
| NLP      | 0 | 1    | 0     | 0    | 0        | 0   | 0      | 1 |
| flying   | 0 | 0    | 1     | 0    | 0        | 0   | 0      | 1 |
| .        | 0 | 0    | 0     | 0    | 1        | 1   | 1      | 0 |

# Problems with simple co-occurrence vectors

- Increase in size with vocabulary
- Very high dimensional: require a lot of storage
- Subsequent classification models have sparsity issues
- → Models are less robust



# Solution: Low dimensional vectors

- Idea: store “most” of the important information in a fixed, small number of dimensions: a dense vector
- Usually around 25 – 1000 dimensions
- How to reduce the dimensionality?

# Method 1: Dimensionality Reduction on $X$

- Singular Value Decomposition of co-occurrence matrix  $X$ .

$$\begin{array}{ccccc}
 \begin{array}{c} m \\ \boxed{\phantom{X}} \\ n \\ X \end{array} & = & \begin{array}{c} r \\ \boxed{\begin{array}{c} | \quad | \quad | \quad \cdots \\ U_1 U_2 U_3 \cdots \\ | \quad | \quad | \quad \cdots \end{array}} \\ n \\ U \end{array} & \begin{array}{c} r \\ \boxed{\begin{array}{c} S_1 \quad \quad \quad 0 \\ \quad S_2 \quad S_3 \quad \cdots \\ 0 \quad \quad \quad \cdots \quad S_r \end{array}} \\ r \\ S \end{array} & \begin{array}{c} m \\ \boxed{\begin{array}{c} \text{---} V_1 \text{---} \\ \text{---} V_2 \text{---} \\ \text{---} V_3 \text{---} \\ \vdots \end{array}} \\ r \\ V^T \end{array} \\
 \\
 \begin{array}{c} m \\ \boxed{\phantom{\hat{X}}} \\ n \\ \hat{X} \end{array} & = & \begin{array}{c} k \\ \boxed{\begin{array}{c} | \quad | \quad | \quad \cdots \\ U_1 U_2 U_3 \cdots \\ | \quad | \quad | \quad \cdots \end{array}} \\ n \\ \hat{U} \end{array} & \begin{array}{c} k \\ \boxed{\begin{array}{c} S_1 \quad \quad \quad 0 \\ \quad S_2 \quad S_3 \quad \cdots \\ 0 \quad \quad \quad \cdots \quad S_k \end{array}} \\ k \\ \hat{S} \end{array} & \begin{array}{c} m \\ \boxed{\begin{array}{c} \text{---} V_1 \text{---} \\ \text{---} V_2 \text{---} \\ \text{---} V_3 \text{---} \\ \vdots \end{array}} \\ k \\ \hat{V}^T \end{array}
 \end{array}$$

best rank  $k$  approximation to  $X$ , in terms of least squares.

# Simple SVD word vectors in Python

- Corpus:
- I like deep learning. I like NLP. I enjoy flying.


```
import numpy as np
la = np.linalg
words = ["I", "like", "enjoy",
         "deep", "learnig", "NLP", "flying", "."]
X = np.array([[0, 2, 1, 0, 0, 0, 0, 0],
              [2, 0, 0, 1, 0, 1, 0, 0],
              [1, 0, 0, 0, 0, 0, 1, 0],
              [0, 1, 0, 0, 1, 0, 0, 0],
              [0, 0, 0, 1, 0, 0, 0, 1],
              [0, 1, 0, 0, 0, 0, 0, 1],
              [0, 0, 1, 0, 0, 0, 0, 1],
              [0, 0, 0, 0, 1, 1, 1, 0]])

U, s, Vh = la.svd(X, full_matrices=False)
```

# Simple SVD word vectors in Python

- Corpus: I like deep learning. I like NLP. I enjoy flying.
- Printing first two columns of U corresponding to the 2 biggest singular values

```
for i in xrange(len(words)):  
    plt.text(U[i,0], U[i,1], words[i])
```



# Word meaning is defined in terms of vectors


- In most deep learning models, a word is represented as a dense vector

$$\textit{linguistics} = \begin{pmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 0.109 \\ -0.542 \\ 0.349 \\ 0.271 \end{pmatrix}$$

# Hacks to X


- Problem: function words (the, he, has) are too frequent  
→ syntax has too much impact. Some fixes:
  - $\min(X, t)$ , with  $t \sim 100$
  - Ignore them all
- Use Pearson correlations instead of counts, then set negative values to 0

# Interesting semantic patterns emerge in the vectors



- An Improved Model of Semantic Similarity Based on Lexical Co-Occurrence (Rohde et al. 2005)


# Interesting semantic patterns emerge in the vectors



- An Improved Model of Semantic Similarity Based on Lexical Co-Occurrence (Rohde et al. 2005)



# Interesting semantic patterns emerge in the vectors



- An Improved Model of Semantic Similarity Based on Lexical Co-Occurrence (Rohde et al. 2005)

# Problems with SVD

- Computational cost scales quadratically for  $n \times m$  matrix:
  - $O(mn^2)$  flops (when  $n < m$ )
  - → Bad for millions of words or documents
- Hard to incorporate new words or documents
- Different learning regime than other DL models

# Idea: Directly learn low-dimensional word vectors

- Old idea. Relevant for this lecture & deep learning:
  - Learning representations by back-propagating errors.  
(Rumelhart et al., 1986)
  - A neural probabilistic language model (Bengio et al., 2003)
    - Multilayer perceptron
  - NLP (almost) from Scratch (Collobert & Weston, 2008)
    - CNN
  - An even simpler and faster model:
    - word2vec (Mikolov et al. 2013) → intro now

# Distributed Representations


- This is usually called **distributed representations** in the context of deep learning
  - Vector representation does not represent a distribution, but distributed over the space
  - Term widely used in connectionism (Learning distributed representations of concepts, Hinton (1986))
    - “In the componential approach each concept is simply a set of features and so a neural net can be made to implement a set of concepts by assigning a unit to each feature and setting the strengths of the connections between units so that each concept corresponds to a stable pattern of activity distributed over the whole network.”
- Compared to distributional semantics
  - The **distributional hypothesis** in linguistics is derived from the semantic theory of language usage, i.e. words that are used and occur in the same contexts tend to purport similar meanings.

# Main Idea of word2vec


- Instead of capturing co-occurrence counts directly,
- Predict surrounding words of every word
- Faster and can easily incorporate a new sentence/document or add a word to the vocabulary

# Represent the meaning of word – word2vec

- 2 basic neural network models:
  - Continuous Bag of Word (CBOW): use a window of word to predict the middle word
  - Skip-gram (SG): use a word to predict the surrounding ones in window.




**CBOW**




**Skip-gram**


# Word2vec – Continuous Bag of Word


- E.g. “The cat sat on floor”
  - Window size = 2























We can consider either  $W$  or  $W'$  as the word's representation. Or even take the average.

# Approximations

- With large vocabularies this objective function is not scalable and would train too slowly! → Why?
- Idea: approximate the normalization or
- Define negative prediction that only samples a few words that do not appear in the context
- Similar to focusing on mostly positive correlations
- More reading
  - [https://canvas.ust.hk/courses/16504/files/1444107?module\\_item\\_id=214783](https://canvas.ust.hk/courses/16504/files/1444107?module_item_id=214783)

# Linear Relationships in word2vec

- These representations are *very good* at encoding dimensions of similarity!
- Analogies testing dimensions of similarity can be solved quite well just by doing vector subtraction in the embedding space
  - Syntactically
    - *apple – apples  $\approx$  car – cars  $\approx$  family – families*
    - Similarly for verb and adjective morphological forms
  - Semantically (Semeval 2012 task 2)
    - *shirt – clothing  $\approx$  chair – furniture*
    - *king – man  $\approx$  queen – woman*



# Word Analogies

- Test for linear relationships, examined by Mikolov et al. (2014)

a:b :: c:?



$$d = \arg \max_x \frac{(w_b - w_a + w_c)^T w_x}{||w_b - w_a + w_c||}$$

man:woman :: king:?


+ king [ 0.30 0.70 ]

- man [ 0.20 0.20 ]


+ woman [ 0.60 0.30 ]

---


queen [ 0.70 0.80 ]




# Word analogies




# Glove Visualizations



# Glove Visualizations: Company - CEO



# Glove Visualizations: Superlatives



# More Examples

- “word2vec Parameter Learning Explained”, Xin Rong
  - <https://ronxin.github.io/wevi/>
- Word2Vec Tutorial - The Skip-Gram Model
  - <http://mccormickml.com/2016/04/19/word2vec-tutorial-the-skip-gram-model/>