

# Software Security: Exploitation

Shuai Wang



香港科技大學

THE HONG KONG UNIVERSITY OF SCIENCE AND TECHNOLOGY

# Software bugs that can be potentially exploited

- The so-called “**vulnerability**”
  - Crash Causing Defects
  - Null pointer dereference
  - Use after free
  - Double free
  - Array indexing errors
  - Mismatched array new/delete
  - Stack overflow
  - Heap overflow
  - Return pointers to local variables
  - Logically inconsistent code
  - Uninitialized variables
  - Invalid use of negative values
  - Passing large parameters by value
  - Underallocations of dynamic data
  - Memory leaks
  - File handle leaks
  - Network resource leaks
  - Unused values
  - Unhandled return codes
  - Use of invalid iterators

But note that many of the (logic) bugs are not exploitable and therefore are not “**vulnerabilities**”

# Common Software Vulnerabilities

- Buffer overflows ← this time
- Format string problems
- Integer overflows
- Use after free ← next time
- Type confusion
- Race conditions...

Attacks marked in red are the most popular ones.

# Buffer Overflow

- Used in 1988's Morris Internet Worm
- Alphe One's "*Smashing The Stack For Fun And Profit*" in 1996 popularizes stack buffer overflows
  - Your reading article today..
- Still extremely common today
  - Over 50% of advisories published by CERT (computer security incident report team) are caused by various buffer overflows



# Computer Memory

Your own memory may look like this:

wake up; have breakfast; need to  
buy milk; turn off the lights; go to  
class; that man has a strange shirt;  
fall asleep; wake up

A web server's memory may look like this:

Bob requests main page; Atta wants  
reply "Cat"; Li sets password to  
"sup3rsekr1t"; Kate wants image  
"derpy\_cat"; Poe sets secret key; ...

# Memory Read



Please reply "Cat"  
(3 letters).

Please reply "Cat"  
(5 letters).

Bob requests main page; Atta wants  
reply "Cat"; Li sets password to  
"sup3rsekr1t"; Kate wants image  
"derpy\_cat"; Poe sets secret key; ...

Memory


Cat



Cat";

# Under the hood

```
char buffer[100];
```



Buffer is a data storage area inside computer memory

- Intended to hold pre-defined amount of data
- If more data is stored/required from it, it tampers the adjacent memory

# Memory/Buffer Overread

Bob requests main page; Atta wants reply "Cat"; Li sets password to "sup3rsekr1t"; Kate wants image "derpy\_cat"; Poe sets secret key; ...

Memory

Please reply "Cat"  
(100 letters).

Cat"; Li sets password to  
"sup3rsekr1t"; Kate wants  
image "derpy\_cat"; Poe  
sets secret key; ...





# Exploiting Memory For Fun and Profit

- All right, we talked about how buffer can be **overread**
  - The “Heartbleed” attack on OpenSSL
- Then, what about buffer **overwritten**?
  - Change the memory content?
  - What else?
    - Note that “**code**” can also be written in the memory... → executing arbitrary code!




# Buffer Overwrite Scenario

- Users enter data into a Web form
- Web form is sent to server
- Server writes data to array called buffer, **without checking length of input data**
- Data “overflows” buffer
  - Such overflow might enable an attack
  - If so, attack could be carried out by anyone with Internet access

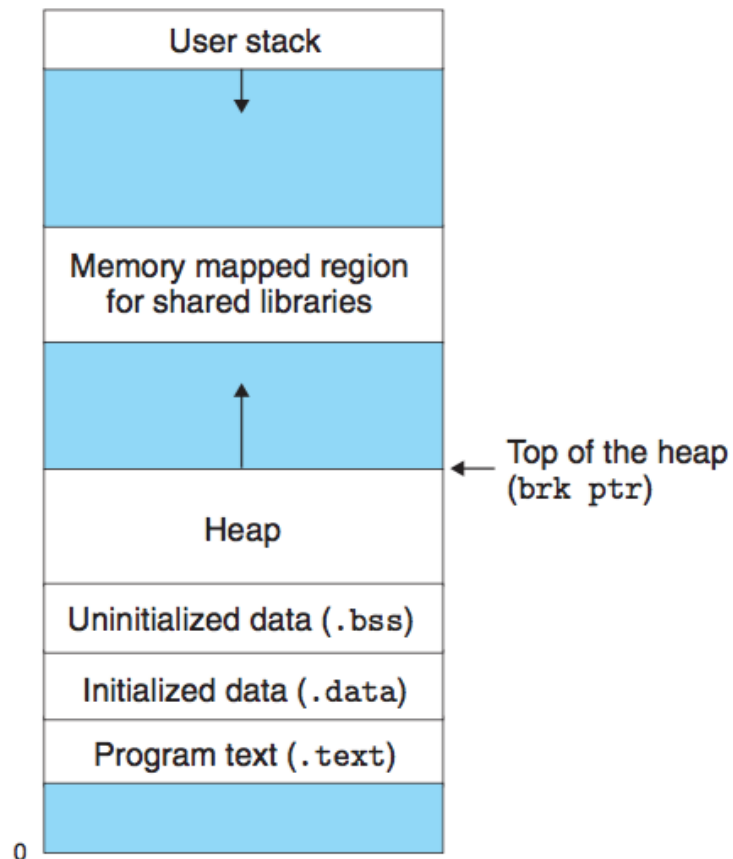
# Simple Buffer Overflow

- Consider boolean flag to check password correct or not
- Buffer overflow could overwrite flag allowing anyone to authenticate




# How is memory implemented?

- Linux standard memory layout for a process but Mac/Windows are very similar.




# A Stack Frame



- A stack consists of logical stack frames that are pushed when calling a function and popped when returning.
- When a function is called, the return address, stack frame pointer and the variables are pushed on the stack.
- The stack grows from high address to low address.
  - When we overflow the buffer, the return address will be overwritten.

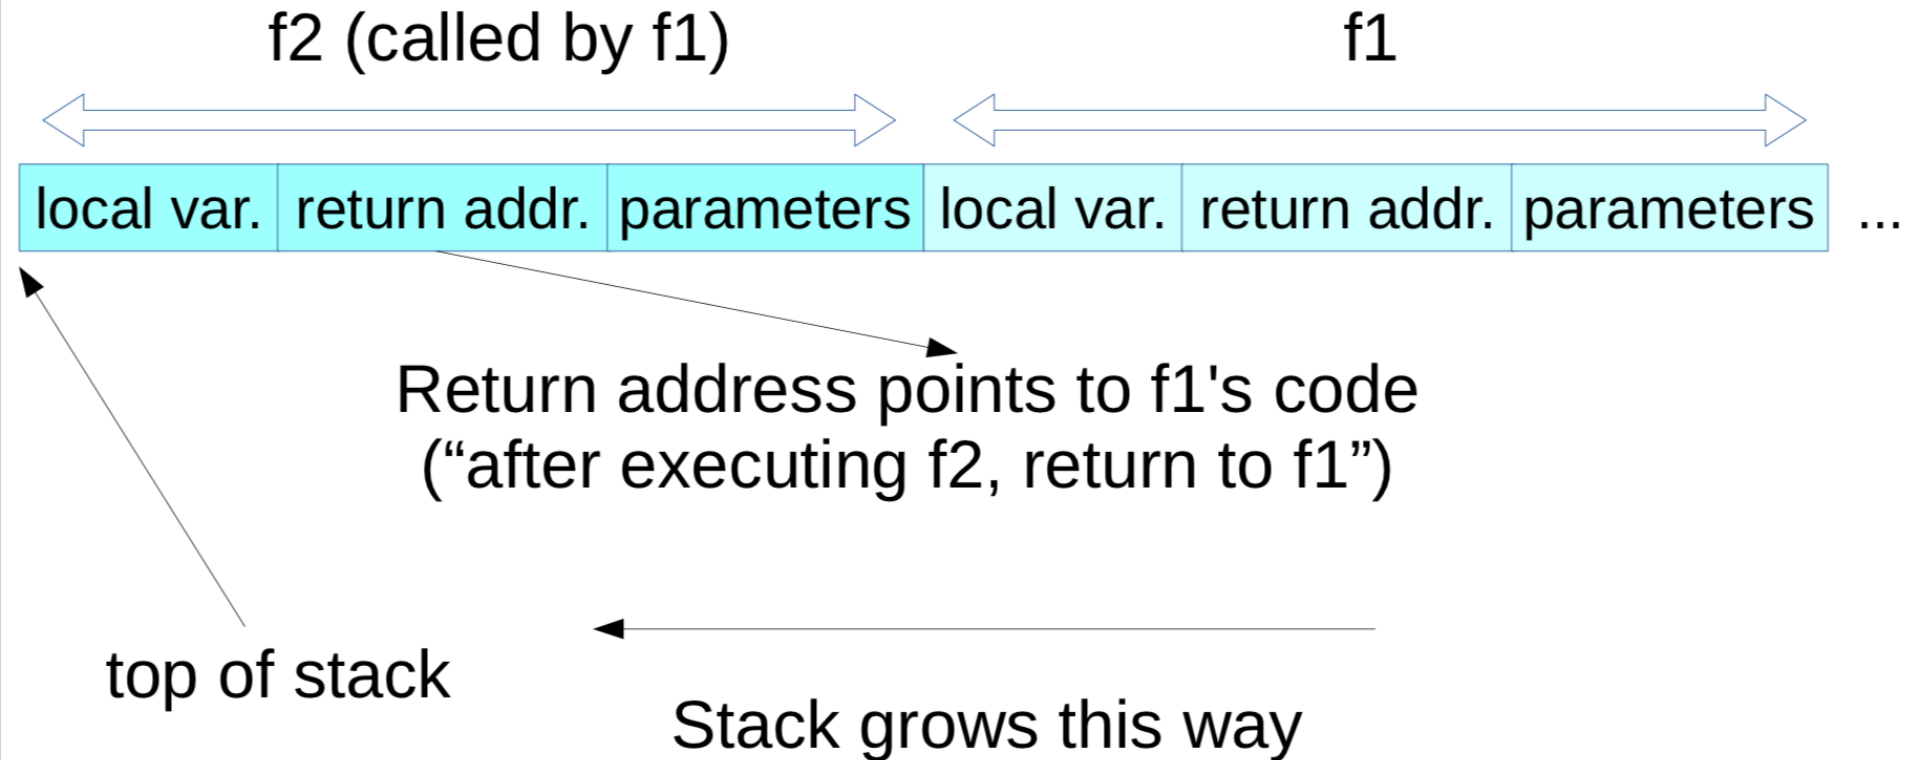
# Stack Example

```
void func(int a, int b){  
    char buffer[10];  
    return;  
}  
void main(){  
    func(1,2);  
    int c = 1 + 1;  
}
```



# A Sequence of Stack Frames

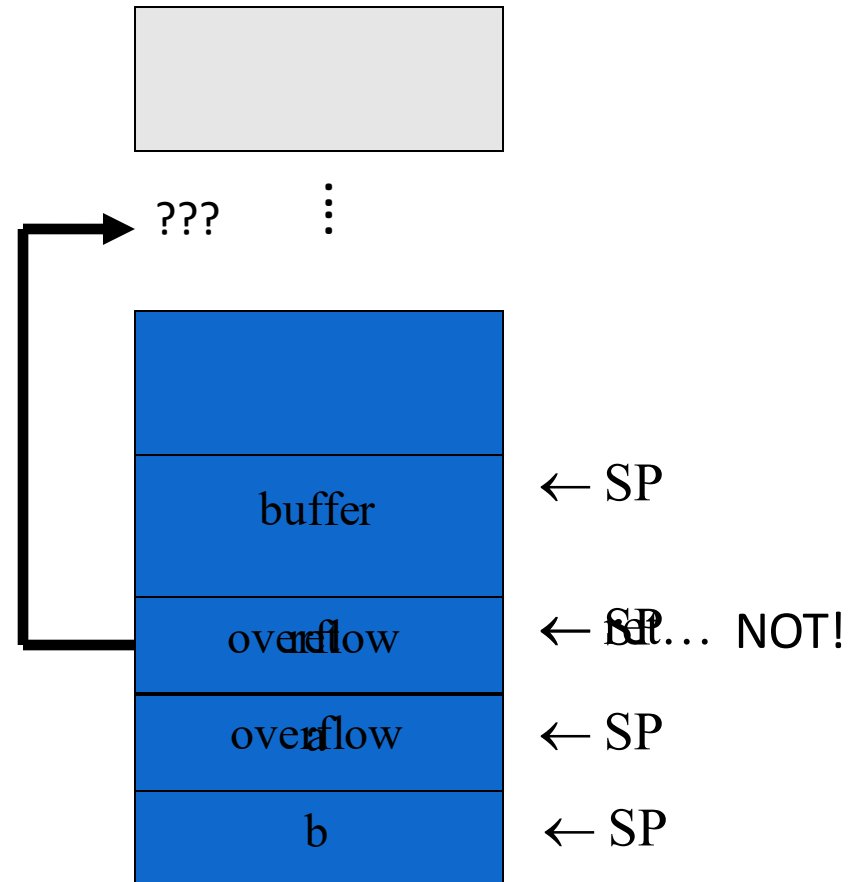
**A simplified function stack (ignored stack frame pointer)!**



# Smashing the Stack

**Ignored stack frame pointer since it does not affect our discussion!**

- ❑ What happens if buffer overflows?
- ❑ Program “returns” to wrong location
- ❑ A crash is likely

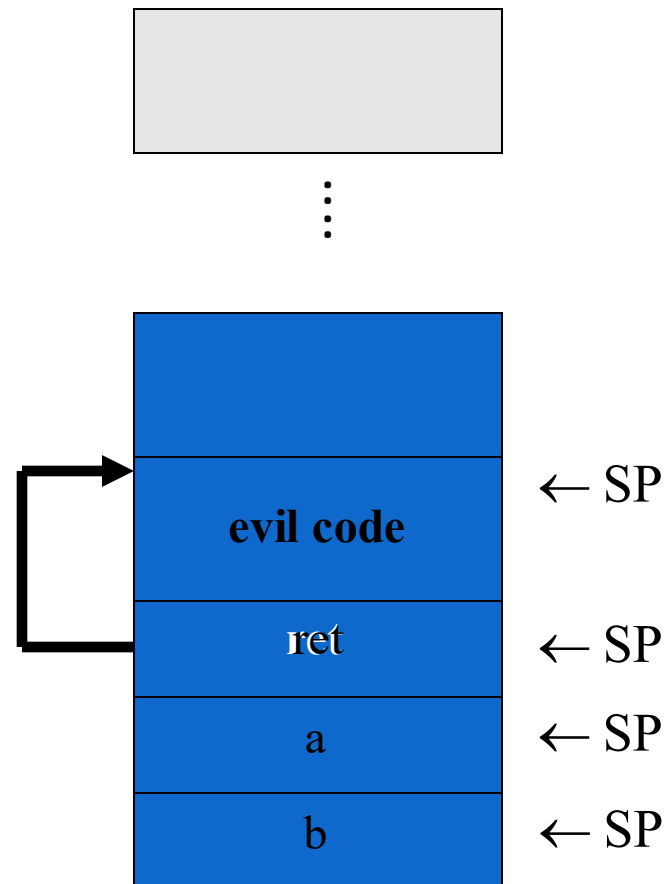




# Smashing the Stack

**Ignored stack frame pointer since it does not affect our discussion!**

- ❑ An even better idea...
- ❑ **Code injection**
- ❑ Attacker can run arbitrary code on your machine
  - ❑ And even more ← later



# Buffer Overflow Attack

- Also known as “stack smashing”, “buffer overrun”
  - Also known as a classic example of **code injection attack**, **control-flow hijacking attack**
- A buffer overflow must exist in the code
  - But not all buffer overflows are exploitable
- If exploitable, attacker can **inject code**
- Trial and error is likely required
  - Suppose we need to guess the address of **evil code**
  - But we only need one success anyway...
- Stack overflow is “attack of the decade” ...
  - Also heap overflow and so on.

# Buffer Overflow Attack

- Buffer grows **backwards** on the stack and therefore **overwrites** the **return address**.
- Any thoughts on the mitigation...?
- What if buffer grows forwardly?
  - No, it does not solve the problem.
  - See my notes.

# Unsafe C/C++ Lib Functions

`strcpy (char *dest, const char *src)`

`gets (char *s)`

`strcat (char *dest, const char *src)`

`scanf ( const char *format, ... )`

`sprintf (const char *format, ... )`

`...`

# Buffer Overflow Attack due to Unsafe C/C++ Library Functions


- Suppose Web server contains this function

```
void func(char *str) {  
    char buf[126];  
    strcpy(buf, str);  
}
```

Allocate local buffer  
(126 bytes reserved on stack)

Copy argument into local buffer

- When this function is invoked, a new **frame** with local variables is pushed onto the stack



# Buffer Overflow Attack due to Unsafe C/C++ Library Functions

- Memory pointed to by str is copied onto stack...

```
void func(char *str) {  
    char buf[126];  
    strcpy(buf, str);  
}
```

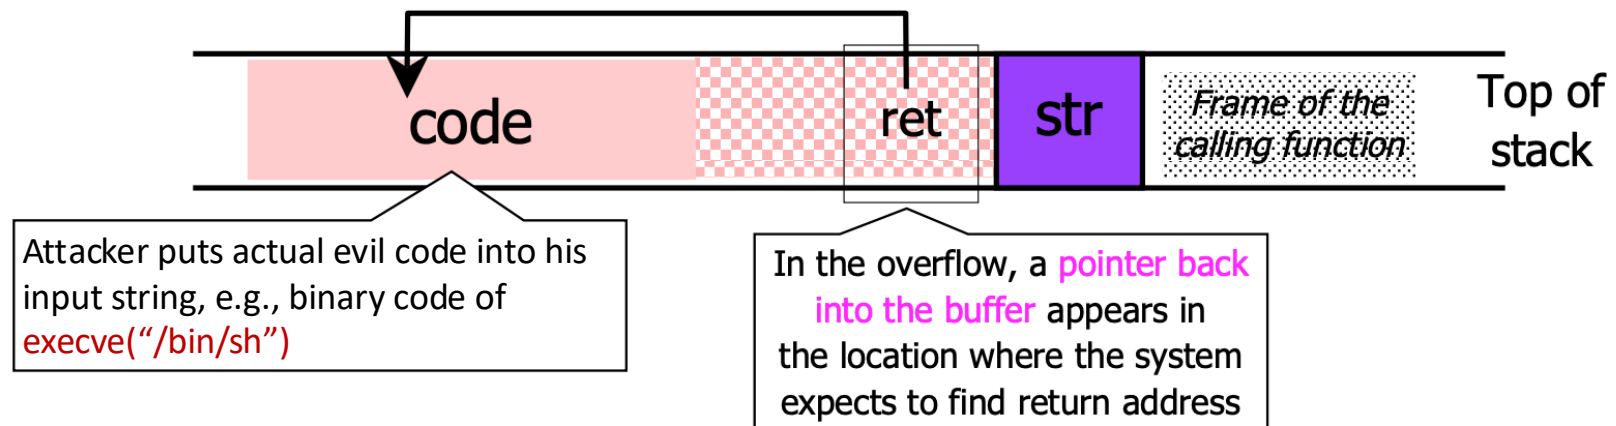
strcpy does NOT check whether the string at \*str contains fewer than 126 characters

- If a string longer than 126 bytes is copied into buffer, it will overwrite adjacent stack locations



# Buffer Overflow Attack due to Unsafe C/C++ Library Functions

- Suppose buffer contains attacker-created string
  - For example, `*str` contains a string received from the network as input to some network service daemon



- When function exits, code in the buffer will be executed, giving attacker a shell
  - **Root shell** if the victim program is `setuid root`

# Buffer Overflow Attack due to Unsafe C/C++ Library Functions

```
void input_username(...) {  
    char username[16];  
    printf("Enter username:");  
    gets(username);  
    ...  
}
```

gets does not check bounds!

[ ] [1200]

...

username[16]

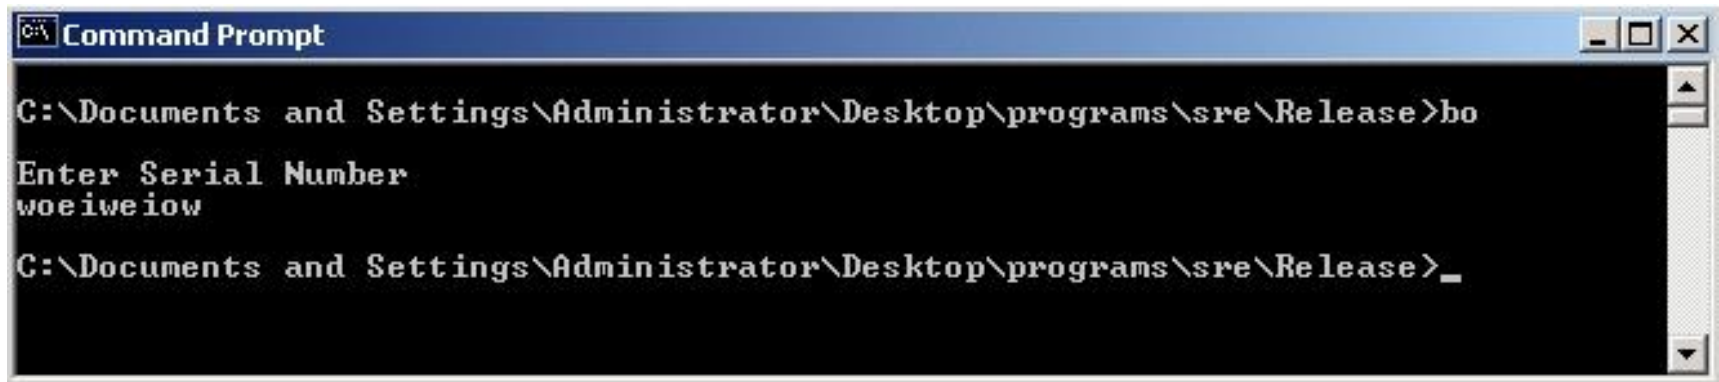
return addr.

Parameters



# Stack Overflow Attack Example

- Suppose program asks for a serial number that you do not know
- Also, you do **not** have source code
- You only has the executable (exe)




```
C:\Documents and Settings\Administrator\Desktop\programs\sre\Release>bo
Enter Serial Number
woeiweiw
C:\Documents and Settings\Administrator\Desktop\programs\sre\Release>_
```

- ❑ Program quits on incorrect serial number

# Buffer Overflow Present?

- By trial and error, attacker discovers apparent buffer overflow



- ❑ Note that 0x41 is ASCII for "A"
- ❑ Looks like **ret** overwritten by 2 bytes!

# Disassemble Code

- Next, disassemble bo.exe to find


```
.text:00401000
.text:00401000
.text:00401003
.text:00401008
.text:0040100D
.text:00401011
.text:00401012
.text:00401017
.text:0040101C
.text:0040101E
.text:00401022
.text:00401027
.text:00401028
.text:0040102D
.text:00401030
.text:00401032
.text:00401034
.text:00401039
.text:0040103E

sub     esp, 1Ch
push    offset aEnterSerialNum ; "\nEnter Serial Number\n"
call    sub_40109F
lea     eax, [esp+20h+var_1C]
push    eax
push    offset aS                ; "%S"
call    sub_401088
push    8
lea     ecx, [esp+2Ch+var_1C]
push    offset aS123n456 ; "S123N456"
push    ecx
call    sub_401050
add     esp, 18h
test    eax, eax
jnz     short loc_401041
push    offset aSerialNumberIs ; "Serial number is correct.\n"
call    sub_40109F
add     esp, 4
```

- ❑ The goal is to exploit buffer overflow to jump to address 0x401034

# Buffer Overflow Attack


- Find that, in ASCII, 0x401034 is “@^P4”



- ❑ Byte order is reversed?
  - ❑ X86 processors are “little-endian”

# Overflow Attack, Take 2

- Reverse the byte order to “4^P@a” and...



```

C:\Documents and Settings\Administrator\Desktop\programs\sre\Release>bo
Enter Serial Number
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA4^P
Serial number is correct.
C:\Documents and Settings\Administrator\Desktop\programs\sre\Release>_

```

- ❑ Success! We've bypassed serial number check by exploiting a buffer overflow
- ❑ What just happened?
  - Overwrote return address on the stack

# Buffer Overflow

- Attacker did **not** require access to the source code
- Only tool used was a disassembler to determine address to jump to
- Find desired address by trial and error?
  - Necessary if attacker **does not have exe**
  - For example, a remote attack

# Buffer Overflow Defense

- Never execute code on stack
  - W^X (write XOR Execute)
- Detect overflow
  - Use a **canary**
- Randomize stack
  - Address space layout randomization (**ASLR**)
- Don't use C/C++
  - Use **safe languages** (Java, C#)
- Use **safer C functions**
  - Additional checks on the buf length.

# Marking stack as non-execute

- Basic stack exploit can be prevented by marking stack segment as non-executable.
  - W^X (write XOR Execute)
  - Support in Windows. Code patches exist for Linux.

## Problems:

- Some software need writeable .text section or executable stack section
  - “self-modifying code”
  - Just-in-time compilation



# Randomization: Motivations

- Buffer overflow exploits need to know the (virtual) address to which pass control
  - Address of attack code in the buffer
- Same address is used on many machines
- Idea: introduce **diversity**
  - Make stack addresses unpredictable and different from machine to machine

# Address Space Layout Randomization (ASLR)


- Arranging the positions of key data areas randomly in a process' address space.
  - e.g., the base of the executable and position of the stack
  - Attacks:
    - Repetitively guess randomized address
    - Spraying injected attack code
- Windows has this enabled, software packages available for Linux and other UNIX variants

# Buffer Overflow Defense

**A simplified function stack (ignored stack frame pointer)!**

- **Canary**

- Run-time stack check
- Push canary onto stack
- Canary value:
  - Constant 0x000aff0d
  - Or, random value decided during runtime (see my example)



# Canary

- Supported by most platforms.
  - Minimal performance effects: 8% slow down for Web servers.
- Note: Canaries don't offer fool-proof protection.
  - Some attacks can leave canaries untouched.