

# CSIT 5740 Introduction to Software Security

Note set 4B

Dr. Alex LAM



DEPARTMENT OF  
**COMPUTER SCIENCE & ENGINEERING**

The set of note is adopted and converted from a software security course at the Purdue University by Prof. Antonio Bianchi

# *Address Space Layout Randomization (ASLR)*

# Guessing Addresses

- Overflow a buffer → Overwrite a pointer (e.g., return address on the stack)  
→ Make the execution jump to a memory region which:
  - 1) we control the content of
  - 2) is executable (more on this later)
  - 3) **we know where it is located**

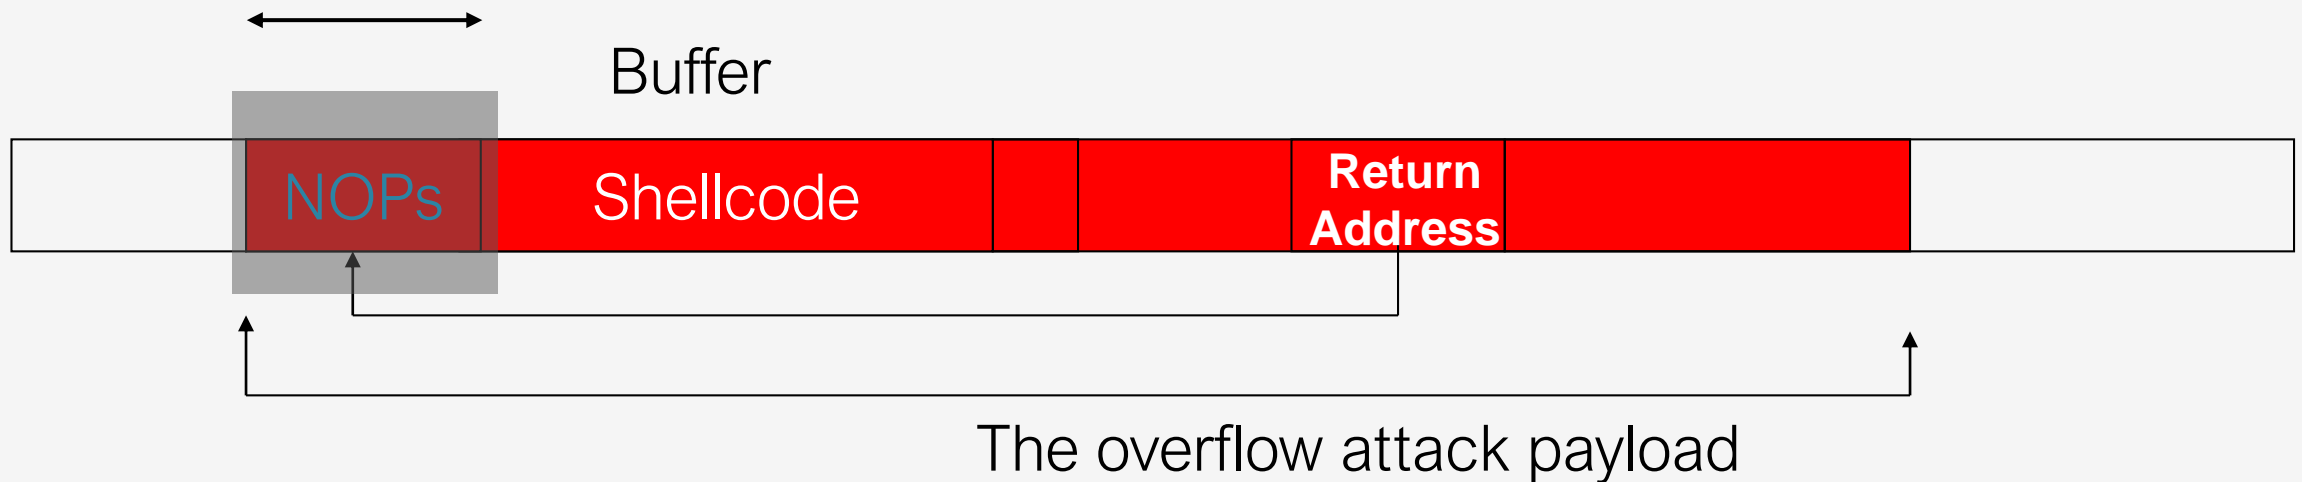
# Guessing Addresses

- In most cases the address of the buffer is not known
- It has to be “guessed” (and the guess must be VERY precise)
- The stack address of a program can be obtained by using gdb
  - For now we assume: no memory layout randomization
- Given the same environment and knowing the size of command-line parameters, the address of the stack can be roughly guessed
- However the exact value may be slightly different when a program is started by gdb or not (due to different environment variables)
- That’s why we use NOP sleds for writing the shellcode

# NOP Sled

- Just like what have mentioned at the beginning of note set 4A. To increase the probability of success, we can add, at the beginning of our controlled buffer (containing the shellcode) a series of NOP instructions
  - In Intel assembly, NOP is 0x90

**If we guess a value within this range the shellcode will be executed successfully**



# Address Space Layout Randomization (ASLR)

- An Operating System mitigation technique → makes exploitation harder
- Randomizes the position of the:
  - Heap
  - Stack
  - Dynamically-linked libraries
  - Program's main code, if compiled as position independent (PIE).  
Enabled/Disabled with compiler options `-fPIE/-no-pie`  
slightly slower

# *Address Space Layout Randomization (ASLR)*

---

- On 32-bit Linux: only 20 bits of entropy
  - Still vulnerable to brute-force attack, if unlimited attempts are possible
  - NOP-sled can help the attack significantly
- 64-bit architectures: more virtual addresses available  
→ more randomization → much more secure

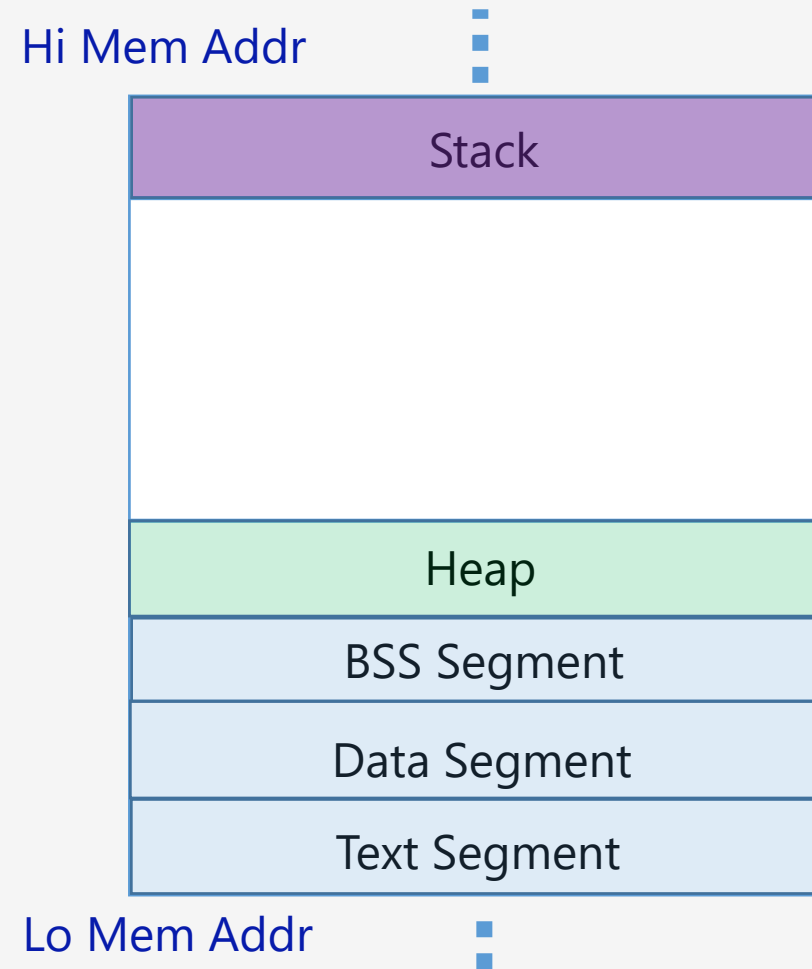
# ASLR in Linux

- ASLR can be disabled by:
  - Setting the associated kernel variable to 0
  - `echo "0" > /proc/sys/kernel/randomize_va_space`
    - Requires root
- ASLR can be enabled by (it is enabled by default)
  - `echo "2" > /proc/sys/kernel/randomize_va_space`
    - Also requires root
- Running a program in GDB
  - makes the program ignoring the setuid bit
  - disables randomization by default



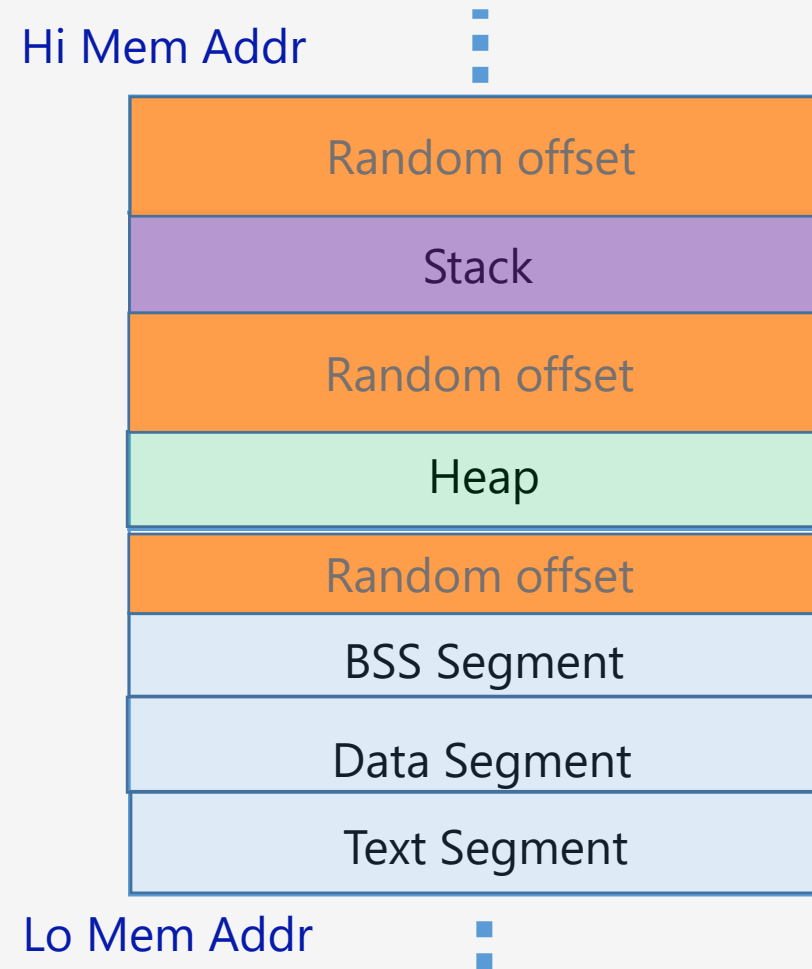
# Example: ASLR disabled

When there is no ASLR, x86 memory looks like this



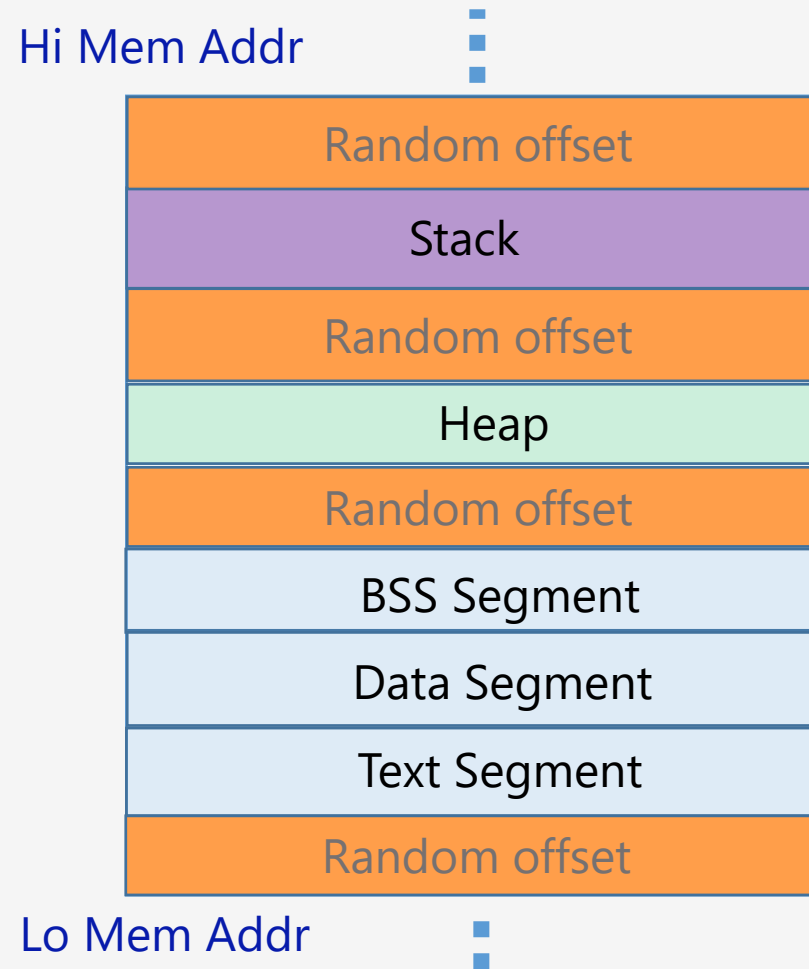
# Example: ASLR enabled, non-PIE binary

When there is there ASLR, but PIE is turned off



# Example: ASLR enabled, PIE binary

When there is there ASLR, PIE is also turned on



# ASLR Exploitation

- If the randomization is limited (32 bit)
  - Bruteforcing
  - NOP Sled
  - ...
- fork does not re-randomize the layout → it may help bruteforcing
  - e.g., in the case of a “forking server”
  - fork does not change the memory locations used by the child process → after fork, the memory layout of the child will be equal to the memory layout of the parent process

# ASLR Exploitation

- If bruteforcing is not feasible, we need “a leak”
  - For instance, we can force the program to print the value of a pointer → we can use this value to know where the code has been loaded at
  - Alternative, we can exploit a format string vulnerability (more on this later)
- If we need a leak, an exploit becomes more complex.  
At least 2 steps are required, example:
  - 1) Read memory → Leak a pointer → Compute where our controlled buffer is
    - Note that the “relative distance” between many pointers normally remain constant even in the presence of ASLR  
→ Leaking one pointer allow to recover the address of other pointers
  - 2) Write memory → Overwrite a function pointer → Jump to the controlled buffer

***Format String  
attack to leak  
information***

# Leaking Memory with printf

- `int printf(const char *format, ...);`
- Notice that “...” above? It means the printf() function takes in an variable number of arguments (i.e. variadic).
  - How does it know how many arguments are passed to it?
  - It infers from the first argument \*format, which is a string
  - For example: `printf("Hello %s, you are %d years old\n", name, age);`
- What happens when we supply just one argument instead of two?
  - `printf("Hello %s, you are %d years old\n", name);`
  - `printf()` will output the string (%s) according to the `rdi` value (the computer has copied the data in the name variable into it), and then it will output the `rsi` value (for the unsupplied %d)

# *printf Internals*

- To find out why, we must first understand what happens when we call
- `printf("%d", 1);` // How does this work?
- `printf` assumes that a parameter was passed, so it accesses the next location that should have one: in this case, `rsi` (second parameter)
- Keep in mind the calling convention:
  - The arguments of the function calls are copied to : `rdi, rsi, rdx, rcx, r8, r9, < stack. . . >`



---

`printf("%d", 1)`

<b>rdi:</b>	addr of the format string "%d"
<b>rsi:</b>	1
<b>rdx:</b>	garbage
<b>rcx:</b>	garbage
<b>r8:</b>	garbage
<b>r9:</b>	garbage

---

```
printf("%d%d", 1, 2)
```

rdi:	addr of the format string "%d%d"
rsi:	1
rdx:	2
rcx:	garbage
r8:	garbage
r9:	garbage

---

```
printf("%d%d%d", 1, 2, 3)
```

rdi:	addr of the format string "%d%d%d"
rsi:	1
rdx:	2
rcx:	3
r8:	garbage
r9:	garbage

---

```
printf("%d%d%d%d", 1, 2, 3, 4)
```

rdi:	addr of the format string "%d%d%d%d"
rsi:	1
rdx:	2
rcx:	3
r8:	4
r9:	garbage

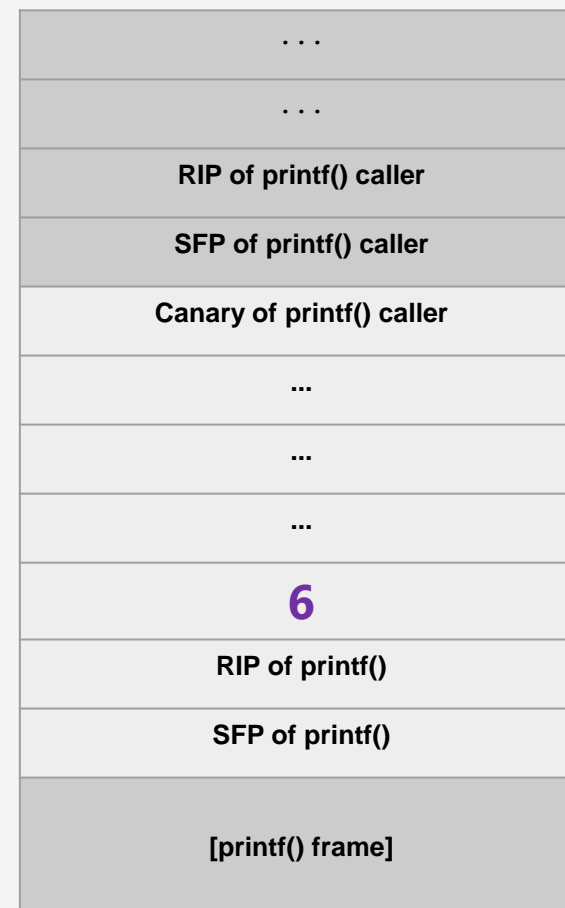
---

```
printf("%d%d%d%d%d", 1, 2, 3, 4, 5)
```

rdi:	addr of the format string "%d%d%d%d%d"
rsi:	1
rdx:	2
rcx:	3
r8:	4
r9:	5

`printf("%d%d%d%d%d%d", 1, 2, 3, 4, 5, 6)`

<b>rdi:</b>	addr of the format string "%d%d%d%d%d%d"
<b>rsi:</b>	1
<b>rdx:</b>	2
<b>rcx:</b>	3
<b>r8:</b>	4
<b>r9:</b>	5



Starting from the 6<sup>th</sup> argument, printf() starts to read from the stack

arg 6

`printf("%d%d%d%d%d%d%d", 1, 2, 3, 4, 5, 6, 7)`

<b>rdi:</b>	addr of the format string "%d%d%d%d...%d"
<b>rsi:</b>	1
<b>rdx:</b>	2
<b>rcx:</b>	3
<b>r8:</b>	4
<b>r9:</b>	5

...
...
RIP of printf() caller
SFP of printf() caller
Canary of printf() caller
...
...
7
6
RIP of printf()
SFP of printf()
[printf() frame]

Starting from the 6<sup>th</sup> argument, printf() starts to read from the stack

arg 7  
arg 6

# What if we do not supply all the args

```
printf("%d%d%d%d%d%d%d", 1, 2, 3, 4, 5, 6)
```

rdi:	addr of the format string "%d%d%d%d...%d"
rsi:	1
rdx:	2
rcx:	3
r8:	4
r9:	5

...
...
RIP of printf() caller
SFP of printf() caller
Canary of printf() caller
...
...
data on stack
6
RIP of printf()
SFP of printf()
[printf() frame]

printf() will still get to the corresponding stack position to find its arg 7

arg 7  
arg 6

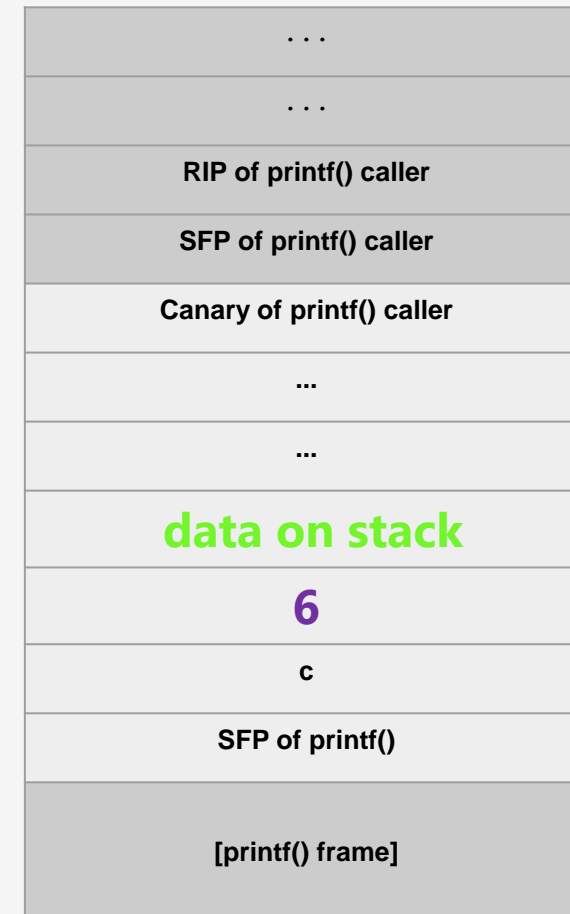


# What if we do not supply all the args

```
printf("%d%d%d%d%d%d%d", 1, 2, 3, 4, 5, 6)
```

So, the above printf will print:

123456data\_on\_stack



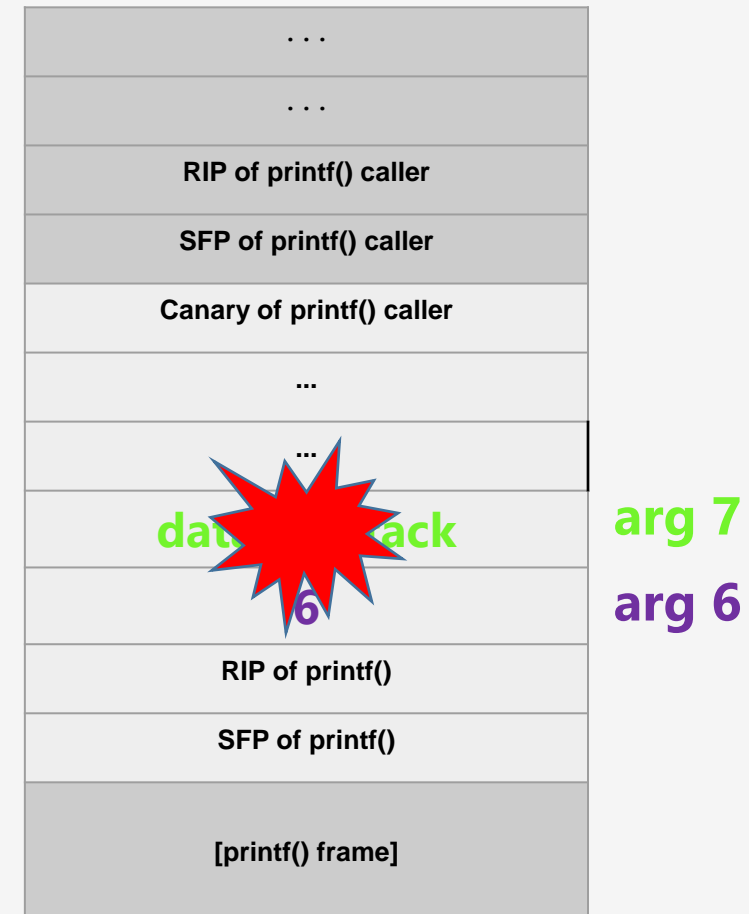
arg 7

arg 6

# What if we do not supply all the args

```
printf("%d%d%d%d%d%d%d", 1, 2, 3, 4, 5, 6)
```

So, printf can be exploited to leak important data on the stack!



# *printf: Exploitation*

- If a program passes unsanitized user input to `printf`, we can cause it to utilize values from the stack that were not intended as parameters `printf`
- This sort of attacks is called: **format string attacks**
- Using this, we can dump stack information, e.g. via `%p`:
  - 1) Create `buf` as a global char array
  - 2) `buf = "rsi: %p, rdx: %p, rcx: %p, r8: %p, r9: %p, stack: %p %p";`
  - 3) `printf(buf);`
- This will print `rsi`, `rdx`, `rcx`, `r8`, `r9`, and two stack `qwords`  
→ This allows us to leak data from the program

# *printf: Exploitation*

- If a program passes unsanitized user input to `printf`, we can cause it to utilize values from the stack that were not intended as parameters `printf`
- It does not give direct control over the program, however
  - `%p` and other formats only let us **read** from the stack, not to write
    - It can be used to leak sensitive information from the program or defeat randomization
  - `%n`, can be used to **write** data!
    - disabled in many modern libc implementations

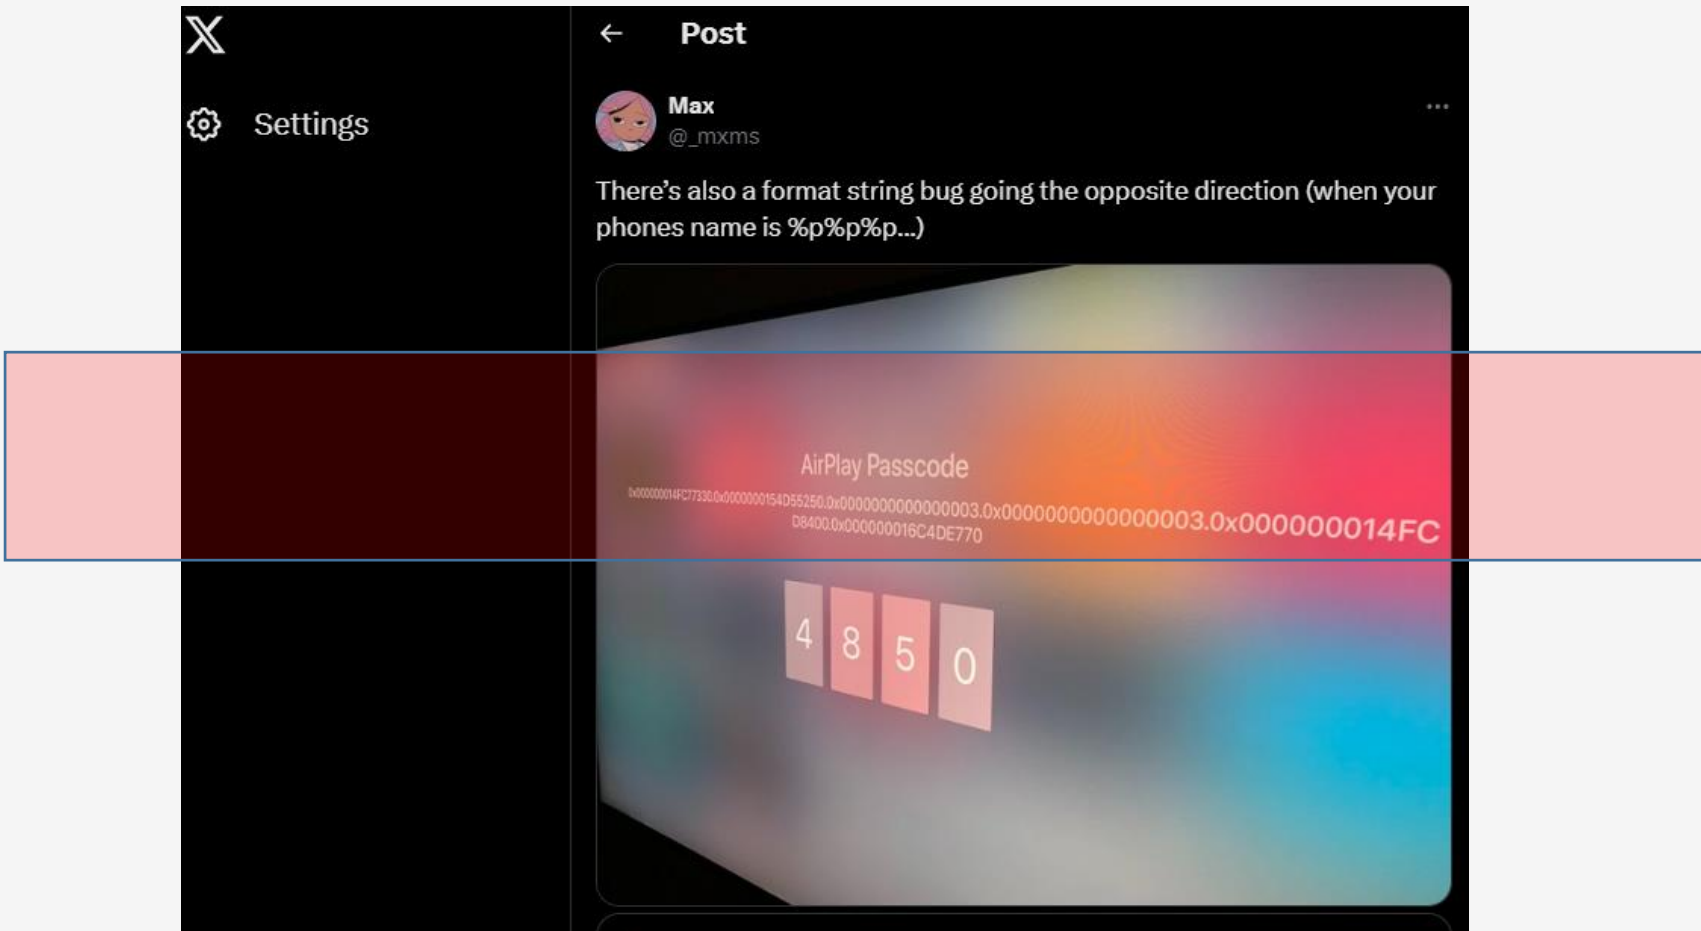
# *printf: Exploitation*

- `printf` supports the following format to print the  $n^{\text{th}}$  argument in as a 32bit/64bit pointer on 32bit/64bit machines
  - `%n$p`
  - For example: `printf("%9$p")` will print the 9<sup>th</sup> argument
  - By using the above we will show how using the format string attack to defeat a stack memory protection mechanism called “stack canary”, which was supposed to protect the system against stack smashing

# *printf: Exploitation*

- If we allow the user to supply arbitrary format string to `printf` , then the user can launch the format string attack. The following is one such scenario
  - Declare a char array, say, `buffer[]`
  - Get the user input into `buffer`
  - Let `printf` print the user input directly (i.e. `printf(buffer)`)
  - Then the user can provide strings like “`%p%p%p`” or “`%9$p`” to get the program information
- To mitigate that, one just need to use:
  - **`printf("%s", buffer)`**
  - The format string in the `buffer` will be considered static strings and will be output as the static strings “`%p%p%p`” or “`%9$p`”. It will not print any other data.

# A real world example of format string attack on Apple Airplay in 2019



# *Stack Canaries*



# Stack Canaries

- Add a **random** value before the saved returned address, and check it before the function returns
- If an overflow happens the canary's value change → the program's execution is halted before exploitation is possible
- Implemented by default in modern compilers by modifying the prologue/epilogue to add/verify the canary on the stack
  - Optimization: only used in functions using a buffer on the stack
  - Can be disabled with: `-fno-stack-protector`

# Stack Canaries

- Miners protect themselves against toxic gas buildup in the mine with a canary
  - Canary: A beautiful bird that is sensitive to toxic gas
  - If toxic gas builds up, the canary dies first
  - The miners notice that the canary has died and evacuate the mine



# Stack Canaries

- A canary value is unique every time the program **runs** but the same for all functions **within a run**
- A canary value uses a NULL byte as the first byte to mitigate string-based attacks (since it terminates any string before it ( i.e. **the last byte, or the “end” is \x00** )
  - Example: A format string vulnerability with `%s` might try to print everything on the stack
    - The null byte in the canary will mitigate the damage by stopping the print earlier.

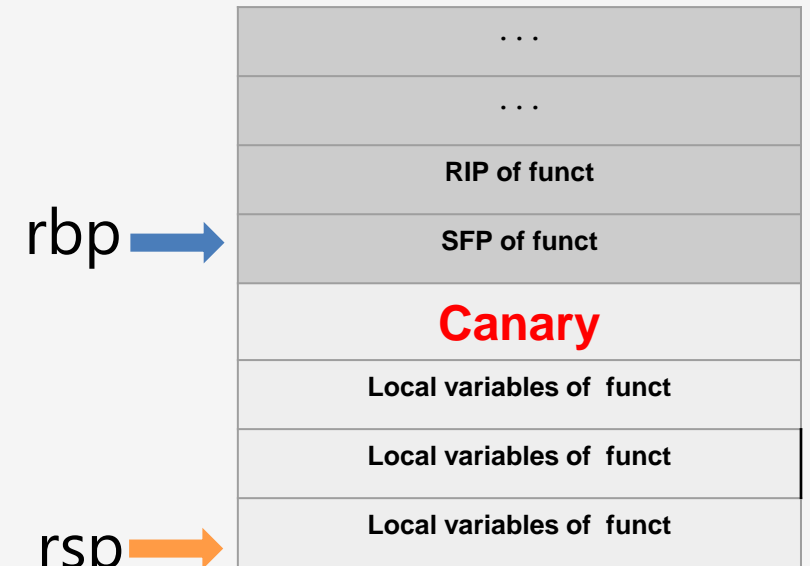
# Stack Canaries

## Prologue

```
push    rbp
mov     rbp, rsp
sub     rsp, 0x20
mov     rax, QWORD PTR fs:0x28 ; fs:0x28 contains the canary
mov     QWORD PTR [rbp-0x8], rax
```

## Epilogue

```
mov     rax, QWORD PTR [rbp-0x8]
xor     rax, QWORD PTR fs:0x28c ; if canary on stack same as the real
                                   ; canary, rax will be 0
je      400d4f; 0x400d4f is the address the leave instruction below
call    4007e0 <__stack_chk_fail@plt>; prints an error message and immediately terminates the program
leave
ret
```



# Bypassing Canaries

- Memory Leaks (using format string attack)
- Arbitrary write
  - $a[x] = y$
  - just controlling  $x$  could be enough
  - $x$  could be negative
  - $x$  could be set as huge number and “wraps around”
- Forking Servers and Bruteforcing

# Exposing Canary through format string attacks

```
void funct() {  
    char answer[24]="%6$p,%7$p,%8$p,%9$p";  
    printf(answer)  
}  
  
int main(int argc, char *argv[]) {  
  
    funct();  
    return 0;  
}
```

arg 9

arg 8

arg 7


arg 6

...
...
RIP of funct
SFP of funct
<b>Canary of funct</b>
Upper 8 bytes of answer[]
Middle 8 bytes of answer[]
Lower 8 bytes of answer[]
RIP of printf()
SFP of printf()
[printf() frame]

# Exposing Canary through format string attacks

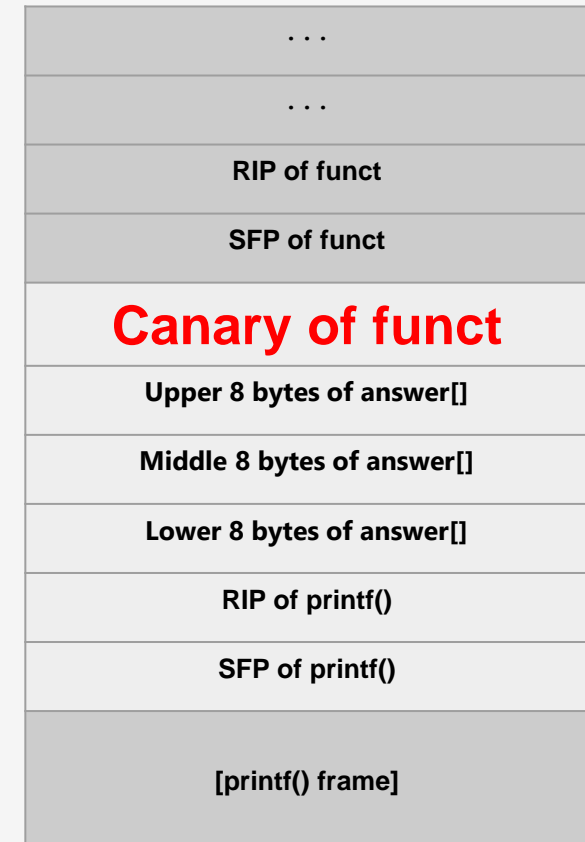
```
void funct() {  
    char answer[24]="%6$p,%7$p,%8$p,%9$p";  
    printf(answer)  
}  
  
int main(int argc, char *argv[]) {  
    funct();  
    return 0;  
}
```

arg 9  
arg 8  
arg 7  
arg 6



# Exposing Canary through format string attacks

```
void funct() {  
    char answer[24]="%6$p,%7$p,%8$p,%9$p";  
    printf(answer)  
}  
  
int main(int argc, char *argv[]) {  
    funct();  
    return 0;  
}
```



arg 9

The output from the program could be:

0x2437252c70243625, 0x252c702438252c70, 0x702439, **0xd963a0cd2106600**



# Canaries and Forking Servers

- As for ASLR, an interesting case is forking servers
  - fork does not change the layout of the memory (since it just “clones” the current process)
- the stack canary will not change when a program forks
- if we can overflow 1 byte at a time, we can try
  - to guess the canary byte-by-byte (max 256 tries per byte)
  - by checking if the child process crashed or not (e.g., the connection dropped)
  - Then in total,  $256+256+256$  guesses for 32-bit canary, and  $256+256+\dots+256 = 7 * 256 = 1792$  guesses for 64-bit canary!

# Guessing canary byte by byte (32-bit)

---

...
...
RIP of funct
SFP of funct
<b>Canary of funct</b>

# Guessing canary byte by byte (32-bit)

---

...			
...			
RIP of funct			
SFP of funct			
d9	63	a0	00

# Guessing canary byte by byte (32-bit)

...			
...			
RIP of funct			
SFP of funct			
d9	63	a0	41
41414141			
41414141			
41414141			



# Guessing canary byte by byte (32-bit)


...			
...			
RIP of funct			
SFP of funct			
d9	63	a0	42
41414141			
41414141			
41414141			



After at most 256 guesses...

# Guessing canary byte by byte (32-bit)

...			
...			
RIP of funct			
SFP of funct			
d9	63	a0	00
41414141			
41414141			
41414141			



# Guessing canary byte by byte (32-bit)

...			
...			
RIP of funct			
SFP of funct			
d9	63	41	00
41414141			
41414141			
41414141			






After at most 256 guesses...

# Guessing canary byte by byte (32-bit)


...			
...			
RIP of funct			
SFP of funct			
d9	63	a0	00
41414141			
41414141			
41414141			



After at most 256 guesses...

# Guessing canary byte by byte (32-bit)


...			
...			
RIP of funct			
SFP of funct			
d9	63	a0	00
41414141			
41414141			
41414141			



After at most 256 guesses...

# Guessing canary byte by byte (32-bit)

...			
...			
RIP of funct			
SFP of funct			
d9	63	a0	00
41414141			
41414141			
41414141			



# Canaries and Forking Servers

- Now the complete 32-bit canary is revealed
- Each byte will take at most 256 guesses
- The right most byte is always \x00, so no guess is needed (though we do show it in the animation to make it appears to be more “complete”) to guess the canary byte-by-byte (max 256 tries per byte)
  - In total, at most  $256+256+256 = 768$  guesses for 32-bit canary!